

# CO224 Lecture Notes

Computer Architecture

**Dr. Isuru Nawinne**

Department of Computer Engineering

University of Peradeniya

November 20, 2025



# Contents

Preface . . . . .	6
Learning Methods . . . . .	7
Introduction . . . . .	8
<b>1 Fundamentals</b>	<b>9</b>
1.1 Lecture 1: Computer Abstractions and Technology . . . . .	9
1.1.1 Introduction . . . . .	9
1.1.2 The Big Picture of Computer Systems . . . . .	9
1.1.3 Instruction Set Architecture (ISA) - The Key Interface . . . . .	11
1.1.4 From Problem to Execution - The Translation Chain . . . . .	11
1.1.5 Writing Programs at Different Levels . . . . .	12
1.1.6 Microarchitecture Details . . . . .	13
1.1.7 Abstraction Concept . . . . .	14
1.1.8 Performance Theme . . . . .	16
1.1.9 Key Takeaways . . . . .	16
1.1.10 Summary . . . . .	17
1.2 Lecture 2: Technology Trends, Moore's Law, and Computer System Organization	17
1.2.1 Introduction . . . . .	17
1.2.2 Moore's Law - Foundation of Computer Technology Evolution . . . . .	17
1.2.3 Technology Scaling - Historical Data . . . . .	19
1.2.4 Feature Size Scaling - Lithography Improvements . . . . .	22
1.2.5 Technology Roadmaps - ITRS Predictions . . . . .	24
1.2.6 Why Smaller Transistors Improve Performance . . . . .	27
1.2.7 Clock Rate Trends - The Power Wall . . . . .	29
1.2.8 Shift to Multi-Core Processors . . . . .	34
1.2.9 Computer System Organization - Three Layers . . . . .	41
1.2.10 From High-Level Code to Machine Code - The Translation Process . . . . .	46
1.2.11 Program Execution - Inside the CPU . . . . .	49
1.2.12 Real CPU Layout - AMD Barcelona Example . . . . .	54
1.2.13 Key Takeaways . . . . .	60
1.2.14 Summary . . . . .	61
1.3 Lecture 3: Understanding Performance . . . . .	62
1.3.1 Introduction . . . . .	62
1.3.2 Defining and Measuring Performance . . . . .	62
1.3.3 CPU Time and Performance Factors . . . . .	63
1.3.4 Understanding CPI in Detail . . . . .	65
1.3.5 Performance Optimization Principles . . . . .	66
1.3.6 Complete Performance Analysis . . . . .	67
1.3.7 Practical Performance Considerations . . . . .	69
1.3.8 Key Takeaways . . . . .	70
1.3.9 Summary . . . . .	71

<b>2</b>	<b>ARM Assembly Programming</b>	<b>73</b>
2.1	Lecture 4: Introduction to ARM Assembly	73
2.1.1	Introduction	73
2.1.2	ARM Architecture Overview	73
2.1.3	ARM Instruction Format	76
2.1.4	Basic ARM Instructions	77
2.1.5	Memory Access Instructions	79
2.1.6	Assembly Program Structure	80
2.1.7	ARM Development Tools	81
2.1.8	Programming in ARM Assembly	83
2.1.9	Key Takeaways	84
2.1.10	Summary	85
2.2	Lecture 5: Number Representation and Instruction Encoding	85
2.2.1	Introduction	85
2.2.2	Number Representation Systems	85
2.2.3	ARM Instruction Encoding	88
2.2.4	Logical Operations	92
2.2.5	Practical Bit Manipulation Examples	96
2.2.6	Key Takeaways	97
2.2.7	Summary	98
2.3	Lecture 6: Branching and Control Flow	98
2.3.1	Introduction	98
2.3.2	Fundamentals of Conditional Execution	98
2.3.3	Comparison Instructions	100
2.3.4	Conditional Branch Instructions	101
2.3.5	Labels in Assembly	103
2.3.6	Implementing Control Structures	104
2.3.7	Array Access in Loops	107
2.3.8	PC-Relative Addressing	108
2.3.9	Conditional Execution (Alternative to Branching)	109
2.3.10	Basic Blocks	111
2.3.11	Key Takeaways	112
2.3.12	Summary	112
2.4	Lecture 7: Function Call and Return	112
2.4.1	Introduction	113
2.4.2	Function Calling Fundamentals	113
2.4.3	ARM Register Conventions	114
2.4.4	Function Call Instructions	115
2.4.5	Parameter Passing	116
2.4.6	Return Values	117
2.4.7	The Stack	118
2.4.8	Stack Operations	119
2.4.9	Register Preservation	120
2.4.10	Nested Function Calls (Non-Leaf Functions)	121
2.4.11	Recursion Example: Factorial	123
2.4.12	Memory Layout and Stack vs. Heap	124
2.4.13	Key Takeaways	126
2.4.14	Summary	126
2.5	Lecture 8: Memory Access and String Operations	126
2.5.1	Introduction	126
2.5.2	Character Data and Encoding	127

2.5.3	Byte Load/Store Operations	128
2.5.4	Half-Word Load/Store Operations	130
2.5.5	String Copy Example (strcpy)	131
2.5.6	Library Functions: scanf and printf	133
2.5.7	Compilation, Linking, and Loading	135
2.5.8	Exercises	140
2.5.9	Key Takeaways	141
2.5.10	Summary	141
<b>3</b>	<b>Processor Design</b>	<b>143</b>
3.1	Lecture 9: Microarchitecture and Datapath Design	143
3.1.1	Introduction	143
3.1.2	Course Context and MIPS ISA	143
3.1.3	Digital Logic Review	145
3.1.4	CPU Execution Stages	149
3.1.5	R-Type Instruction Datapath	152
3.1.6	I-Type Instruction Datapath	153
3.1.7	Load/Store Instruction Datapath	154
3.1.8	Branch Instruction Datapath	156
3.1.9	Complete Single-Cycle Datapath	157
3.1.10	Key Takeaways	160
3.1.11	Summary	161
3.2	Lecture 10: Processor Control	161
3.2.1	Introduction	161
3.2.2	Control Unit Overview	161
3.2.3	ALU Operations for Different Instructions	162
3.2.4	ALU Control Signal	164
3.2.5	Two-Stage ALU Control Generation	164
3.2.6	Main Control Signals	166
3.2.7	Control Signal Truth Table	169
3.2.8	Control Unit Implementation	172
3.2.9	Why Separate MemRead and MemWrite?	173
3.2.10	Complete Datapath with Control	175
3.2.11	Key Takeaways	176
3.2.12	Summary	177

## **Preface**

This document contains comprehensive lecture notes for CO224 - Computer Architecture course, covering fundamental concepts of computer systems, ARM assembly programming, processor design, memory hierarchy, and multiprocessor systems.

## Learning Methods

These lecture notes are designed to support active learning. Students are encouraged to:

- Read the material before class
- Take additional notes during lectures
- Practice with the provided examples
- Complete assignments and exercises

## Introduction

Computer Architecture is the science and art of selecting and interconnecting hardware components to create computers that meet functional, performance, and cost goals. This course explores the design principles and implementation of modern computer systems.



# Chapter 1

## Fundamentals

### 1.1 Lecture 1: Computer Abstractions and Technology

*By Isuru Nawinne*

#### 1.1.1 Introduction

This lecture introduces the fundamental concepts of computer system abstractions, exploring the relationship between hardware and software while providing an overview of the lecture series structure and topics. We examine how computer systems are built as hierarchies of abstractions, each hiding complexity while providing services to the levels above.

#### 1.1.2 The Big Picture of Computer Systems

##### Cross-Section of a Computer System (Top to Bottom)

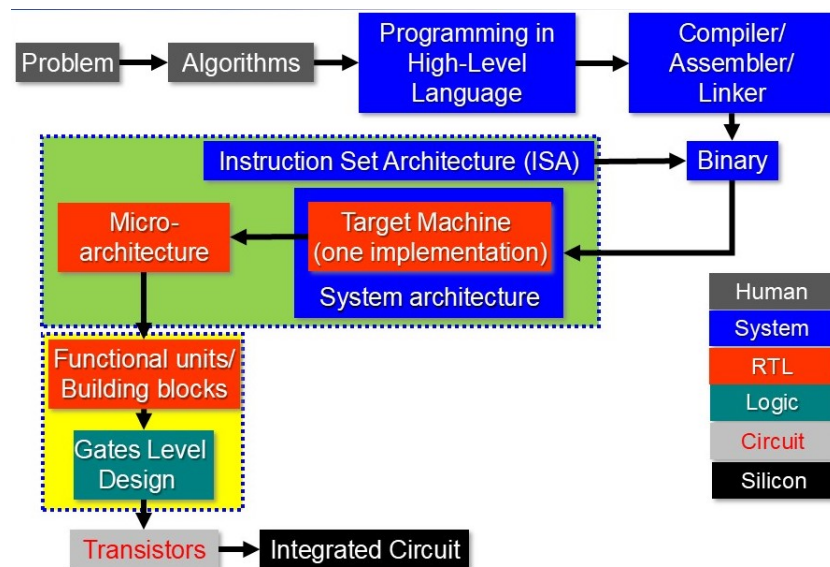


Figure 1.1: Computer System Abstraction Layers

*The diagram above illustrates the complete hierarchy from problems and algorithms at the human level, through the compilation toolchain (Compiler/Assembler/Linker), down to the ISA, microarchitecture (RTL), functional units, logic gates, transistors, and finally the silicon substrate. Each colored layer represents a different abstraction level.*

**Human-Related Level (Gray)**

- **Problems:** Real-world challenges to be solved
- **Algorithms:** Step-by-step solutions to problems
- **Programming Languages:** Tools to express algorithms

**System Level (Blue)**

- **Compilers:** Translate high-level code to assembly
- **Assemblers:** Convert assembly to machine code
- **Linkers:** Combine programs with libraries
- **Instruction Set Architecture (ISA):** The hardware-software interface

**RTL (Register Transfer Level) - Red/Orange**

- **Microarchitecture:** The processor's internal organization
- **Functional Units:** Building blocks that perform operations

**Logic Level (Green)**

- **Gate-level circuits:** Digital logic implementations
- **Logic gates:** AND, OR, NAND, NOR, XOR, etc.

**Circuit Level (Light Gray)**

- **Transistors:** BJT, CMOS devices
- **Voltage levels and currents:** Electrical signals

**Substrate Level (Black)**

- **Semiconductors:** Base materials
- **P-type and N-type semiconductors:** Doped materials
- **Electron currents:** Physical phenomena

**Purpose of Computer Systems**

- Built to solve problems (like any engineering system)
- Process: **Problems** → **Algorithms** → **Programs** → **Machine Code** → **Execution**
- Each level provides services to the level above
- Each level hides complexity from the level above

### 1.1.3 Instruction Set Architecture (ISA) - The Key Interface

#### What is an ISA?

##### Definition:

- A specification defining what the computer will understand
- Contains a list of basic instructions the processor can execute
- Examples: ARM version 8, MIPS, x86
- The critical interface between hardware and software

#### Example Instructions in an ISA

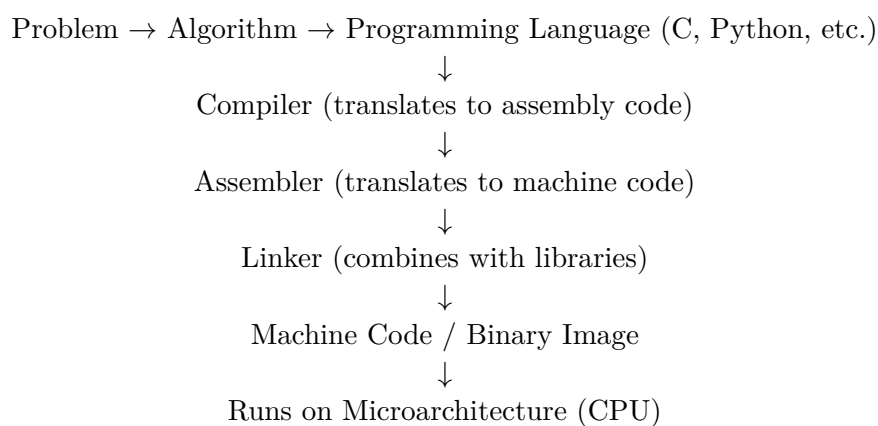
- Add two numbers together
- Subtract one number from another
- Multiply two numbers
- Load a number from memory into CPU
- Store a number from CPU into memory
- All basic operations are well-defined in the ISA

#### Importance of ISA

- Microarchitecture is built to support a specific ISA
- Programs must be written using instructions from the target ISA
- Compilers translate high-level code to ISA instructions
- ISA is the key point combining software with hardware

### 1.1.4 From Problem to Execution - The Translation Chain

#### High-Level Process



## Tool Chain Components

### Compiler

- **Function:** Converts high-level language to assembly language
- **Complexity:** Complex task requiring optimization
- **Optimizations:** Performance and memory optimizations
- **Example:** ARM GCC compiler for ARM processors

### Assembler

- **Function:** Converts assembly to machine code
- **Integration:** Built into the tool chain
- **Output:** Produces binary image (ones and zeros)

### Linker

- **Function:** Combines program with libraries
- **Output:** Creates final executable
- **Process:** Resolves external references

## Architecture-Specific Compilation

- If targeting ARM processor: Use ARM toolchain
- If targeting MIPS processor: Use MIPS toolchain
- Machine code is specific to the target ISA
- Cannot run ARM code on MIPS processor directly

### 1.1.5 Writing Programs at Different Levels

#### Machine Code (Binary)

##### Characteristics:

- Ones and zeros
- Directly executable by processor
- Very difficult for humans to write
- Error-prone and time-consuming

#### Assembly Language

##### Characteristics:

- Textual representation of machine instructions
- Example: "ADD R1, R2, R3" instead of binary
- One-to-one mapping with machine code
- Easier than machine code but still difficult for large programs
- Used in CO224 labs for ARM assembly programming

## High-Level Languages (C, Python, etc.)

### Characteristics:

- Easier to write and understand
- Good for large programs and general-purpose applications
- Requires compiler to translate to assembly/machine code
- Provides abstractions hiding hardware details

### 1.1.6 Microarchitecture Details

#### What is Microarchitecture?

##### Definition:

- A digital logic circuit built to support a given ISA
- Processes binary image (machine code)
- Understands meaning of ones and zeros
- Performs operations in actual hardware

#### Hierarchy of Microarchitecture Components

##### Microarchitecture Level

- Manipulates instructions
- Built using functional units and gate-level logic

##### Functional Units Level

- **Purpose:** Manipulates numbers
- **Examples:**
  - Adders (ripple carry, half adders, full adders)
  - Multiplexers
  - Encoders
  - Decoders
- Built using logic gates

##### Logic Gate Level

- **Purpose:** Manipulates logic levels (1s and 0s, HIGH and LOW)
- **Gates:** AND, OR, NAND, NOR, XOR, NOT
- Built using transistors

**Transistor Level**

- **Purpose:** Manipulates voltages and currents
- **Types:** BJT, CMOS
- Built using semiconductors

**Semiconductor Level**

- Deals with electron currents
- P-type and N-type semiconductors
- Combined to create transistors

**1.1.7 Abstraction Concept****What is an Abstraction?****Key Principles:**

- A conceptual entity hiding internal details
- Provides interface to higher levels
- Hides complexity underneath
- Each level doesn't worry about details above or below
- Encapsulates details and defines specific characteristics

**Hardware Abstraction Hierarchy (Bottom to Top)****1. Substrate (Silicon, Germanium)**

- Base semiconductor material

**2. Transistors**

- Built using semiconductor substrate
- Deal with voltage levels

**3. Logic Gates**

- Built using transistors
- Deal with logic levels (HIGH/LOW, 1/0)

**4. Functional Units**

- Built using logic gates
- Deal with numbers
- Examples: Adders, multiplexers

## 5. Microarchitecture

- Built using functional units and logic elements
- Deals with instructions
- Understands machine instructions

## Software Abstraction Hierarchy (Bottom to Top)

### 1. Machine Instructions (Binary)

- Ones and zeros
- Collection of logic levels
- Executable by microarchitecture

### 2. Assembly Instructions

- Textual representation of machine code
- One-to-one mapping with machine instructions
- Easier for humans to read

### 3. Programs / Source Code

- Written in high-level languages
- Collections of instructions
- Represent algorithms

### 4. Algorithms and Data Structures

- Conceptual entities
- Represent solutions to problems
- Highest level abstraction

## Relationships Between Hardware and Software Abstractions

### Voltage Levels $\leftrightarrow$ Logic Levels

- **Logic 1:** Higher voltage range (e.g., 4-5V)
- **Logic 0:** Lower voltage range (e.g., 0-1V)
- Ranges depend on transistor type (TTL vs CMOS)

### Logic Levels $\leftrightarrow$ Numbers

- Numbers represented as strings of binary digits
- Collections of logic levels form numbers

**Numbers  $\leftrightarrow$  Instructions**

- Instructions represented as binary numbers
- Microarchitecture interprets these numbers

**Summary of Relationships**

- **Transistors**  $\leftrightarrow$  Voltages (deal with)
- **Logic Gates**  $\leftrightarrow$  Logic Levels (deal with)
- **Functional Units**  $\leftrightarrow$  Numbers (deal with)
- **Microarchitecture**  $\leftrightarrow$  Instructions (understands)

**Complete System**

- All abstractions together create "the computer"
- Can deconstruct algorithm down to voltage levels
- Can deconstruct microarchitecture down to silicon
- Tight coupling between hardware and software abstractions
- Computer systems are everywhere due to these abstractions

**1.1.8 Performance Theme****Throughout the Lecture Series**

Performance is a recurring theme that will be touched upon in every topic:

- How efficiently can CPU do things?
- How fast can operations be performed?
- How can performance be improved?
- Hardware-based improvements
- Software-based improvements

**1.1.9 Key Takeaways**

1. Computer systems are built as hierarchies of abstractions
2. Each abstraction level hides complexity and provides services to levels above
3. Instruction Set Architecture (ISA) is the critical interface between hardware and software
4. Hardware hierarchy: Substrate  $\rightarrow$  Transistors  $\rightarrow$  Gates  $\rightarrow$  Functional Units  $\rightarrow$  Microarchitecture
5. Software hierarchy: Machine Code  $\rightarrow$  Assembly  $\rightarrow$  Programs  $\rightarrow$  Algorithms
6. Tight coupling exists between hardware and software abstractions
7. Voltages  $\rightarrow$  Logic Levels  $\rightarrow$  Numbers  $\rightarrow$  Instructions (relationships between levels)



## 1.2. LECTURE 2: TECHNOLOGY TRENDS, MOORE'S LAW, AND COMPUTER SYSTEM ORGANIZATION

8. Covers ISA, microarchitecture, memory hierarchy, and system organization
9. Labs involve ARM assembly programming and building processor using Verilog
10. Understanding the complete system picture is essential for computer engineers
11. All computer systems, regardless of complexity, are built on these fundamental abstractions
12. Performance optimization is a central theme throughout the lecture series

### 1.1.10 Summary

Computer systems represent one of the most sophisticated examples of hierarchical abstraction in engineering. From the physical movement of electrons in semiconductors to high-level programming languages, each layer builds upon and hides the complexity of the layers below. The Instruction Set Architecture serves as the critical bridge between hardware and software, enabling programmers to write code without worrying about transistor-level details while allowing hardware designers to optimize implementations without breaking software compatibility.

Throughout this lecture series, we will explore these abstractions in depth, learning not just what they are, but why they exist and how they enable the remarkable computing capabilities we rely on every day. By understanding both hardware and software perspectives, computer engineers gain the ability to design, optimize, and innovate across the entire computing stack.

## 1.2 Lecture 2: Technology Trends, Moore's Law, and Computer System Organization

*By Dr. Isuru Nawinne*

### 1.2.1 Introduction

The evolution of computer technology over the past 50 years has been nothing short of revolutionary. From room-sized scientific calculators to powerful smartphones in our pockets, this transformation has been guided by a prediction made by Intel co-founder Gordon Moore. This lecture examines the technological trends that enabled this revolution, the physical limitations that eventually constrained traditional scaling approaches, and the architectural innovations that emerged in response.

We will trace the exponential growth in transistor density, explore how smaller feature sizes enabled both more complex circuits and faster operation, understand why clock frequencies stopped increasing around 2004, and see how the industry pivoted to multi-core architectures. Finally, we'll examine how computer systems are organized into three layers (hardware, system software, and application software) and follow the complete translation process from high-level code to binary execution.

### 1.2.2 Moore's Law - Foundation of Computer Technology Evolution

#### Who Was Gordon Moore?

#### Background and Influence:

- Co-founder of Intel Corporation, historically the biggest manufacturer of computer chips/processors
- Most personal computers and high-end servers use Intel processors

- Made a prediction that shaped the entire semiconductor industry

**Intel's Dominance:**

- Established industry standards for processor design
- Set pace for computational advancement
- Influenced competing manufacturers
- Created benchmark for technology expectations

**Moore's Law Definition****The Prediction:**

Moore's Law is NOT a physical law like the law of gravity. It is an observation and prediction about technology trends:

**"The number of transistors that can be placed on a standard computer chip will double every two years."**

**Practical Interpretation:**

- Roughly translates to: Computational power doubles every two years
- Started in the 1950s and held true for many decades
- Based on continuous demand for increasing computational power
- Self-fulfilling prophecy driven by industry investment

**Historical Context:**

- Initial observation made in mid-1960s
- Revised and refined over subsequent decades
- Became guiding principle for semiconductor industry
- Influenced research priorities and manufacturing investments

**Impact of Moore's Law****Computer Evolution Enabled:**

Computers transformed from room-sized scientific calculators to:

- **Personal Computers:** Desktop and laptop systems in every home
- **Mobile Devices:** Smartphones with computational power exceeding 1990s supercomputers
- **Embedded Systems:** Computational intelligence in everyday objects
- **Wearables:** Smartwatches and fitness trackers

**Revolutionary Applications:**

Moore's Law made computationally intensive applications possible:

1. **Human Genome Decoding:**

## 1.2. LECTURE 2: TECHNOLOGY TRENDS, MOORE'S LAW, AND COMPUTER SYSTEM ORGANIZATION

- Massive computational requirements
- Processing billions of genetic sequences
- Pattern recognition across enormous datasets

### 1. World Wide Web and Internet Search:

- Millisecond response times for complex queries
- Indexing billions of web pages
- Real-time information retrieval

### 1. Artificial Intelligence and Machine Learning:

- Neural networks with billions of parameters
- Real-time image and speech recognition
- Autonomous systems and decision-making

### 1. Complex Simulations and Scientific Computing:

- Weather prediction and climate modeling
- Molecular dynamics simulations
- Astrophysical calculations

### Societal Impact:

- Computer software became ubiquitous and unavoidable
- Changed how we work, communicate, and learn
- Enabled digital transformation of industries
- Created new fields and destroyed old ones

### 1.2.3 Technology Scaling - Historical Data

#### Transistor Count Growth (1970-2010)

##### Chart Analysis:

The historical data shows remarkable consistency with Moore's prediction:

- **Vertical Axis:** Number of transistors ( $10^5$  to  $10^9$  - millions to billions)
- **Horizontal Axis:** Time period (1970 to 2010)
- **Blue Dotted Line:** Doubling every 18 months (aggressive prediction)
- **Red-Brown Line:** Doubling every 24 months (Moore's actual prediction)

##### Real Intel Processor Models:

Tracking actual transistor counts across processor generations:

- **Early Processors:** 4004, 8008, 8080 (thousands of transistors)




Figure 1.2: Moore's Law: Transistor Count Growth

- **1989 Milestone - 8086:** Crossed 1 million transistors
- **Middle Era:** Pentium and Itanium series
- **2008 Achievement:** Crossed 1 billion transistors (quad-core processors)
- **Trend Validation:** Actual counts closely followed the "doubling per 2 years" curve

#### Significance:

- Prediction held remarkably accurate for 40+ years
- Enabled long-term planning for semiconductor industry
- Guided investment in manufacturing technology
- Set performance expectations for consumers

### The x86 Architecture

#### Origin and Naming:

- **8086 Processor:** First x86 architecture processor (notice "86" in the name)
- Established instruction set architecture (ISA) standard
- Created foundation for backward compatibility

#### x86 Architecture Family:

The architecture evolved through multiple generations while maintaining compatibility:

- **80286:** Enhanced memory management
- **80386:** True 32-bit processor (author's first computer in 1993, 16 MHz)

## 1.2. LECTURE 2: TECHNOLOGY TRENDS, MOORE'S LAW, AND COMPUTER SYSTEM ORGANIZATION

- **80486:** Integrated floating-point unit
- **Pentium Series:** Brand name change, performance leap
- **Modern Processors:** Core i3, i5, i7, i9 series

### **AMD's Adoption:**

- AMD also uses x86 architecture
- Compatible instruction set
- Competitive alternative to Intel
- Drives innovation through competition

### **Evolution Strategy:**

- Architecture evolved significantly over decades
- Maintained backward compatibility throughout
- Old programs run on new processors
- Balanced innovation with stability

## **Historical Context**

### **Early Computing Era (1985-1990):**

- 1985: 80386 computers first arrived on market
- No graphical user interfaces (GUIs) existed
- Black screen with text-only displays
- DOS operating system (text-based console)
- Command-line interaction only

### **Transformation Period (Mid-Late 1990s):**

- GUIs emerged (Windows 95 and similar systems)
- Point-and-click interfaces replaced command lines
- Multimedia capabilities became standard
- Internet connectivity became widespread

### **User Experience Revolution:**

- Significant transformation in how people interacted with computers
- Democratized computing beyond technical experts
- Enabled productivity for non-technical users
- Set expectations for modern computing

### 1.2.4 Feature Size Scaling - Lithography Improvements

#### What Made Transistor Count Increase Possible?

##### The Answer: Smaller Transistors

The exponential growth in transistor count was enabled primarily by reducing transistor size through improved manufacturing processes.

##### Lithography Process:

- Etching transistors onto silicon wafer using photolithographic techniques
- Patterns created using light masks and photosensitive materials
- Feature size: Measure of transistor dimensions in nanometers (nm)
- Smaller features = more transistors per unit area

##### Feature Size Timeline:

The relentless march toward smaller dimensions:

- **2004:** 90 nanometer manufacturing process
- **2006:** 65 nanometer
- **2008:** 45 nanometer (very famous generation, many developments)
- **Continuing:** 32 nm, 22 nm processes
- **2013 Actual:** 22 nm achieved
- **2015 Target:** 16 nm achieved
- **2019 Target:** 12 nm achieved
- **2023 Target:** 7 nm achieved
- **2028 Target:** 5 nm exceeded
- **Future Roadmap:** 3nm and 2nm are currently in production, future is 1nm and sub-1nm

#### What is "Feature Size"?

##### Original Definition:

- Originally represented physical measurement: minimum distance between source and drain of transistor
- Also called channel width, gate size, or half-pitch
- Directly related to transistor dimensions

##### Modern Reality:

- **NOT a precisely defined physical measurement anymore**
- More of a **marketing term** in current usage
- General measure of manufacturing process advancement
- Smaller number suggests more advanced technology

## 1.2. LECTURE 2: TECHNOLOGY TRENDS, MOORE'S LAW, AND COMPUTER SYSTEM ORGANIZATION

### Alternative Names:

Different terms referring to approximately the same concept:

- Gate size
- Channel width
- Half-pitch
- Process node
- Technology node

### Why Ambiguity Developed:

- Manufacturing processes became more complex
- Multiple dimensions define transistor performance
- 3D structures don't have simple linear measurements
- Marketing convenience over physical precision

### How Tiny Are Transistors?

#### Mind-Boggling Scale:

Putting modern transistor sizes in perspective:

- **45 nanometer technology:** Can fit **30 million transistors** on the head of a pin
- **Across human hair:** Over **1,000 transistors** fit across the width of a single human hair
- **Comparison to past:** Incredibly small compared to transistors 40-50 years ago

#### Manufacturing Precision:

- Requires cleanroom environments cleaner than surgical suites
- Dust particle can destroy multiple chips
- Atomic-level precision required
- Remarkable engineering achievement

### Transistor Structure

#### Basic Components:

- **Silicon Substrate:** Base semiconductor material
- **Source and Drain:** Two metal contacts on either side
- **Gate:** Control electrode positioned between source and drain
- **Insulator:** Separates gate from channel

#### Feature Size Definition:

- Distance between drain and source (channel width)
- Critical dimension for transistor operation
- Determines electrical characteristics

**Electrical Properties:**

- **Capacitance Load:** Inherent property based on semiconductor material and structure
- Affects switching speed and power consumption
- Function of transistor geometry and materials
- Critical parameter for circuit performance

### 1.2.5 Technology Roadmaps - ITRS Predictions

**ITRS Organization****International Technology Roadmap for Semiconductors:**

- **Established:** Around 2001
- **Purpose:** Predict feature size scaling for next 10 years
- **Membership:** Major semiconductor manufacturers and research institutions
- **Methodology:** Based on technology capabilities and market demand

**Prediction Basis:**

The roadmaps considered multiple factors:

- Demand for computational power
- Available manufacturing technology
- Potential technological improvements
- Economic feasibility
- Physical limitations

**Regular Updates:**

- Produced updated roadmaps regularly
- Adjusted predictions based on actual progress
- Guided industry research priorities
- Dissolved in 2015 due to paradigm shift



### Original Roadmap (2001)

#### Optimistic Projections:

The initial roadmap predicted steady exponential decrease in feature size:

- **2001 Baseline:** 130 nm technology in production
- **2006 Target:** 65 nm
- **2008 Target:** 45 nm
- **2012 Projection:** Continuing decrease following Moore's Law

#### Assumptions:

- Linear continuation of historical trends
- Traditional planar transistor scaling
- Continued improvements in lithography
- Economic sustainability of smaller features

### Revised Roadmap (2013)

#### Adjusted Expectations:

By 2013, reality required revised predictions:

- **2013 Actual:** 22 nm achieved
- **2015 Target:** 16 nm predicted
- **2019 Target:** 12 nm
- **2023 Target:** 7 nm
- **2028 Target:** 5 nm

#### Key Observations:

- **Rate of reduction slowed down** compared to original predictions
- Still following exponential trend but slower pace
- Physical and economic challenges becoming apparent
- Need for alternative approaches emerging

### Final Roadmap (2015)

#### Dramatic Shift in Direction:

The 2015 roadmap marked a fundamental change:

- **2015 Status:** Still around 25-24 nm (behind 2013 predictions)
- **Near-term Projection:** Fast improvements predicted to reach 10 nm
- **2021 Target:** 10 nm technology
- **Long-term Direction:** Feature size would **NOT** decrease further beyond 10 nm

- **Plateau:** Would stick with 10 nm for foreseeable future

**Significance:**

- Sudden departure from decades of continuous scaling
- Recognition of fundamental physical limits
- Industry acknowledgment of new paradigm
- End of traditional Moore's Law scaling

**Why the Change? - 3D Technology**

**Major Paradigm Shift (2013-2015):**

The industry pivoted to a fundamentally different approach:

**Traditional Approach (Before):**

- Single layer of transistors on silicon surface
- Scaling by making transistors smaller
- Two-dimensional planar structures

**New Approach (After):**

- **3D Chips:** Multiple layers of transistors stacked vertically
- **3D FinFET Technology:** Transistor fins extending upward from surface
- **Vertical Integration:** Third dimension for density increase

**Impact on Moore's Law:**

- Transistor count **still increasing** (Moore's Law continues)
- But **NOT by making individual transistors smaller**
- Instead: **Stacking transistors on top of each other**
- Adds thickness dimension to chip design

**Technical Innovations:**

- Gate-all-around (GAA) transistors
- Through-silicon vias (TSVs) for vertical connections
- Advanced packaging techniques
- Thermal management solutions

### Dissolution of ITRS (2015)

#### Reasons for Dissolution:

- **Technology Divergence:** Multiple paths to increase transistor density
- **End of Simple Scaling:** No longer just reducing feature size
- **3D Stacking:** Fundamentally different approach
- **Heterogeneous Integration:** Combining different technologies on same chip

#### Multiple Methods for Transistor Density:

Modern approaches include:

- 3D stacking of transistor layers
- FinFET and GAA transistor structures
- Chiplet architectures
- Advanced packaging technologies
- Heterogeneous integration

#### Moore's Law Status:

- Transistor count **still doubling every 2 years** (as of 2021)
- But through **different means** than traditional scaling
- More complex and diverse strategies
- Higher costs per transistor (economic Moore's Law ending)

### 1.2.6 Why Smaller Transistors Improve Performance

#### Reason 1: More Complex Circuits

##### Increased Transistor Budget:

More transistors available on chip enables more sophisticated functionality.

##### Comparison Example:

##### Limited Transistor Count (100 transistors):

- Can only build simple functional units
- Complex tasks must be broken down into simple operations
- Must use simple functional units repeatedly
- Sequential processing of sub-tasks
- **Result: SLOWER overall execution**

##### Abundant Transistors (1 billion):

- Can build extremely complex circuits
- Perform complex operations in single step

- Don't need to decompose into simple operations
- Dedicated hardware for sophisticated functions
- **Result: FASTER overall execution**

**Architectural Implications:**

- Larger caches for better hit rates
- More sophisticated branch predictors
- Wider execution units (SIMD)
- More parallel functional units
- Hardware accelerators for specific tasks

**Reason 2: Faster Switching****Electrical Advantages of Smaller Size:**

Smaller transistors possess superior electrical characteristics:

**Lower Operating Voltage:**

- Smaller channel width requires less voltage to switch transistor
- Voltage scaling: From 5V (1980s) to 1V (modern)
- Reduces power consumption
- Enables higher switching frequencies

**Reduced Impedance:**

- Lower resistance in transistor channel
- Faster current flow
- Quicker charging/discharging of capacitances

**Faster State Changes:**

- Can switch transistor on/off faster
- Less time needed for signal propagation
- Shorter gate delays

**Overall Impact:**

- Faster transistor switching → Higher possible clock rate
- Higher clock rate → More operations per second
- Faster overall computation

**Physical Explanation:**

The relationship between size and speed involves:

- Reduced gate capacitance (smaller area)
- Shorter carrier transit time (shorter channel)
- Lower RC time constants
- Improved frequency response

### 1.2.7 Clock Rate Trends - The Power Wall

#### Clock Rate Increases (1982-2004)

##### Exponential Growth Era:

Processor clock frequencies increased dramatically for over two decades:

##### Historical Progression:




Figure 1.3: Clock Rate Trends and the Power Wall

- **286 (1982):** 12.5 MHz
- **386 (1985):** 16 MHz (author's first computer)
- **486 (Early 1990s):** 25-33 MHz
- **Pentium (Mid-1990s):** 60-200 MHz
- **Pentium 4 (2001):** 2 GHz (2000 MHz) - **First to break 2 GHz barrier**
- **Pentium 4 Prescott (2004):** 3.6 GHz (3600 MHz) - **Peak of single-core era**

##### Growth Rate:

- Nearly 300× increase in 20 years
- Roughly doubled every 18-24 months
- Parallel to Moore's Law for transistor count
- Consumer expectation of continuous frequency increases

#### The Turning Point (2004-2007)

##### Sudden Deceleration:

Around 2004, the decades-long trend dramatically changed:

- Clock rate increase **slowed dramatically**
- Reached peak around **3.6-4 GHz**
- Settled and plateaued at that level
- **Despite transistors continuing to get smaller**

**The Paradox:**

- Manufacturing processes still improving
- More transistors available
- Smaller, potentially faster transistors
- **But clock frequencies stopped increasing**

**Industry Recognition:**

- Fundamental limitation encountered
- Alternative approaches needed
- Architectural innovation required
- End of "free" performance scaling

**The Power Wall Problem**

Power Consumption Growth Crisis:

As clock rates increased, power consumption grew unsustainably:

**Pentium 4 Prescott Example:**

- Required more than **100 watts** of power
- Power supply could provide the necessary electrical power
- **But HEAT became the critical limiting issue**

**The Thermal Crisis:**

Physical reality of heat generation:

**1. Heat Generation Mechanism:**

- Billions of transistors switching billions of times per second
- Each switching event involves current flow
- Current through resistance generates heat ( $I^2R$  losses)
- Accumulated heat from all transistors

**1. Heat Dissipation Challenge:**

- Heat generation outpaced heat removal capability
- Chips would overheat and potentially burn
- Thermal damage to silicon
- Reliability concerns and failure modes

**The 100-Watt Rule of Thumb:**

Industry consensus emerged:

- **Maximum practical limit: 100 watts per chip**

## 1.2. LECTURE 2: TECHNOLOGY TRENDS, MOORE'S LAW, AND COMPUTER SYSTEM ORGANIZATION

- Cooling solutions couldn't effectively handle more
- Would not cross that boundary for desktop processors
- Required alternative approaches to improve performance

### **Attempted Solutions (All Insufficient):**

Various cooling methods were tried:

- **Improved Air Cooling:**
  - Larger heatsinks
  - More powerful fans
  - Better thermal interface materials
- **Liquid Cooling:**
  - Water cooling systems (like car radiators)
  - More efficient heat transfer
  - Complex and expensive
- **Exotic Solutions:**
  - Phase-change cooling
  - Thermoelectric coolers
  - Ultimately impractical for consumer systems

### **None Sufficient:**

- Couldn't overcome fundamental heat generation problem
- Cost and complexity prohibitive
- Reliability concerns
- Not scalable to mass market

## **Dynamic Power Equation**

The Physics of Power Consumption:

Dynamic power consumption follows this relationship:

Power = Capacitance Load  $\times$  Voltage<sup>2</sup>  $\times$  Frequency

### **Factor Analysis (1982-2004):**

#### **Capacitance Load:**

- **Relatively Constant** per transistor
- Inherent to transistor structure and materials
- Determined by semiconductor physics
- Cannot be arbitrarily reduced

**Voltage Reduction:**

- Decreased from **5V to 1V**
- **5× voltage reduction**
- Squared effect: **25× power reduction** contribution
- Significant mitigation strategy

**Frequency Increase:**

- Increased **300× (12 MHz to 3600 MHz)**
- Direct linear effect on power
- **300× power increase** contribution
- Overwhelmed voltage reduction benefits

**Net Effect Calculation:**

$$\begin{aligned}
 \text{Power Scaling} &= (\text{Capacitance}) \times (\text{Voltage}^2) \times (\text{Frequency}) \\
 &= (1\times) \times \left(\frac{1}{5}\right)^2 \times (300\times) \\
 &= (1\times) \times \left(\frac{1}{25}\right) \times (300\times) \\
 &= 12 \times \text{ power increase}
 \end{aligned}$$

**Key Insight:**

- Despite aggressive voltage scaling (25× reduction in  $V^2$  term)
- Frequency increase (300×) overwhelmed the benefit
- Net result: **Massive power increase**
- Power grew faster than could be managed thermally
- Fundamental limitation reached

**Why Voltage Couldn't Scale Further:**

- Transistor threshold voltages have physical limits
- Signal-to-noise ratio requirements
- Reliability constraints
- Leakage current increases at lower voltages



## Overclocking Phenomenon

Marketing and User Community Response:

Emerged prominently around early 2000s during the MHz wars:

### Manufacturer Approach:

- **"Official" Specifications:** Conservative clock speed (e.g., 3.6 GHz)
- **Actual Capability:** Could run at higher speeds without guarantees
- **Marketing Tactic:** Appeal to gamers and power users
- **Risk Disclaimer:** No warranty at higher speeds

### User Overclocking:

Users could manually increase clock speed beyond rated specification:

### Process:

- Change BIOS/UEFI settings
- Increase multiplier or bus speed
- Often increase voltage
- Improve cooling solutions

### Risks:

- **Generate More Heat:** Exceed thermal design power (TDP)
- **Potential Damage:** Could permanently destroy processor
- **Instability:** System crashes and data corruption
- **Reduced Lifespan:** Accelerated aging of components
- **Voided Warranty:** No manufacturer support

### Target Audience:

- **Gamers:** Seeking maximum frame rates
- **Enthusiasts:** Hobbyists and competitors
- **Overclockers:** Specialized community
- **Benchmarkers:** Competitive performance testing

### Industry Impact:

- Created enthusiast market segment
- Influenced product differentiation (K-series Intel chips)
- Added revenue from premium products
- Many processors destroyed but market remained

### 1.2.8 Shift to Multi-Core Processors

#### The Challenge

##### The Industry Dilemma:

By mid-2000s, the semiconductor industry faced a paradox:

##### Available Resources:

- Moore's Law still valid: More transistors available every generation
- Manufacturing processes continuing to improve
- Silicon area increasing or transistor density growing

##### Constraints:

- **Cannot use all transistors simultaneously** (power wall/heat problem)
- **Cannot increase clock rate** (thermal limitations)
- Traditional performance scaling broken

##### Critical Questions:

- How to utilize available transistors?
- How to continue improving computational power?
- How to maintain Moore's Law performance benefits?

#### Solution: Multiple Processor Cores

Paradigm Shift (2004-2008):

Industry pivoted from single-core to multi-core architectures:

##### Core Concept:

Instead of one powerful processor, put **multiple complete processors on same chip**:

- Each core is a complete CPU
- Cores share cache and memory interface
- Can execute different programs simultaneously
- Parallel execution at thread/process level

##### Early Multi-Core Processors:

##### AMD Barcelona (2007):

- **4 cores** on single die
- Shared L3 cache
- Integrated memory controller

##### Intel Core Series:

- Multiple models with **4 cores**
- Hyperthreading technology (2 threads per core)

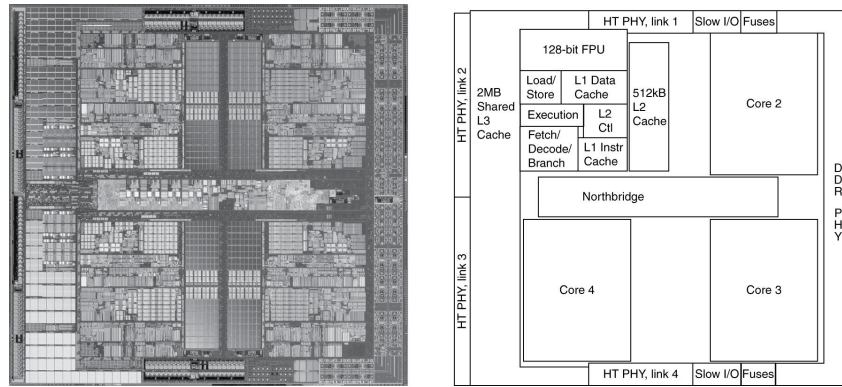


Figure 1.4: AMD Barcelona Quad-Core Processor

- Competitive performance

#### IBM Processors:

- Server-oriented multi-core designs
- High core counts for enterprise
- Power-efficient designs

#### Extreme Designs:

- Some manufacturers: **8 cores** per chip
- Specialized server processors with more
- Graphics processors (GPUs) with hundreds of simple cores

#### Power Management:

- **Dynamic Power Allocation:** Cores powered on/off as needed
- **Turbo Boost:** Temporarily increase frequency of active cores
- **Per-Core Voltage/Frequency Scaling:** Independent control
- **Power Gating:** Completely shut down unused cores
- **Thermal Management:** Distribute heat across die

#### The Plan

##### Initial Industry Vision:

Following Moore's Law principle for core counts:

##### Projected Growth:

- **Every 2 years:** Double the number of cores per chip
- Use increased transistor budget for more cores
- Each generation:  $2\times$  cores, same power envelope

##### Timeline Projection:

- **2006:** 4 cores
- **2008:** 8 cores
- **2010:** 16 cores
- **2012:** 32 cores
- **2014:** 64 cores
- **By 2021:** Should have **hundreds of cores** in consumer processors

**Theoretical Benefits:**

- Continuous performance improvement
- Utilizing Moore's Law transistor growth
- Working within power constraints
- Parallel computing becoming mainstream

**Reality Check:**

- **This did NOT happen**
- Current consumer chips: Typically **4-16 cores** (2021)
- Server processors: Up to 64-128 cores
- Not the hundreds predicted
- Growth much slower than initial projections

**Why Multi-Core Growth Slowed****The Fundamental Problem: Parallel Programming Difficulty****Software Challenge:**

Multi-core processors require fundamentally different programming approach:

**Sequential Programming (Traditional):**

- Single thread of execution
- One operation after another
- Natural mental model
- Straightforward debugging
- Predictable behavior

**Parallel Programming (Required for Multi-Core):**

- Multiple simultaneous threads of execution
- Programmer must **explicitly** write code for multiple processors
- Must think: "I'm writing for 4, 8, or 16 processors"
- Coordinate and synchronize multiple processes/threads

## 1.2. LECTURE 2: TECHNOLOGY TRENDS, MOORE'S LAW, AND COMPUTER SYSTEM ORGANIZATION

- Manage shared resources and data

### **Available Parallel Programming Techniques:**

#### **Multi-Threading:**

- **POSIX Threads (pthreads)** in C/C++
- Java threading primitives
- Python threading/multiprocessing
- Operating system thread scheduling

#### **Multiple Processes:**

- Fork/join models
- Message passing (MPI for scientific computing)
- Process pools

#### **Communication Mechanisms:**

- Shared memory
- Message queues
- Pipes and sockets
- Synchronization primitives (mutexes, semaphores, barriers)

#### **Language Support:**

- Available in most major programming languages
- Library support varies in quality
- Language-level primitives vs library-based approaches

#### **Inherent Difficulties:**

##### **1. Parallel Programming is HARD:**

- Much more difficult than sequential programming
- Different mental model required
- Non-deterministic behavior
- Difficult to reproduce bugs
- Race conditions and deadlocks

##### **2. Requires Deep Understanding:**

- **Hardware Architecture:** How cores communicate, cache coherency
- **Processor Organization:** Memory hierarchy, interconnects
- **Communication Overhead:** Cost of data transfer between cores

- **Synchronization Overhead:** Cost of coordinating execution

#### Key Technical Challenges:

##### Load Balancing:

- **Problem:** Distribute work evenly across all cores
- **Bad Scenario:** One processor idle while another overloaded
- **Requirement:** Dynamic or static work distribution
- **Complexity:** Workload often unknown until runtime
- **Solution Difficulty:** NP-hard problem in general case

##### Communication Optimization:

- **Problem:** Minimize data transfer between cores
- **Reality:** Communication takes time (overhead)
- **Amdahl's Law:** Communication is sequential bottleneck
- **Cache Coherency:** Hardware protocol overhead
- **Solution:** Locality-aware algorithms, minimize sharing

##### Synchronization:

- **Problem:** Coordinate execution between cores
- **Bad Scenario:** One thread waiting indefinitely for another
- **Overhead:** Synchronization primitives have cost
- **Deadlock Risk:** Circular dependencies can halt system
- **Solution:** Careful design, lock-free algorithms

##### Performance Consequences:

If parallel programming not done well:

- **Wasting Available Hardware:** Cores sitting idle
- **No Performance Gain:** Sequential sections dominate
- **Worse Performance:** Overhead exceeds benefits
- **Unpredictable Results:** Race conditions cause incorrect output

## Instruction-Level Parallelism vs Multi-Core Parallelism

### Instruction-Level Parallelism (ILP):

#### Characteristics:

- **Hardware-Based Solution:** Processor automatically finds parallelism
- **Automatic Execution:** Fetches multiple instructions simultaneously
- **Out-of-Order Execution:** Reorders for efficiency
- **Compiler Support:** Helps but not required
- **Transparent to Programmer:** No special code needed
- **Automatic and Hidden:** Works without programmer awareness

#### Techniques:

- Superscalar execution
- Out-of-order execution
- Register renaming
- Speculative execution
- Branch prediction

#### Benefits:

- Free performance improvement
- Works on existing sequential code
- No programmer burden
- Automatic optimization

### Multi-Core Parallelism:

#### Characteristics:

- **Explicit Programming Required:** Programmer must manually parallelize
- **Not Automatic:** No hardware magic
- **Much More Difficult:** Requires expertise
- **Programmer Responsibility:** Must handle all coordination

#### Programmer Must:

- Break program into parallel threads
- Distribute work across cores
- Handle inter-core communication
- Manage synchronization

- Deal with race conditions
- Avoid deadlocks
- Balance load
- Minimize communication overhead

**Contrast:**

Aspect	ILP	Multi-Core		————	————	————		Who does work	
Hardware	Programmer	Transparency	Invisible	Explicit		Difficulty	Automatic	Hard	
Applicability	All code	Limited patterns		Overhead	Hidden	Must manage			

## Impact on Software Development

### For Regular Programmers:

- **Too Difficult:** Most cannot effectively parallelize
- **Not Worth Effort:** For many applications
- **Sequential Sufficient:** Many programs don't need parallel performance
- **Training Gap:** Most programmers not trained in parallel programming

### For Computer Engineers:

- **Essential Skill:** Must learn parallel programming
- **Career Requirement:** High-performance computing demands it
- **Necessary Understanding:** Must understand hardware deeply
- **Specialized Constructs:** Must master threading, synchronization
- **Architecture Knowledge:** Must understand cache coherency, memory models

### Application Domains:

#### High-Performance Applications Requiring Parallelism:

- Scientific computing and simulations
- Video encoding/decoding
- Machine learning training
- Real-time graphics rendering
- Big data processing
- Financial modeling

#### Applications That Remain Sequential:

- Many business applications
- Simple utilities
- I/O-bound programs



## 1.2. LECTURE 2: TECHNOLOGY TRENDS, MOORE'S LAW, AND COMPUTER SYSTEM ORGANIZATION

- Interactive applications
- Legacy software

### **Education Impact:**

- Computer science curricula adding parallel programming courses
- Need for hardware architecture understanding
- Gap between industry needs and graduate preparation
- Specialized training for HPC (high-performance computing)

### 1.2.9 Computer System Organization - Three Layers

#### **Hardware Layer (Bottom)**

##### **Physical Components:**

##### **Processor (CPU):**

- Central Processing Unit
- Executes machine instructions
- Contains control and datapath
- Includes registers and functional units

##### **Microarchitecture:**

- Internal organization of processor
- Pipeline structure
- Execution units
- Cache organization
- Bus interfaces

##### **Memory Hierarchy:**

- **Level 1 Cache (L1):**

- Smallest, fastest
- Separate instruction and data caches
- On-core, immediate access
- Typically 32-64 KB per core

- **Level 2 Cache (L2):**

- Larger, slightly slower
- May be shared or per-core
- Typically 256 KB - 1 MB per core

- **Level 3 Cache (L3):**
  - Largest, slower than L2
  - Shared across all cores
  - Typically several MB
- **Main Memory (RAM):**
  - Dynamic RAM (DRAM)
  - Several GB capacity
  - Much slower than cache
  - Volatile storage

**Input/Output Controllers:**

- USB controllers
- Network interfaces
- Display adapters
- Storage controllers

**Secondary Storage Interfaces:**

- SATA for hard drives/SSDs
- NVMe for fast SSDs
- External storage connections

**Purpose:**

- Actual physical components that execute computation
- Store and retrieve data
- Interact with peripherals and external world
- Provide computational substrate

**System Software Layer (Middle)**

Tool Chain Components:

**Compiler:**

- **Function:** Translates high-level language to assembly
- **Input:** Source code (C, Java, Python, etc.)
- **Output:** Assembly language or intermediate representation
- **Optimization:** Improves performance, reduces size
- **Examples:** GCC, Clang, MSVC, Javac

**Assembler:**

- **Function:** Translates assembly to machine code
- **Input:** Assembly language (human-readable mnemonics)
- **Output:** Object files (binary machine code)
- **Tasks:** Symbol resolution, address assignment
- **Examples:** GNU Assembler (as), NASM

**Linker:**

- **Function:** Combines object files and libraries
- **Tasks:** Resolves external references, creates executable
- **Output:** Complete executable program
- **Link Types:** Static linking, dynamic linking
- **Examples:** GNU ld, MSVC linker

**Purpose of Tool Chain:**

- Support application development
- Bridge high-level abstractions to machine code
- Enable programmer productivity
- Provide optimization opportunities

**Operating System:**

**Core Responsibilities:**

**Resource Management:**

- CPU time allocation
- Memory space allocation
- I/O device arbitration
- Storage space management

**Memory Management:**

- Virtual memory implementation
- Page tables and address translation
- Memory protection between processes
- Swap space management

**Storage Management:**

- File system implementation

- Directory structures
- File permissions and security
- Disk block allocation

**Input/Output Handling:**

- Device drivers
- Interrupt handling
- Buffering and caching
- Asynchronous I/O

**Task Scheduling:**

- Process scheduling algorithms
- Thread scheduling
- Priority management
- Time-slicing and preemption

**Resource Sharing:**

- Prevents conflicts between programs
- Enforces isolation
- Provides controlled sharing mechanisms

**Why Operating System Needed:****Trust and Security:**

- **Cannot trust application software**
- Programs can be malicious or buggy
- Programs don't consider other programs
- Need supervision and enforcement

**Coordination and Protection:**

- Prevents programs from breaking hardware
- Enforces rules set by hardware (privileged instructions)
- Provides abstraction hiding hardware details
- Mediates access to shared resources

**Programmer Benefits:****Abstractions Provided:**

Programmers don't need to worry about:

## 1.2. LECTURE 2: TECHNOLOGY TRENDS, MOORE'S LAW, AND COMPUTER SYSTEM ORGANIZATION

- Where program code resides in physical memory
- Where variables are stored in RAM
- Hardware resource conflicts
- Direct hardware access
- Physical device characteristics

### **OS Guarantees:**

- Safe hardware usage
- Process isolation
- Consistent interfaces
- Reliable file storage
- Network communication

### **Example Services:**

- File I/O without knowing disk geometry
- Memory allocation without physical addresses
- Network communication without protocol details
- Device I/O without hardware specifics

## **Application Software Layer (Top)**

### **User-Level Programs:**

- Programs written by application programmers
- Solve specific problems or provide services
- Interact with users
- Implement business logic

### **High-Level Programming Languages:**

#### **Popular Languages:**

- **C:** Systems programming, performance-critical
- **Java:** Enterprise applications, portability
- **Python:** Scripting, data science, machine learning
- **R:** Statistical analysis, data science
- **JavaScript:** Web development, client-side
- **C++:** Performance with abstraction
- **Go:** Concurrent systems, cloud services

- **Rust:** Systems programming, memory safety

**Language Characteristics:  
Hundreds/Thousands Available:**

- Each optimized for specific application domains
- Different paradigms (imperative, functional, object-oriented)
- Trade-offs between performance and productivity
- Community and ecosystem considerations

**Domain Optimization:**

- **Machine Learning:** Python (NumPy, TensorFlow, PyTorch), R
- **Systems Programming:** C, C++, Rust
- **Enterprise Applications:** Java, C#
- **Web Development:** JavaScript, TypeScript, PHP, Ruby
- **Scientific Computing:** Python, Julia, MATLAB, Fortran
- **Mobile Development:** Swift, Kotlin, Java
- **Game Development:** C++, C#

**Level of Abstraction:**

- Represents algorithms and solutions to problems
- Closest to problem domain
- Furthest from hardware details
- Highest productivity for programmers
- Requires compilation/interpretation to execute

### 1.2.10 From High-Level Code to Machine Code - The Translation Process

#### Example: Swap Function in C

**Source Code:**

```
void swap(int v[], int k) { int temp; temp = v[k]; v[k] = v[k+1]; v[k+1] = temp;
```

**Function Purpose:**

- **Operation:** Swap two values in array
- **Parameters:**
  - `v[]`: Array pointer (base address)
  - `k`: Index of first element to swap
- **Elements Swapped:** Positions `k` and `k+1`
- **Method:** Uses temporary variable

- **Simplicity:** Basic operation used frequently in sorting algorithms

**Algorithm:**

1. Store  $v[k]$  in temporary variable
2. Copy  $v[k+1]$  to  $v[k]$
3. Copy temporary to  $v[k+1]$

**After Compilation - MIPS Assembly Code**

**Assembly Translation:**

The compiler generates 7 MIPS instructions to implement the swap function:

```
MUL  $2, $5, 4      # Multiply k by 4 (array index to byte offset)
ADD  $2, $4, $2      # Add base address to offset (address of v[k])
LW   $15, 0($2)     # Load v[k] into register $15 (temp = v[k])
LW   $16, 4($2)     # Load v[k+1] into register $16
SW   $16, 0($2)     # Store v[k+1] to v[k]
SW   $15, 4($2)     # Store temp to v[k+1]
```

**Translation Analysis:**

- **5 C statements**  $\rightarrow$  **7 assembly instructions**
- Expansion due to instruction granularity
- Each assembly instruction is simple operation

**Key Operations Explained:**

**1. Address Calculation:**

- **Multiply by 4:** Each integer occupies 4 bytes in memory
- **Index k** must be converted to **byte offset** ( $k \times 4$ )
- Calculate memory address of  $v[k]$

**2. Memory Addressing:**

- Base address of array in register \$4
- Offset calculated and added to base
- Results in absolute memory address

**3. Register Usage:**

- **\$4:** Base address of array  $v$  (parameter)
- **\$5:** Value of  $k$  (parameter)
- **\$2:** Temporary register for address calculation
- **\$15:** Temporary storage for  $v[k]$  value
- **\$16:** Temporary storage for  $v[k+1]$  value

**Instruction Set Details:**

- **MIPS ISA** used in example (not ARM, but similar concepts)
- Load-Store architecture
- Register-to-register operations
- Explicit memory addressing

**After Assembly - Machine Code****Binary Representation:**

Each assembly instruction translates to 32-bit binary instruction:

```
00000000101000100001000000011000 # MUL $2, $5, 4
00000000100000100001000000100001 # ADD $2, $4, $2
10001100010011110000000000000000 # LW $15, 0($2)
10001100010100000000000000000100 # LW $16, 4($2)
10101100010100000000000000000000 # SW $16, 0($2)
10101100010011110000000000000100 # SW $15, 4($2)
```

**One-to-One Mapping:**

- Each assembly instruction → Exactly one 32-bit machine instruction
- No information lost or gained
- Deterministic translation
- Assembly is human-readable form of machine code

**Instruction Format:**

Different instruction types have different bit field layouts:

**R-Type (Register) Format:**

[Opcode 6 bits][Rs 5 bits][Rt 5 bits][Rd 5 bits][Shamt 5 bits][Funct 6 bits]

**I-Type (Immediate) Format:**

[Opcode 6 bits][Rs 5 bits][Rt 5 bits][Immediate 16 bits]

**Instruction Components Specify:**

- **Opcode:** Operation category
- **Destination Register:** Where result goes
- **Source Registers:** Where operands come from
- **Immediate Values:** Constant values (like 4 in multiply)
- **Function Code:** Specific operation for R-type

**Example Analysis:**

In the immediate value 4:

- Appears in specific bit positions
- Encoded in binary (00000000000100)



## 1.2. LECTURE 2: TECHNOLOGY TRENDS, MOORE'S LAW, AND COMPUTER SYSTEM ORGANIZATION

- Part of instruction encoding

### **Binary Image:**

- Complete program represented as sequence of 32-bit words
- Called **executable** or **binary image**
- Stored in secondary storage (hard disk, SSD)
- Loaded into memory when program executes
- CPU fetches and executes instructions sequentially

### 1.2.11 Program Execution - Inside the CPU

#### **Block Diagram of Computer**

##### **System Components:**

##### **Compiler/Tool Chain:**

- Translates human-written program to machine code
- Optimization and code generation
- Produces executable binary

##### **Memory:**

- Stores program instructions
- Stores program data
- Hierarchical (cache, RAM, disk)

##### **CPU (Central Processing Unit):**

- Executes machine instructions
- Performs arithmetic and logic
- Controls program flow

##### **Input/Output:**

- Peripherals (keyboard, display, network)
- Storage devices (disk, SSD)
- Communication interfaces

##### **Program Execution Flow:**

##### **1. Compile Stage:**

- Source code  $\rightarrow$  Assembly  $\rightarrow$  Machine code
- Performed once (or when code changes)
- Output: Executable binary file

**2. Store Stage:**

- Machine code saved to secondary storage
- Persistent storage (survives power off)
- Typically on hard disk or SSD

**3. Load Stage:**

- Machine code loaded into main memory (RAM) when program runs
- Operating system performs loading
- Program becomes "process"

**4. Execute Stage:**

- CPU fetches instructions from memory one by one
- Executes each instruction in sequence (or out-of-order)
- Updates registers and memory

**5. Results Stage:**

- Computed values stored back in memory
- Output sent to I/O devices
- Results displayed or saved

**Inside the CPU - Two Main Components****Datapath:****Structure:**

- Collection of logic circuits interconnected
- Forms a path through CPU
- Instruction and data travel through this path
- Sequential stages of processing

**Components:**

- Functional units (adders, multipliers, shifters, logic units)
- Registers for temporary storage
- Multiplexers for routing
- Buses for data transfer

**Function:**

- Instruction travels from one logic circuit to another

## 1.2. LECTURE 2: TECHNOLOGY TRENDS, MOORE'S LAW, AND COMPUTER SYSTEM ORGANIZATION

- Each circuit performs specific operation on data
- Transforms inputs to outputs
- Executes the computational work

### **Examples of Functional Units:**

- Arithmetic Logic Unit (ALU)
- Floating-Point Unit (FPU)
- Load-Store Unit
- Branch Unit

### **Control:**

#### **Structure:**

- Another logic circuit (or set of circuits)
- Generates control signals
- Coordinates datapath operation

### **Function:**

- Governs instruction/data flow through datapath
- Ensures instructions execute correctly
- Selects appropriate functional units
- Controls multiplexers and enables

### **Responsibilities:**

- Decode instructions
- Generate appropriate control signals
- Coordinate timing
- Handle exceptions and interrupts

### **Interaction:**

- **Control** tells **Datapath** what to do
- **Datapath** performs the actual computation
- **Control** monitors **Datapath** status
- Together implement instruction execution

**Execution Process (Conveyor Belt Analogy)****Instruction Execution Cycle:****1. Fetch:**

- Instructions stored in memory
- Control fetches one instruction at a time
- Brings instruction into CPU
- Increments program counter

**2. Decode:**

- Instruction enters datapath
- Control decodes instruction
- Determines operation type
- Identifies operands

**3. Execute:**

- Instruction travels through logic circuits in datapath
- Operations performed on data
- Functional units activated
- Intermediate results produced

**4. Memory:**

- Memory accesses performed if needed (load/store)
- Data read from or written to memory
- Address calculation completed

**5. Writeback:**

- Results generated
- Written back to registers
- Results may be sent to memory or I/O

**6. Repeat:**

- Cycle repeats for next instruction
- Like conveyor belt: continuous flow
- One instruction after another (in simple model)

**Conveyor Belt Metaphor:**

- Instructions like items on conveyor belt
- Each station performs specific operation
- Continuous movement through system
- Pipelining overlaps multiple instructions (discussed in later lectures)

## Cache Memory

### Purpose and Motivation:

#### The Performance Gap:

- CPU can process data very fast
- Main memory access is relatively slow
- Speed mismatch creates bottleneck
- CPU would waste time waiting for memory

#### Cache Solution:

- Fast memory located on CPU chip
- Very close to processor core physically
- Stores copies of frequently used instructions and data
- Exploits locality of reference

### Cache Hierarchy:

#### Level 1 Cache (L1):

- Smallest capacity (32-64 KB)
- Fastest access (1-2 cycles)
- Closest to core
- Often split: L1-I (instruction), L1-D (data)

#### Level 2 Cache (L2):

- Medium capacity (256 KB - 1 MB)
- Medium access time (4-10 cycles)
- May be per-core or shared
- Unified (instructions and data)

#### Level 3 Cache (L3):

- Largest capacity (several MB)
- Slower access (20-40 cycles)
- Shared across all cores
- Last level cache (LLC)

### Performance Impact:

- Cache hit: Data found in cache (fast)
- Cache miss: Must access main memory (slow)

- Hit rate critical for performance
- Well-designed cache can achieve >95% hit rate

**Will Learn in Lecture:**

- Cache organization
- Mapping strategies (direct-mapped, set-associative)
- Replacement policies
- Write policies
- Cache coherency in multi-core

### 1.2.12 Real CPU Layout - AMD Barcelona Example

#### Overview

**AMD Barcelona Processor:**

- Released around 2007
- Quad-core processor (4 cores on single die)
- 65nm manufacturing process
- Actual chip much smaller than magnified images
- Can visually identify individual components

**Die Photo Analysis:**

- Optical or electron microscope image
- Shows physical layout of components
- Different functional units visible
- Reveals organizational decisions
- Educational value for understanding architecture

#### Four Processor Cores

**Core Distribution:**

Physical layout shows clear quadrant organization:

- **Core 1:** Upper left area of die
- **Core 2:** Upper right area of die
- **Core 3:** Lower left area of die
- **Core 4:** Lower right area of die

**Layout Strategy:**

- **Mirror Image Layouts:** Cores identical but mirrored
- **Symmetry:** Simplifies design and manufacturing
- **Thermal Distribution:** Spreads heat across die
- **Interconnect Balance:** Equal distances to shared resources

### Inside Each Core

#### Floating-Point Unit (FPU):

##### Characteristics:

- **Large Component:** Significant silicon area in each core
- **Complex Circuitry:** Handles IEEE 754 floating-point arithmetic
- **High Transistor Count:** Precision requires many gates

##### Operations:

- Addition, subtraction of floating-point numbers
- Multiplication of floating-point numbers
- Division of floating-point numbers
- Square root and other mathematical functions

##### Why So Large:

- Floating-point math more complex than integer
- Requires normalization, rounding, exception handling
- Multiple pipeline stages
- High precision demands

#### Load-Store Unit:

##### Function:

- Handles all memory operations
- Loads data from memory to CPU registers
- Stores data from CPU registers to memory
- Critical for data transfer

##### Operations:

- Address calculation
- Cache access
- TLB (Translation Lookaside Buffer) lookup
- Memory ordering and consistency

#### Integer Execution Unit:

##### Characteristics:

- **Smaller than FPU:** Integer operations generally simpler
- **High Frequency:** Often faster than floating-point

**Operations:**

- Integer arithmetic (add, subtract, multiply, divide)
- Bitwise logical operations (AND, OR, XOR, NOT)
- Shifts and rotates
- Comparisons

**Why Smaller:**

- Simpler algorithms
- No normalization needed
- Exact arithmetic (no rounding)
- Fewer pipeline stages

**Fetch and Decode Unit:****Responsibilities:****Instruction Fetch:**

- Fetches instructions from memory (via I-cache)
- Predicts branch targets
- Manages instruction buffer

**Instruction Decode:**

- Makes sense of binary instruction encoding
- Determines instruction type
- Identifies operands
- Generates micro-ops (for CISC architectures)

**Pipeline Frontend:**

- Prepares instructions for execution
- Handles instruction-level parallelism
- Feeds execution units

**Level 1 Data Cache (L1 D-Cache):****Characteristics:**

- Stores frequently used **data** (not instructions)
- Very fast access (1-2 cycle latency)
- Close to execution units
- Separate from instruction cache (Harvard architecture)



**Typical Specifications:**

- 32-64 KB capacity
- 8-way set associative
- Write-through or write-back policy

**Level 1 Instruction Cache (L1 I-Cache):**

**Characteristics:**

- Stores frequently used **instructions** (program code only)
- Very fast access
- Feeds fetch unit
- Separate from data cache

**Benefits of Separation:**

- No structural hazards (simultaneous instruction fetch and data access)
- Optimized for different access patterns
- Simpler control logic

**Level 2 Unified Cache (L2 Cache):**

**Characteristics:**

- **Larger than L1:** Typically 512 KB per core in Barcelona
- Stores **both instructions and data** (unified)
- Further from execution units (higher latency)
- Victim cache for L1 misses

**Architecture:**

- Dedicated control logic for coherency
- Interface to L3 cache or memory
- May use different associativity than L1

**Shared Components**

**North Bridge (Central Hub):**

**Location:**

- Central/middle area of chip
- Strategic position for communication

**Functions:**

- **L2-to-Memory Connection:** Connects all L2 caches to main memory

- **Inter-Core Communication:** Coordinates between cores
- **Memory Controller:** May include integrated memory controller
- **Cache Coherency:** Maintains coherency protocol between cores

**Critical Role:**

- Central communication circuit
- Bandwidth bottleneck if not designed well
- Affects multi-core scaling

**DDR PHY (Physical Controller):**

**DDR Memory:**

- **DDR:** Dual Data Rate SDRAM
- Transfers data on both rising and falling clock edges
- Industry-standard memory interface

**PHY (Physical Layer):**

- **PHY:** Physical layer controller
- Interfaces CPU to DDR RAM modules
- Handles physical signaling

**Responsibilities:**

- Electrical interface to memory chips
- Signal timing and termination
- Training and calibration
- Error detection/correction

**HyperTransport Controllers:**

**HyperTransport Technology:**

- High-speed interconnect technology (AMD proprietary)
- Point-to-point serial communication
- Replaces legacy parallel buses
- High bandwidth, low latency

**Connections:**

- **External Devices:** Graphics cards, other processors
- **Chipset Communication:** Northbridge, southbridge links
- **I/O Device Connectivity:** Network, storage, peripherals

**Benefits:**

- Scalable bandwidth
- Lower pin count than parallel buses
- NUMA (Non-Uniform Memory Access) support for multi-socket systems

### **Additional Information**

**WikiChip Database:** <https://en.wikichip.org>

**Comprehensive Processor Information:**

**Major Manufacturers Covered:**

- **Intel Processors:** x86 architecture, Core series, Xeon servers
- **AMD Processors:** x86 architecture, Ryzen, EPYC, Threadripper
- **ARM Processors:** Mobile devices, embedded systems, servers
- **Samsung Exynos:** Smartphones and tablets
- **Apple A-Series:** iPhone and iPad processors
- **Apple M-Series:** Mac computers (ARM-based)
- **Qualcomm:** Snapdragon mobile processors
- **NVIDIA, Broadcom, Texas Instruments, and many more**

**Available Information:**

**Visual Content:**

- Processor die photographs and diagrams
- Block diagrams showing architecture
- Cache hierarchy visualizations
- Microarchitecture pipeline diagrams

**Technical Specifications:**

- Manufacturing process (nm technology)
- Transistor counts and density
- Transistor types and structures
- Die size and area
- Power consumption (TDP)
- Clock speeds (base and turbo)
- Core counts and threading
- Cache sizes and organization

**Advanced Topics:**

- 3D stacking technology details
- FinFET and GAA transistor structures
- Packaging technologies
- Memory interface specifications

- I/O capabilities

**Current Technology Landscape (2021):  
Mainstream Manufacturing:**

- **10 nm and 7 nm** processes in volume production
- Multiple manufacturers at this node

**Future Direction:**

- **Next Few Years:** Shift to 5 nm and 3 nm
- 2 nm and 1 nm in research

**Important Clarification:**

- **Numbers don't represent actual gate size anymore**
- **Marketing terms** more than physical measurements
- **Example:** 5 nm transistors may have wider channels than 10 nm
- **Density Increase Through:**
  - 3D stacking (vertical integration)
  - FinFET and GAA structures
  - Improved layouts and design rules
  - Multi-patterning lithography

### 1.2.13 Key Takeaways

1. **Moore's Law predicted transistor doubling every 2 years** - remarkably accurate for over 40 years, guiding semiconductor industry planning and investment
1. **Smaller transistors enabled by improved lithography** - progression from 90nm → 45nm → 22nm → 7nm → 5nm through advancing manufacturing processes
1. **Feature size now marketing term rather than physical measurement** - modern processes use 3D structures making simple linear dimensions misleading
1. **Smaller transistors provide dual benefits** - enable more complex circuits (more transistors available) and faster switching (lower voltage, reduced impedance)
1. **Clock rate increased exponentially until 2004** - grew from 12.5 MHz (1982) to 3.6 GHz (2004), then hit fundamental thermal limitations
1. **Power wall halted frequency scaling** - heat generation ( $P = CV^2f$ ) exceeded cooling capability, establishing 100W practical limit for consumer processors
1. **Dynamic power equation explains the crisis** - despite  $25\times$  power reduction from voltage scaling,  $300\times$  frequency increase overwhelmed the benefit
1. **Overclocking emerged as risky performance technique** - users could exceed rated speeds at risk of destroying processors, popular among gaming enthusiasts

1. **Industry pivoted to multi-core processors** - solution to utilize Moore's Law transistors without exceeding power limits, starting 2004-2008
1. **Multi-core growth slowed due to programming difficulty** - initial projection of hundreds of cores didn't materialize; parallel programming remains challenging
1. **Parallel programming requires explicit management** - unlike automatic instruction-level parallelism, multi-core requires programmers to handle threads, synchronization, communication
1. **Three major parallel programming challenges** - load balancing across cores, minimizing communication overhead, optimizing synchronization
1. **3D chip technology changed scaling paradigm (2013-2015)** - industry shifted from pure 2D shrinking to vertical stacking of transistor layers
1. **ITRS dissolved in 2015** - technology roadmap organization ended as multiple paths to density replaced simple feature size scaling
1. **Computer systems organized in three layers** - hardware (physical components), system software (OS, compilers, tools), application software (user programs)
1. **System software provides abstraction and protection** - OS prevents malicious programs from damaging hardware, hides complexity from application programmers
1. **Program translation is multi-stage process** - high-level language → assembly language → machine code through compiler, assembler, linker
1. **CPU contains datapath and control** - datapath performs computation by routing data through functional units; control coordinates execution and generates signals
1. **Cache memory critical for performance** - fast on-chip memory (L1, L2, L3) stores frequently accessed data/instructions, hiding main memory latency
1. **Real CPUs have complex layouts** - die photos reveal intricate organization with multiple cores, cache hierarchies, shared interconnects, memory controllers

### 1.2.14 Summary

This lecture provides a comprehensive examination of computer technology evolution from the 1970s to present day. Moore's Law, predicting transistor count doubling every two years, serves as the guiding principle for the semiconductor industry and enables the transformation of computers from room-sized machines to powerful pocket devices.

The progression of manufacturing technology steadily reduced feature sizes from 90 nanometers to current 7nm and 5nm processes. Smaller transistors provided two key advantages: more transistors per chip enabling complex functionality, and faster switching speeds enabling higher clock frequencies. Clock rates grew exponentially from 12.5 MHz in 1982 to 3.6 GHz in 2004.

However, around 2004, the industry encountered the power wall - a fundamental thermal limitation. The dynamic power equation ( $P = CV^2f$ ) revealed that despite aggressive voltage scaling, the massive frequency increases caused power consumption and heat generation to exceed cooling capabilities. The 100-watt limit for consumer processors could not be overcome by improved cooling solutions.

The solution was multi-core processors: placing multiple complete CPU cores on a single chip. This allowed continued performance improvement within power constraints by exploiting

thread-level parallelism. However, the initial vision of exponentially growing core counts didn't materialize due to the difficulty of parallel programming. Unlike automatic instruction-level parallelism, multi-core requires programmers to explicitly manage threads, balance loads, minimize communication, and handle synchronization - a significantly more challenging paradigm.

Around 2013-2015, the industry made another major shift to 3D chip technology. Instead of only shrinking transistors in two dimensions, manufacturers began stacking transistor layers vertically using FinFET and similar technologies. This represented such a fundamental change that the International Technology Roadmap for Semiconductors (ITRS) dissolved in 2015, as simple feature-size predictions no longer captured the diverse approaches to increasing transistor density.

The lecture concluded by examining computer system organization across three layers: hardware (processor, memory, I/O), system software (compilers, assemblers, operating system), and application software (programs written in high-level languages). We traced the complete journey from high-level code through compilation and assembly to binary machine code, and explored how programs execute through the interaction of control and datapath components within the CPU. Cache memory's critical role in hiding main memory latency was emphasized, and real-world processor layouts illustrated the complex organization of modern multi-core chips.

Understanding these technology trends and architectural responses provides essential context for studying computer architecture and explains why processors are organized as they are today.

## 1.3 Lecture 3: Understanding Performance

*By Dr. Isuru Nawinne*

### 1.3.1 Introduction

Understanding computer performance is fundamental to computer architecture and system design. This lecture explores how performance is measured, the factors that influence it, and the principles that guide performance optimization. We examine the metrics used to evaluate systems, the mathematical relationships between performance factors, and Amdahl's Law—a critical principle for understanding the limits of performance improvements.

### 1.3.2 Defining and Measuring Performance

#### Response Time vs. Throughput

##### Response Time (Execution Time)

- Time to complete a single task
- Includes all overhead and waiting time
- User-perceived performance metric
- Example: Time for a program to run from start to finish

##### Throughput (Bandwidth)

- Number of tasks completed per unit time
- Measures system capacity
- Important for servers and data centers
- Example: Number of transactions processed per second

**Relationship Between Metrics**

- Improving response time often improves throughput
- Improving throughput doesn't always improve response time
- Different optimization strategies for each metric
- System design must balance both considerations

**Performance Definition****Mathematical Definition**

Performance =  $1 / \text{Execution Time}$

**Performance Comparison**

- If System A is faster than System B:
- $\text{Execution Time}_A < \text{Execution Time}_B$
- $\text{Performance}_A > \text{Performance}_B$

**Relative Performance**

$\text{Performance}_A / \text{Performance}_B = \text{Execution Time}_B / \text{Execution Time}_A$

Example: If System A is 2× faster than System B:

- $\text{Performance}_A / \text{Performance}_B = 2$
- $\text{Execution Time}_B / \text{Execution Time}_A = 2$
- System A takes half the time of System B

**1.3.3 CPU Time and Performance Factors****Components of Execution Time****Total Execution Time**

- CPU time: Time CPU spends computing the task
- I/O time: Time waiting for input/output operations
- Other system activities: OS overhead, other programs

**CPU Time Focus**

- Primary metric for processor performance
- Excludes I/O and system effects
- Directly reflects processor and memory system performance
- Most relevant for comparing processor architectures

## The CPU Time Equation

### Basic Formula

CPU Time = Clock Cycles  $\times$  Clock Period

Or equivalently:

CPU Time = Clock Cycles / Clock Rate

### Key Relationships

- Clock Period =  $1 / \text{Clock Rate}$
- Clock Rate measured in Hz (cycles/second)
- Clock Cycles = total cycles to execute program
- Higher clock rate  $\rightarrow$  shorter clock period  $\rightarrow$  faster execution

### Example Calculation

Program requires 10 billion cycles Processor runs at 4 GHz ( $4 \times 10^9$  Hz)

CPU Time =  $10 \times 10^9$  cycles / ( $4 \times 10^9$  cycles/sec) = 2.5 seconds

## Instruction Count and CPI

### Cycles Per Instruction (CPI)

- Average number of clock cycles per instruction
- Varies by instruction type and implementation
- Key microarchitecture metric

### Extended CPU Time Equation

CPU Time = Instruction Count  $\times$  CPI  $\times$  Clock Period

Or:

CPU Time = (Instruction Count  $\times$  CPI) / Clock Rate

### Three Performance Factors

1. **Instruction Count:** Number of instructions executed
2. **CPI:** Average cycles per instruction
3. **Clock Rate:** Speed of the processor clock

### Factor Dependencies

- Instruction Count: Determined by algorithm, compiler, ISA
- CPI: Determined by processor implementation (microarchitecture)
- Clock Rate: Determined by hardware technology and organization



### 1.3.4 Understanding CPI in Detail

#### CPI Variability

##### Different Instructions, Different CPIs

- Simple operations: May complete in 1 cycle (ADD, AND)
- Memory operations: May take multiple cycles (LOAD, STORE)
- Branch instructions: Variable cycles (depends on prediction)
- Multiply/Divide: Often take many cycles

##### Calculating Average CPI

Average CPI =  $\Sigma (\text{CPI}_i \times \text{Instruction Count}_i) / \text{Total Instruction Count}$

Where:

- $\text{CPI}_i$  = cycles per instruction for instruction type i
- $\text{Instruction Count}_i$  = number of times instruction i executed

##### CPI Example Calculation

###### Given:

- Program executes 100,000 instructions
- 50,000 ALU operations (CPI = 1)
- 30,000 load instructions (CPI = 3)
- 20,000 branch instructions (CPI = 2)

###### Calculation:

Total Cycles =  $(50,000 \times 1) + (30,000 \times 3) + (20,000 \times 2) = 50,000 + 90,000 + 40,000 = 180,000$  cycles

Average CPI =  $180,000 / 100,000 = 1.8$

##### Instruction Classes

##### Common Instruction Categories

1. **Integer arithmetic:** ADD, SUB, AND, OR
2. **Data transfer:** LOAD, STORE
3. **Control flow:** BRANCH, JUMP, CALL
4. **Floating-point:** FADD, FMUL, FDIV

##### CPI Characteristics by Class

- Integer arithmetic: Usually 1 cycle
- Data transfer: 1-3 cycles (cache hit) or more (cache miss)
- Control flow: 1-2 cycles (correct prediction) or more (misprediction)
- Floating-point: 2-20+ cycles depending on operation

### 1.3.5 Performance Optimization Principles

#### Make the Common Case Fast

##### Core Principle

- Optimize frequent operations rather than rare ones
- Greater impact on overall performance
- Focus resources where they matter most

##### Examples

- Optimize ALU operations (common) over division (rare)
- Fast cache for recent data (commonly accessed)
- Branch prediction for likely paths
- Simple instructions execute quickly

##### Application in Design

- Identify common operations through profiling
- Allocate hardware resources accordingly
- Accept slower performance for rare cases
- Trade-offs guided by usage patterns

#### Amdahl's Law

**The Fundamental Principle** The speedup that can be achieved by improving a particular part of a system is limited by the fraction of time that part is used.

##### Mathematical Formula

$$\text{Speedup\_overall} = 1 / [(1 - P) + (P / S)]$$

Where:

- $P$  = Proportion of execution time that can be improved
- $S$  = Speedup of the improved portion
- $(1 - P)$  = Proportion that cannot be improved

##### Alternative Formulation

$$\text{Execution Time\_new} = \text{Execution Time\_old} \times [(1 - P) + (P / S)]$$

#### Amdahl's Law Examples

##### Example 1: Multiply Operation Speedup

Given:

- Multiply operations take 80% of execution time
- New hardware makes multiplies  $10\times$  faster

Calculation:

$P = 0.80$  (80% can be improved)  $S = 10$  (10× speedup)

$\text{Speedup\_overall} = 1 / [(1 - 0.80) + (0.80 / 10)] = 1 / [0.20 + 0.08] = 1 / 0.28 = 3.57\times$

**Key Insight:** Despite 10× improvement in multiplies, overall speedup is only 3.57× because 20% of time is unaffected.

#### Example 2: Limited Improvement Fraction

Given:

- Only 30% of execution can be improved
- Improvement is 100× faster

Calculation:

$P = 0.30$   $S = 100$

$\text{Speedup\_overall} = 1 / [(1 - 0.30) + (0.30 / 100)] = 1 / [0.70 + 0.003] = 1 / 0.703 = 1.42\times$

**Key Insight:** Even with 100× improvement, overall speedup is only 1.42× because only 30% of execution benefits.

### Implications of Amdahl's Law

#### Limitations of Parallelization

- Serial portions limit parallel speedup
- As parallelism increases, serial portion dominates
- Cannot achieve infinite speedup regardless of cores

#### Optimization Strategy

- Focus on largest contributors to execution time
- Consider what fraction can realistically be improved
- Multiple small improvements may beat one large improvement
- Balance improvements across components

#### Example: Multicore Scaling

If 90% of program parallelizes perfectly: 2 cores: Speedup = 1.82× 4 cores: Speedup = 3.08× 8 cores: Speedup = 4.71× 16 cores: Speedup = 6.40× ∞ cores: Speedup = 10.00× (maximum possible)

The 10% serial portion ultimately limits speedup to 10×.

### 1.3.6 Complete Performance Analysis

#### The Complete Performance Equation

##### Bringing It All Together

$\text{CPU Time} = (\text{Instruction Count} \times \text{CPI} \times \text{Clock Period})$

Expanded:

$\text{CPU Time} = (\text{Instructions}) \times (\text{Cycles/Instruction}) \times (\text{Seconds/Cycle})$

#### What Affects Each Factor

##### Instruction Count:

- Algorithm: Efficient algorithms execute fewer instructions

- Programming language: High-level vs low-level
- Compiler: Optimization quality
- ISA: Instruction complexity and capabilities

**CPI:**

- ISA: Instruction complexity
- Microarchitecture: Pipeline depth, branch prediction
- Cache performance: Hit rates affect memory access CPI
- Instruction mix: Distribution of instruction types

**Clock Period (or Clock Rate):**

- Technology: Transistor speed (nm process)
- Organization: Pipeline depth, critical path length
- Power constraints: Higher frequency requires more power
- Cooling limitations: Heat dissipation capacity

**Performance Comparison Example**

**Scenario:** Compare two implementations of the same ISA

- System A: Clock Rate = 2 GHz, CPI = 2.0
- System B: Clock Rate = 3 GHz, CPI = 3.0
- Same program with 1 million instructions

**System A:**

$\text{CPU Time}_A = (1 \times 10^6 \text{ instructions}) \times (2.0 \text{ cycles/instruction}) / (2 \times 10^9 \text{ cycles/sec}) = 2 \times 10^6 \text{ cycles} / (2 \times 10^9 \text{ cycles/sec}) = 0.001 \text{ seconds} = 1 \text{ millisecond}$

**System B:**

$\text{CPU Time}_B = (1 \times 10^6 \text{ instructions}) \times (3.0 \text{ cycles/instruction}) / (3 \times 10^9 \text{ cycles/sec}) = 3 \times 10^6 \text{ cycles} / (3 \times 10^9 \text{ cycles/sec}) = 0.001 \text{ seconds} = 1 \text{ millisecond}$

**Result:** Both systems have identical performance despite different clock rates and CPIs.

**Trade-offs in Design****Clock Rate vs. CPI Trade-off**

- Higher clock rate may require deeper pipeline
- Deeper pipeline often increases CPI (more stalls)
- Must balance frequency gains against CPI losses

**Instruction Count vs. CPI Trade-off**

- Complex instructions reduce instruction count

- But complex instructions may increase CPI
- CISC vs RISC architecture debate

#### **Power vs. Performance**

- Higher clock rate increases power consumption
- $\text{Power} = \text{Capacitance} \times \text{Voltage}^2 \times \text{Frequency}$
- Mobile systems prioritize power over peak performance

### **1.3.7 Practical Performance Considerations**

#### **Benchmarking**

##### **Purpose of Benchmarks**

- Measure real-world performance
- Compare different systems objectively
- Standard workloads for reproducibility

##### **Types of Benchmarks**

- Synthetic: Artificial programs (e.g., Dhrystone, Whetstone)
- Application: Real programs (e.g., SPEC CPU, databases)
- Workload: Representative task mixes

##### **Benchmark Pitfalls**

- May not represent your workload
- Can be optimized for unfairly
- Need multiple benchmarks for complete picture

#### **Performance Metrics in Practice**

##### **MIPS (Million Instructions Per Second)**

$$\text{MIPS} = \text{Instruction Count} / (\text{Execution Time} \times 10^6) = \text{Clock Rate} / (\text{CPI} \times 10^6)$$

##### **Limitations of MIPS:**

- Doesn't account for instruction complexity
- Different ISAs have different instruction capabilities
- Higher MIPS doesn't guarantee better performance
- "Meaningless Indication of Processor Speed"

##### **Better Metrics:**

- Execution time for specific workloads
- Throughput for server applications
- Energy efficiency (performance per watt)
- Performance per dollar

## Power and Energy Considerations

### Power Wall

- Cannot increase clock rate indefinitely
- Power consumption limits frequency scaling
- Led to multi-core era

#### Dynamic Power Equation

Power = Capacitance  $\times$  Voltage<sup>2</sup>  $\times$  Frequency

#### Energy Equation

Energy = Power  $\times$  Time

#### Implications:

- Lowering voltage reduces power dramatically (squared effect)
- Higher frequency increases power linearly
- Faster execution may save energy overall (less time)
- Energy efficiency increasingly important metric

### 1.3.8 Key Takeaways

1. **Performance is the inverse of execution time** - faster systems have shorter execution times and higher performance values.
1. **Three key factors determine CPU performance:**
  - Instruction Count (algorithm, compiler, ISA)
  - CPI (microarchitecture, instruction mix)
  - Clock Rate (technology, organization)
1. **Amdahl's Law limits speedup** - the potential speedup from improving any part of a system is limited by how much time that part is used.
1. **"Make the common case fast"** - optimize frequently executed operations for maximum impact on overall performance.
1. **CPI varies by instruction type** - average CPI depends on the mix of instructions and their individual costs.
1. **Trade-offs are fundamental** - improvements in one area (e.g., clock rate) may harm another (e.g., CPI or power consumption).
1. **Benchmarking is essential** - real workloads provide the most meaningful performance measurements.
1. **Power is a critical constraint** - modern performance optimization must consider power and energy efficiency, not just speed.
1. **Multiple factors must be optimized together** - focusing on only one aspect (like clock rate) can be counterproductive.
1. **Understanding performance equations** enables rational design decisions and accurate performance predictions.

### 1.3.9 Summary

Performance analysis is central to computer architecture, providing the foundation for making informed design decisions. By understanding the relationship between instruction count, CPI, and clock rate, architects can identify optimization opportunities and predict the impact of changes. Amdahl's Law reminds us that the benefit of any improvement is constrained by what fraction of execution time it affects, emphasizing the importance of focusing on the common case. As we design systems, we must balance competing factors—clock rate, CPI, power consumption, and cost—to achieve the best overall performance for target applications. The principles covered in this lecture provide the analytical framework for evaluating processor designs and optimization strategies throughout the study of computer architecture.





## Chapter 2

# ARM Assembly Programming

### 2.1 Lecture 4: Introduction to ARM Assembly

*By Dr. Kisaru Liyanage*

#### 2.1.1 Introduction

This lecture introduces ARM assembly language programming, providing the foundation for understanding how high-level programs translate to machine code. We explore the ARM instruction set architecture (ISA), focusing on its RISC design philosophy, register organization, basic instruction formats, and the toolchain used for development. Understanding assembly language is essential for comprehending how processors execute programs and for optimizing performance-critical code.

#### 2.1.2 ARM Architecture Overview

##### RISC Philosophy

##### Reduced Instruction Set Computer (RISC)

- Simple, uniform instruction format
- Fixed instruction length (32 bits in ARM)
- Load/store architecture (only LOAD/STORE access memory)
- Large number of general-purpose registers
- Few addressing modes
- Hardware simplicity for higher clock rates

##### Contrasted with CISC (Complex Instruction Set Computer)

Feature	RISC		CISC		——— ——— ———		<b>Instruction Format</b>		Simple, uniform format	
	Variable-length instructions			<b>Instruction Complexity</b>		Simple instructions, more instructions per program		<b>Memory Access</b>		Load/store architecture (only LOAD/STORE access memory)
	Memory operands in arithmetic operations			<b>Registers</b>		Large number of general-purpose registers		<b>Hardware Design</b>		Hardware simplicity for higher clock rates
	More complex hardware			<b>Pipelining</b>		Regular structure enables efficient pipelining				More difficult to pipeline

##### ARM Design Principles

- Simplicity enables high performance

- Regular instruction encoding aids decoding
- Load/store architecture simplifies memory access
- Large register file reduces memory traffic
- Consistent design across instruction types

## ARM Registers

### General-Purpose Registers

- **R0 to R15:** 16 registers total
- **32 bits wide:** Can hold integers, addresses, or data
- **R0-R12:** General computation and data storage
- **R13 (SP):** Stack Pointer - points to top of stack
- **R14 (LR):** Link Register - stores return address
- **R15 (PC):** Program Counter - address of next instruction

### Register Usage Conventions

Name	Register number	Usage	Preserved on call?
a1-a2	0-1	Argument / return result / scratch register	no
a3-a4	2-3	Argument / scratch register	no
v1-v8	4-11	Variables for local routine	yes
ip	12	Intra-procedure-call scratch register	no
sp	13	Stack pointer	yes
lr	14	Link Register (Return address)	yes
pc	15	Program Counter	n.a.

Figure 2.1: ARM Register Conventions

### R0-R3: Argument/result registers

- Pass parameters to functions
- Return values from functions
- Scratch registers (not preserved)

### R4-R11: Local variable registers

- Must be preserved across function calls
- Callee saves/restores if used

### R12: Intra-procedure-call scratch register

- Can be corrupted by function calls
- Not preserved

### R13 (SP): Stack Pointer

- Points to top of stack

- Must always be valid

#### **R14 (LR): Link Register**

- Stores return address on function call
- Contains address to return to

#### **R15 (PC): Program Counter**

- Always points to next instruction
- Modifying PC changes execution flow

#### **Why So Many Registers?**

- Reduces memory accesses (faster than cache/RAM)
- Enables register allocation by compiler
- Supports efficient function calls
- Improves performance through locality

### **Memory Organization**

#### **Little-Endian Byte Ordering**

- Least significant byte at lowest address
- Example: 0x12345678 stored as:

Address:	[base+0]	[base+1]	[base+2]	[base+3]
Content:	78	56	34	12

#### **textbfWord Alignment**

- Words are 32 bits (4 bytes)
- Word addresses should be multiples of 4
- Accessing unaligned words may cause errors or slowdown

#### **textbfAddress Space**

- 32-bit addresses can access  $2^{32}$  bytes = 4 GB
- Byte-addressable memory
- Instructions and data in same address space (Von Neumann architecture)

### 2.1.3 ARM Instruction Format

#### Instruction Structure

##### Fixed 32-Bit Length

- Every instruction exactly 32 bits
- Simplifies instruction fetch and decode
- Enables predictable pipeline operation

##### Typical Instruction Fields

[Condition] [Opcode] [Operands]  
4 bits      varies      varies

##### Example: ADD Instruction

ADD R1, R2, R3      ; R1 = R2 + R3
------------------------------------

Encoding includes:

- Condition code (usually "always")
- Opcode for ADD operation
- Destination register (R1)
- Source register 1 (R2)
- Source register 2 (R3)

#### Instruction Types

##### Data Processing Instructions

- Arithmetic: ADD, SUB, RSB (reverse subtract)
- Logical: AND, ORR, EOR (XOR), BIC (bit clear)
- Comparison: CMP, CMN, TST, TEQ
- Move: MOV, MVN (move negated)
- Shift/Rotate: LSL, LSR, ASR, ROR

##### Data Transfer Instructions

- Load: LDR (word), LDRB (byte), LDRH (halfword)
- Store: STR (word), STRB (byte), STRH (halfword)
- Multiple: LDM, STM (load/store multiple registers)

##### Control Flow Instructions

- Branch: B (unconditional), BEQ, BNE, BGE, BLT, etc.
- Function call: BL (branch and link)
- Return: MOV PC, LR

## Operand Types

### Register Operands

```
ADD R0, R1, R2      ; R0 = R1 + R2 (all registers)
```

### Immediate Operands

```
ADD R0, R1, #5      ; R0 = R1 + 5 (# indicates immediate)
MOV R2, #100        ; R2 = 100
```

### Immediate Value Constraints

- Limited to certain patterns due to 32-bit instruction encoding
- 8-bit immediate + 4-bit rotation
- Assembler warns if immediate cannot be encoded

### Shifted Register Operands

```
ADD R0, R1, R2, LSL #2      ; R0 = R1 + (R2 << 2)
SUB R3, R4, R5, LSR #1      ; R3 = R4 - (R5 >> 1)
```

## 2.1.4 Basic ARM Instructions

### Arithmetic Instructions

#### Addition

```
ADD Rd, Rn, Rm          ; Rd = Rn + Rm
ADD Rd, Rn, #imm         ; Rd = Rn + immediate
```

Examples:

```
ADD R0, R1, R2          ; R0 = R1 + R2
ADD R3, R3, #1          ; R3 = R3 + 1 (increment)
```

#### Subtraction

```
SUB Rd, Rn, Rm          ; Rd = Rn - Rm
SUB Rd, Rn, #imm         ; Rd = Rn - immediate
RSB Rd, Rn, #imm         ; Rd = immediate - Rn (reverse subtract)
```

Examples:

```
SUB R0, R1, R2          ; R0 = R1 - R2
SUB R4, R4, #10         ; R4 = R4 - 10 (decrement)
RSB R5, R6, #0          ; R5 = 0 - R6 (negate)
```

#### Multiplication (covered in later tutorials)

```
MUL Rd, Rn, Rm          ; Rd = Rn x Rm (lower 32 bits)
```

### Logical Instructions

#### AND Operation

```
AND Rd, Rn, Rm          ; Rd = Rn AND Rm
AND Rd, Rn, #imm         ; Rd = Rn AND immediate
```

Usage: Bit masking, clearing specific bits

Example:

```
AND R0, R0, #0xFF ; Keep only lower 8 bits
```

### OR Operation

```
ORR Rd, Rn, Rm ; Rd = Rn OR Rm (ORR in ARM)
ORR Rd, Rn, #imm ; Rd = Rn OR immediate
```

Usage: Setting specific bits

Example:

```
ORR R1, R1, #0x80 ; Set bit 7
```

### Exclusive OR

```
EOR Rd, Rn, Rm ; Rd = Rn XOR Rm
EOR Rd, Rn, #imm ; Rd = Rn XOR immediate
```

Usage: Toggling bits, fast comparison

Example:

```
EOR R2, R2, R2 ; R2 = 0 (XOR with itself)
```

### Move and Move Not

```
MOV Rd, Rm ; Rd = Rm
MOV Rd, #imm ; Rd = immediate
MVN Rd, Rm ; Rd = NOT Rm (bitwise complement)
```

Examples:

```
MOV R0, R1 ; Copy R1 to R0
MOV R2, #0 ; Clear R2
MVN R3, R4 ; R3 = ~R4 (invert all bits)
```

## Shift Operations

### Logical Shift Left (LSL)

```
LSL Rd, Rn, #shift ; Rd = Rn << shift
MOV Rd, Rn, LSL #shift
```

Effect: Multiplies by  $2^{\text{shift}}$

Example:

```
LSL R0, R1, #2 ; R0 = R1 x 4
```

### Logical Shift Right (LSR)

```
LSR Rd, Rn, #shift ; Rd = Rn >> shift (unsigned)
MOV Rd, Rn, LSR #shift
```

Effect: Divides by  $2^{\text{shift}}$  (unsigned)

Example:

```
LSR R0, R1, #3 ; R0 = R1 / 8
```

### Arithmetic Shift Right (ASR)

```
ASR Rd, Rn, #shift ; Rd = Rn >> shift (signed)
```

Effect: Divides by  $2^{\text{shift}}$ , preserves sign

Example:

```
ASR R0, R1, #2      ; R0 = R1 / 4 (signed)
```

### Rotate Right (ROR)

```
ROR Rd, Rn, #shift  ; Rotate Rn right by shift
```

Effect: Bits rotated off right end reappear at left

Example:

```
ROR R0, R1, #8      ; Rotate R1 right by 8 bits
```

## 2.1.5 Memory Access Instructions

### Load Instructions

#### Load Word (LDR)

```
LDR Rd, [Rn]          ; Rd = Memory[Rn]
LDR Rd, [Rn, #offset]; Rd = Memory[Rn + offset]
```

Examples:

```
LDR R0, [R1]          ; Load word from address in R1
LDR R2, [R3, #4]       ; Load from address R3+4
LDR R4, [R5, #-8]      ; Load from address R5-8
```

#### Load Byte (LDRB)

```
LDRB Rd, [Rn, #offset]; Load one byte, zero-extend to 32 bits
```

Example:

```
LDRB R0, [R1]          ; R0 = (byte at R1), upper 24 bits = 0
```

#### Load Halfword (LDRH)

```
LDRH Rd, [Rn, #offset]; Load 16 bits, zero-extend to 32 bits
```

Example:

```
LDRH R0, [R1, #2]      ; R0 = (halfword at R1+2), upper 16 bits = 0
```

### Pseudo-Instruction for Loading Addresses

```
LDR Rd, =label        ; Load address of label into Rd
LDR Rd, =value         ; Load 32-bit constant into Rd
```

Examples:

```
LDR R0, =array         ; R0 = address of array
LDR R1, =0x12345678    ; R1 = 0x12345678 (large immediate)
```

### Store Instructions

#### Store Word (STR)

```
STR Rd, [Rn]           ; Memory[Rn] = Rd
STR Rd, [Rn, #offset]; Memory[Rn + offset] = Rd
```

Examples:

```
STR R0, [R1]          ; Store R0 to address in R1
STR R2, [R3, #8]      ; Store R2 to address R3+8
```

### Store Byte (STRB)

```
STRB Rd, [Rn, #offset]; Store lower 8 bits of Rd
```

Example:

```
STRB R0, [R1]          ; Store lower byte of R0 to address R1
```

### Store Halfword (STRH)

```
STRH Rd, [Rn, #offset]; Store lower 16 bits of Rd
```

Example:

```
STRH R0, [R1, #4]      ; Store lower halfword of R0 to R1+4
```

## Addressing Modes

textbfOffset Addressing

```
LDR R0, [R1, #4]       ; R0 = Memory[R1 + 4], R1 unchanged
```

textbfPre-indexed Addressing

```
LDR R0, [R1, #4]!      ; R1 = R1 + 4, then R0 = Memory[R1]
                        ; ! indicates update base register
```

textbfPost-indexed Addressing

```
LDR R0, [R1], #4       ; R0 = Memory[R1], then R1 = R1 + 4
```

textbfRegister Offset

```
LDR R0, [R1, R2]       ; R0 = Memory[R1 + R2]
LDR R0, [R1, R2, LSL #2] ; R0 = Memory[R1 + (R2 << 2)]
```

## 2.1.6 Assembly Program Structure

### Directives

#### Section Directives

```
.text          ; Code section (instructions)
.data          ; Data section (initialized variables)
.bss           ; Uninitialized data section
```

#### Global and External

```
.global main    ; Make symbol visible to linker
.extern printf  ; Declare external symbol
```

#### Data Definition

```
.word value     ; Define 32-bit word
.byte value     ; Define byte
.asciz "string"  ; Define null-terminated string
.space n        ; Reserve n bytes of space
```



## Labels

### Purpose

- Mark locations in code or data
- Provide symbolic names for addresses
- Enable jumps and references

textbfSyntax

```
label:                ; Label for instruction
    MOV R0, #1
    ADD R1, R0, R2

array:                ; Label for data
    .word 1, 2, 3, 4
```

### Simple Program Example

```
.text
.global main

main:
    MOV R0, #5        ; R0 = 5
    MOV R1, #10       ; R1 = 10
    ADD R2, R0, R1    ; R2 = R0 + R1 = 15
    MOV R0, R2        ; R0 = R2 (return value)
    MOV PC, LR        ; Return from main

.data
message:
    .asciz "Hello, ARM!"
```

## 2.1.7 ARM Development Tools

### Toolchain Components

#### Cross-Compiler

- `arm-linux-gnueabi-gcc`: Compiles C to ARM code
- Runs on x86 PC, produces ARM binaries
- Necessary because development machine  $\neq$  target machine

textbfAssembler

- `arm-linux-gnueabi-as`: Assembles ARM assembly to object code
- Part of `binutils` package

textbfLinker

- `arm-linux-gnueabi-ld`: Links object files to executable

- Resolves symbols, combines code sections

textbfEmulator

- `qemu-arm`: Emulates ARM processor on x86
- Allows running ARM binaries on PC
- Useful for testing without ARM hardware

## Compilation Process

### From C to Executable

```
C Source (.c)
    --> [gcc -S]
Assembly (.s)
    --> [as]
Object Code (.o)
    --> [ld]
Executable (a.out)
    --> [qemu-arm]
Execution
```

### Command Examples

```
# Compile C to assembly
arm-linux-gnueabi-gcc -S program.c -o program.s

# Assemble to object code
arm-linux-gnueabi-as program.s -o program.o

# Link to executable
arm-linux-gnueabi-gcc program.o -o program

# Run with emulator
qemu-arm program
\
```

### One-Step Compilation

```
# Compile, assemble, and link in one command
arm-linux-gnueabi-gcc program.c -o program
\
```

## Debugging and Inspection

### GDB (GNU Debugger)

```
# Debug with QEMU and GDB
qemu-arm -g 1234 program &      # Start QEMU, wait for debugger
arm-linux-gnueabi-gdb program  # Start GDB
(gdb) target remote :1234       # Connect to QEMU
(gdb) break main                # Set breakpoint
(gdb) continue                  # Run to breakpoint
(gdb) step                      # Execute one instruction
(gdb) info registers            # Show register values
\
```

**Objdump**

```
# Disassemble binary to assembly
arm-linux-gnueabi-objdump -d program
\
```

**nm**

```
# List symbols in object file
arm-linux-gnueabi-nm program.o
\
```

**2.1.8 Programming in ARM Assembly****Translating C to ARM****C Code:**

```
int a = 5;
int b = 10;
int c = a + b;
\\end{lstlisting}

\\textbf{ARM Assembly:}

\\begin{lstlisting}[language=assembly]
    MOV R0, #5      ; a = 5
    MOV R1, #10     ; b = 10
    ADD R2, R0, R1   ; c = a + b
\\end{lstlisting}

\\textbf{C Code with Array:}

\\begin{lstlisting}[language=c]
int arr[3] = {1, 2, 3};
int x = arr[1];
\\end{lstlisting}

\\textbf{ARM Assembly:}

\\begin{lstlisting}[language=assembly]
.data
arr:
    .word 1, 2, 3

.text
    LDR R0, =arr      ; R0 = address of arr
    LDR R1, [R0, #4]  ; R1 = arr[1] (offset 4 bytes)
\\end{lstlisting}

\\subsubsection{Common Patterns}

\\textbf{Clearing a Register}

\\begin{lstlisting}[language=assembly]
MOV R0, #0           ; Method 1
EOR R0, R0, R0       ; Method 2 (XOR with itself)
\\end{lstlisting}
```

```
\textbf{Negating a Value}

\begin{lstlisting}language=Assembler
RSB R0, R0, #0      ; R0 = 0 - R0
MVN R0, R0          ; R0 = ~R0 (bitwise, not arithmetic)
ADD R0, R0, #1      ; Then add 1 (two's complement)
```

### Multiplying by Powers of 2

```
LSL R0, R1, #3      ; R0 = R1 x 8 (faster than MUL)
```

### Dividing by Powers of 2

```
LSR R0, R1, #2      ; R0 = R1 / 4 (unsigned)
ASR R0, R1, #2      ; R0 = R1 / 4 (signed)
```

### Swapping Two Registers

```
EOR R0, R0, R1      ; XOR-based swap (no temporary)
EOR R1, R0, R1
EOR R0, R0, R1
```

## 2.1.9 Key Takeaways

- **ARM follows RISC principles** - simple instructions, load/store architecture, large register file, fixed instruction length
- **16 registers (R0-R15)** with special purposes: R13 (SP), R14 (LR), R15 (PC), and calling conventions for R0-R11
- **16 general-purpose registers (R0-R15)** with special roles for SP, LR, and PC
- **Three main instruction categories** - data processing (arithmetic/logic), data transfer (load/store), control flow (branches)
- **Fixed 32-bit instruction format** simplifies hardware and enables efficient pipelining
- **Little-endian byte ordering** - least significant byte stored at lowest address
- **Immediate values** indicated by # symbol, with encoding constraints due to fixed instruction size
- **Memory access only through LOAD/STORE** - arithmetic operations work on registers only (load/store architecture)
- **Rich addressing modes** - offset, pre-indexed, post-indexed, register offset with optional shifts
- **Cross-compilation toolchain** - arm-linux-gnueabi-gcc, as, ld, and qemu-arm for development on x86
- **Assembly programming requires understanding** of register allocation, instruction selection, and calling conventions

### 2.1.10 Summary

ARM assembly language provides the low-level interface between software and hardware, revealing how high-level constructs translate to machine operations. The ARM architecture's RISC design emphasizes simplicity and regularity, with a uniform 32-bit instruction format, a generous 16-register set, and a clean separation between computation (using registers) and memory access (through explicit load/store instructions). Understanding ARM assembly is crucial for optimizing performance-critical code, implementing system-level software, and comprehending how processors execute programs. The development toolchain—including cross-compilers, assemblers, linkers, and emulators—enables efficient development and testing of ARM software. Mastering these fundamentals prepares us for more advanced topics including function calling conventions, stack management, and processor microarchitecture implementation.

## 2.2 Lecture 5: Number Representation and Instruction Encoding

*By Dr. Kisaru Liyanage*

### 2.2.1 Introduction

This lecture delves into how computers represent and manipulate data at the binary level. We explore number systems, two's complement representation for signed integers, instruction encoding formats in ARM assembly, and logical operations for bit manipulation. Understanding these fundamentals is essential for programming efficiently in assembly language and comprehending how processors execute arithmetic and logical operations.

### 2.2.2 Number Representation Systems

#### Unsigned Binary Integers

##### Binary System Basics

- Base-2 number system using digits 0 and 1
- Each bit position represents a power of 2
- Rightmost bit is least significant (LSB)
- Leftmost bit is most significant (MSB)

##### Place Value Calculation

Binary: 1011

```
Value = (1 * 2^3) + (0 * 2^2) + (1 * 2^1) + (1 * 2^0)
       = 8 + 0 + 2 + 1
       = 11 (decimal)
```

##### N-Bit Unsigned Range

- N bits can represent  $2^N$  different values
- Range: 0 to  $(2^N - 1)$
- 8 bits: 0 to 255
- 32 bits: 0 to 4,294,967,295

### Binary to Decimal Conversion

Example: 10110101

$$= 1 \cdot 128 + 0 \cdot 64 + 1 \cdot 32 + 1 \cdot 16 + 0 \cdot 8 + 1 \cdot 4 + 0 \cdot 2 + 1 \cdot 1$$

$$= 128 + 32 + 16 + 4 + 1$$

$$= 181$$

### Two's Complement Representation

#### Purpose of Two's Complement

- Represents both positive and negative integers
- Simplifies hardware (same adder for signed/unsigned)
- Unique zero representation
- Natural overflow behavior

#### Sign Bit

- MSB indicates sign
- MSB = 0: Positive number
- MSB = 1: Negative number

#### Positive Numbers

- Same as unsigned binary
- MSB is always 0
- Example: +5 in 8 bits = 00000101

#### Negative Numbers

- Represented as  $2^N - |value|$
- Example: -5 in 8 bits:

$$2^8 - 5 = 256 - 5 = 251 = 11111011$$

### Two's Complement Conversion

Method 1 (Invert and Add):

1. Write positive value in binary
2. Invert all bits (0→1, 1→0)
3. Add 1 to result

Example: -5 in 8 bits

+5:           00000101

Invert:       11111010

Add 1:       11111011 (this is -5)

Method 2 (Subtraction):

$$-5 = 2^8 - 5 = 256 - 5 = 251 = 11111011$$

### N-Bit Signed Range

- Range:  $-(2^{(N-1)})$  to  $+(2^{(N-1)} - 1)$
- 8 bits: -128 to +127
- 32 bits: -2,147,483,648 to +2,147,483,647

### Special Cases

- Zero: 00000000 (unique representation)
- Most negative: 10000000 (-128 in 8 bits)
  - Has no positive counterpart!
  - Negating gives overflow

## Sign Extension

### Purpose

- Extend smaller signed value to larger width
- Preserve numerical value
- Required when loading bytes/halfwords into 32-bit registers

### Process

- Replicate the sign bit (MSB) to fill new bits
- Preserves positive/negative value

### Examples

8-bit to 32-bit:

00000101 (+5) → 00000000 00000000 00000000 00000101 (+5)  
 11111011 (-5) → 11111111 11111111 11111111 11111011 (-5)

### ARM Instructions for Sign Extension

- **LDRH**: Load halfword (16 bits), zero-extend to 32 bits
- **LDRSH**: Load signed halfword, sign-extend to 32 bits
- **LDRB**: Load byte (8 bits), zero-extend to 32 bits
- **LDRSB**: Load signed byte, sign-extend to 32 bits

### Example Usage

```
LDRH R0, [R1]      ; R0 = 0x0000ABCD (zero-extended)
LDRSH R0, [R1]     ; R0 = 0xFFFFABCD (sign-extended if bit 15 = 1)

LDRB R0, [R1]      ; R0 = 0x000000AB (zero-extended)
LDRSB R0, [R1]     ; R0 = 0xFFFFFAB (sign-extended if bit 7 = 1)
```

## Hexadecimal Notation

### Why Hexadecimal?

- Compact representation of binary
- One hex digit = 4 binary bits
- Easier to read than long binary strings
- Common in programming and debugging

### Hex Digits

Binary	Hex	Decimal
-----	-----	-----
0000	0	0
0001	1	1
0010	2	2
0011	3	3
0100	4	4
0101	5	5
0110	6	6
0111	7	7
1000	8	8
1001	9	9
1010	A	10
1011	B	11
1100	C	12
1101	D	13
1110	E	14
1111	F	15

### Conversion Examples

Binary: 1011 0110 1101 0010  
 Hex:     B     6     D     2  
 Result: 0xB6D2

Hex: 0x3F  
 Binary: 0011 1111  
 Decimal: 63

### ARM Hexadecimal Usage

```
MOV R0, #0xFF      ; R0 = 255
MOV R1, #0x100     ; R1 = 256
LDR R2, =0xDEADBEEF ; R2 = 3735928559
```

## 2.2.3 ARM Instruction Encoding

### Fixed-Length Instructions

#### 32-Bit Instruction Format



- Every ARM instruction is exactly 32 bits
- Simplifies instruction fetch and decode
- Enables efficient pipelining

#### **Advantages**

- Predictable instruction boundaries
- Simple PC increment (always +4)
- Fast decode logic

#### **Trade-offs**

- Some instructions may "waste" bits
- Immediate values limited in size
- Code density lower than variable-length (e.g., x86)

### **Data Processing Instruction Format**

#### **Format Structure**

[Cond] [00] [I] [Opcode] [S] [Rn] [Rd] [Operand2]  
 4-bit 2 1 4-bit 1 4 4 12-bit

#### **Field Descriptions**

##### **Condition (4 bits, bits 28-31)**

- Conditional execution feature
- 0000 = EQ (equal, Z=1)
- 0001 = NE (not equal, Z=0)
- 1010 = GE (greater or equal, signed)
- 1110 = AL (always execute, default)

##### **I bit (bit 25)**

- 0 = Operand2 is register
- 1 = Operand2 is immediate value

##### **Opcode (4 bits, bits 21-24)**

- Specifies operation (AND, EOR, SUB, ADD, etc.)
- 0100 = ADD
- 0010 = SUB
- 0000 = AND
- 1100 = ORR

**S bit (bit 20)**

- 0 = Don't update condition flags
- 1 = Update flags (CPSR)

**Rn (4 bits, bits 16-19)**

- First operand register number
- 0000 = R0, 0001 = R1, etc.

**Rd (4 bits, bits 12-15)**

- Destination register number

**Operand2 (12 bits, bits 0-11)**

- If I=0: Shift amount and second register
- If I=1: 8-bit immediate + 4-bit rotation

**Example: ADD R0, R1, R2**

Encoding fields:

- Cond: 1110 (always)
- I: 0 (register operand)
- Opcode: 0100 (ADD)
- S: 0 (don't update flags)
- Rn: 0001 (R1)
- Rd: 0000 (R0)
- Operand2: 0002 (R2, no shift)

Result: 0xE0810002

**Data Transfer Instruction Format****Format Structure**

[Cond] [01] [I] [P] [U] [B] [W] [L] [Rn] [Rd] [Offset]  
 4-bit 2 1 1 1 1 1 1 4 4 12-bit

**Key Fields****L bit (bit 20)**

- 0 = Store (STR)
- 1 = Load (LDR)

**B bit (bit 22)**

- 0 = Word transfer (32 bits)
- 1 = Byte transfer (8 bits)

**P bit (bit 24)**

- 0 = Post-indexed addressing
- 1 = Pre-indexed or offset addressing

**U bit (bit 23)**

- 0 = Subtract offset from base
- 1 = Add offset to base

**W bit (bit 21)**

- 0 = No write-back
- 1 = Write-back (update base register)

**Rn (base register)**

- Contains memory address or base address

**Rd (data register)**

- For Load: Destination register
- For Store: Source register

**Offset (12 bits)**

- Memory address offset
- Can be immediate or register

**Example: LDR R0, [R1, #4]**

Encoding fields:

- Cond: 1110 (always)
- L: 1 (load)
- B: 0 (word)
- P: 1 (offset addressing)
- U: 1 (add offset)
- Rn: 0001 (R1)
- Rd: 0000 (R0)
- Offset: 004 (immediate 4)

Result: 0xE5910004

**Immediate Value Encoding****Challenge**

- 32-bit instruction must fit: opcode, registers, immediate
- Cannot fit full 32-bit immediate

**ARM Solution: 8-bit + 4-bit Rotation**

- Immediate field: 12 bits total

- Lower 8 bits: Immediate value (0-255)
- Upper 4 bits: Rotation amount (0-15)
- Rotation: Right by  $(2 \times rotation\ field)$  bits

### Calculation

Actual Value = Immediate \* ROR (2 \* Rotation)

### Examples

Immediate=0xFF, Rotation=0:

Value = 0xFF ROR 0 = 0x000000FF

Immediate=0xFF, Rotation=8:

Value = 0xFF ROR 16 = 0x00FF0000

Immediate=0xFF, Rotation=12:

Value = 0xFF ROR 24 = 0xFF000000

### Allowed immediates

- Not all 32-bit values can be encoded
- Valid: 0xFF, 0xFF00, 0xFF0000, 0xFF000000
- Valid: 0xFF000000FF (rotation wraps around)
- Invalid: 0x123 (cannot be formed by rotation)

### Assembler Handling

- Assembler checks if immediate is valid
- Gives error if immediate cannot be encoded
- Use LDR pseudo-instruction for arbitrary values:

LDR R0, =0x12345678 ; Loads from literal pool

## 2.2.4 Logical Operations

### Bitwise AND

#### Operation

- Performs logical AND on each bit pair
- Result bit = 1 only if both input bits are 1

#### Truth Table

A	B	A AND B
0	0	0
0	1	0
1	0	0
1	1	1

**ARM Instruction**

```
AND Rd, Rn, Rm      ; Rd = Rn AND Rm
AND Rd, Rn, #imm     ; Rd = Rn AND immediate
```

**Common Uses****Bit Masking (Extract Specific Bits)**

```
; Extract lower 8 bits of R1
MOV R0, R1
AND R0, R0, #0xFF    ; R0 = R1 & 0xFF (keep bits 0-7)

; Extract bits 8-15
MOV R0, R1
AND R0, R0, #0xFF00  ; R0 = R1 & 0xFF00 (keep bits 8-15)
```

**Clearing Specific Bits**

```
; Clear bit 5 of R1
AND R1, R1, #0xFFFFFDF ; Bit 5 mask: ~(1 << 5)
```

**Checking if Bit Set**

```
AND R2, R1, #0x80    ; Check if bit 7 is set
CMP R2, #0            ; Compare with zero
BEQ bit_clear        ; Branch if bit was clear
```

**Bitwise OR****Operation**

- Performs logical OR on each bit pair
- Result bit = 1 if either input bit is 1

**Truth Table**

A	B	A OR B
0	0	0
0	1	1
1	0	1
1	1	1

**ARM Instruction**

```
ORR Rd, Rn, Rm      ; Rd = Rn OR Rm (ORR in ARM)
ORR Rd, Rn, #imm     ; Rd = Rn OR immediate
```

**Common Uses****Setting Specific Bits**

```
; Set bit 3 of R1
ORR R1, R1, #0x08    ; Bit 3 mask: (1 << 3) = 0x08

; Set bits 4 and 5
ORR R1, R1, #0x30    ; Mask: 0x30 = 0b00110000
```

### Combining Values

```
; Combine lower byte of R1 with upper bytes of R2
AND R1, R1, #0xFF          ; Keep only lower byte
AND R2, R2, #0xFFFFF00    ; Keep only upper bytes
ORR R0, R1, R2              ; Combine
```

### Bitwise XOR (Exclusive OR)

#### Operation

- Performs logical XOR on each bit pair
- Result bit = 1 if input bits differ

#### Truth Table

A	B	A XOR B
0	0	0
0	1	1
1	0	1
1	1	0

### ARM Instruction

```
EOR Rd, Rn, Rm          ; Rd = Rn EOR Rm (EOR in ARM)
EOR Rd, Rn, #imm        ; Rd = Rn EOR immediate
```

#### Common Uses

##### Toggling Specific Bits

```
; Toggle bit 2 of R1
EOR R1, R1, #0x04      ; Bit 2 mask: (1 << 2)
; If bit was 0, becomes 1; if was 1, becomes 0
```

#### Fast Zero

```
EOR R0, R0, R0          ; R0 = 0 (XOR with itself)
```

### Comparison

```
; Check if R1 and R2 are equal
EOR R3, R1, R2          ; R3 = R1 XOR R2
CMP R3, #0              ; If R3 = 0, R1 == R2
BEQ values_equal
```

### Swapping Without Temporary

```
; Swap R0 and R1 without using another register
EOR R0, R0, R1
EOR R1, R0, R1
EOR R0, R0, R1
; Now R0 and R1 are swapped
```

**Bitwise NOT****Operation**

- Inverts all bits (0->1, 1->0)
- Also called complement

**ARM Instruction**

```

MVN Rd, Rm          ; Rd = NOT Rm (Move Not)
MVN Rd, #imm        ; Rd = NOT immediate

```

**Common Uses****Creating Bit Masks**

```

; Create mask with all bits set except bit 3
MOV R0, #0x08        ; 0x08 = 0b00001000
MVN R1, R0           ; R1 = 0xFFFFFFFF (all except bit 3)

```

**Negation (with ADD)**

```

; Negate R1 (two's complement)
MVN R1, R1           ; Invert all bits
ADD R1, R1, #1       ; Add 1
; Now R1 = -R1 (original)

```

**Shift Operations****Logical Shift Left (LSL)**

```

LSL Rd, Rn, #shift   ; Rd = Rn << shift
MOV Rd, Rn, LSL #shift

```

- Shifts bits left, fills right with zeros
- Each shift left multiplies by 2
- Example: 0b00001010 LSL 2 = 0b00101000

**Logical Shift Right (LSR)**

```

LSR Rd, Rn, #shift   ; Rd = Rn >> shift (unsigned)
MOV Rd, Rn, LSR #shift

```

- Shifts bits right, fills left with zeros
- Each shift right divides by 2 (unsigned)
- Example: 0b10100000 LSR 2 = 0b00101000

**Arithmetic Shift Right (ASR)**

```

ASR Rd, Rn, #shift   ; Rd = Rn >> shift (signed)

```

- Shifts bits right, fills left with sign bit

- Preserves sign for signed division
- Example: 0b11110000 ASR 2 = 0b11111100 (sign preserved)

### Rotate Right (ROR)

ROR Rd, Rn, #shift ; Rotate Rn right by shift

- Bits shifted out right reappear at left
- No information lost
- Example: 0b10000001 ROR 1 = 0b11000000

### Common Shift Applications

#### Fast Multiplication/Division by Powers of 2

```
LSL R0, R1, #3      ; R0 = R1 * 8 (2^3)
LSR R0, R1, #2      ; R0 = R1 / 4 (unsigned)
ASR R0, R1, #2      ; R0 = R1 / 4 (signed)
```

### Bit Extraction

```
; Extract bits 8-11 from R1
LSR R0, R1, #8      ; Shift bits 8-11 to bits 0-3
AND R0, R0, #0xF    ; Mask to keep only 4 bits
```

### Bit Positioning

```
; Move bit 0 to bit 7
LSL R0, R1, #7      ; Shift left 7 positions
AND R0, R0, #0x80   ; Keep only bit 7
```

## 2.2.5 Practical Bit Manipulation Examples

### Extracting Bit Fields

#### Extract bits 16-23

```
LSR R0, R1, #16     ; Shift right to position
AND R0, R0, #0xFF   ; Mask to 8 bits
```

#### Extract bits 4-9 (6 bits)

```
LSR R0, R1, #4      ; Shift to position 0
AND R0, R0, #0x3F   ; Mask to 6 bits (0b111111)
```

### Setting and Clearing Bits

#### Set bits 8-15

```
ORR R1, R1, #0xFF00 ; Set bits 8-15
```

#### Clear bits 16-23

```
LDR R0, =0xFF00FFFF ; Mask with bits 16-23 clear
AND R1, R1, R0       ; Clear bits 16-23 of R1
```

#### Toggle bits 0-7

```
EOR R1, R1, #0xFF   ; Toggle lower byte
```



### Checking Flags

Check if any of bits 4-7 are set

```

AND R2, R1, #0xF0    ; Mask bits 4-7
CMP R2, #0            ; Check if zero
BNE bits_set         ; Branch if any bit was set

```

Check if specific pattern matches

```

; Check if bits 8-11 are 0b1010
LSR R0, R1, #8        ; Position bits
AND R0, R0, #0xF      ; Mask 4 bits
CMP R0, #0xA          ; Compare with 0b1010
BEQ pattern_match

```

### Color Packing/Unpacking

Pack RGB values (8 bits each)

```

; R0 = Red, R1 = Green, R2 = Blue
LSL R1, R1, #8        ; Green << 8
LSL R2, R2, #16       ; Blue << 16
ORR R3, R0, R1        ; Combine Red and Green
ORR R3, R3, R2        ; Combine with Blue
; R3 now contains 0x00BBGRR

```

Unpack RGB values

```

; R0 contains 0x00BBGRR
AND R1, R0, #0xFF     ; Extract Red
LSR R2, R0, #8        ; Extract Green
AND R2, R2, #0xFF
LSR R3, R0, #16       ; Extract Blue
AND R3, R3, #0xFF

```

#### 2.2.6 Key Takeaways

1. **Unsigned binary integers** represent values from 0 to  $2^N - 1$  using  $N$  bits.
2. **Two's complement** represents signed integers, with MSB as sign bit and range  $-(2^{(N-1)})$  to  $+(2^{(N-1)} - 1)$ .
3. **Sign extension** preserves value when expanding narrower signed values to wider registers.
4. **Hexadecimal notation** provides compact representation with one hex digit per 4 binary bits.
5. **ARM instructions are fixed 32-bit length**, simplifying fetch/decode but limiting immediate values.
6. **Data processing format** includes condition, opcode, source/destination registers, and operand.
7. **Data transfer format** specifies load/store, byte/word, addressing mode, and offset.

8. **Immediate encoding** uses 8-bit value + 4-bit rotation, limiting which constants can be encoded directly.
9. **Bitwise AND** used for masking (extracting specific bits) and clearing bits.
10. **Bitwise OR** used for setting specific bits and combining values.
11. **Bitwise XOR** used for toggling bits, fast zero, and comparisons.
12. **Shift operations** enable fast multiplication/division by powers of 2 and bit positioning.
13. **Bit manipulation** is fundamental for low-level programming, hardware control, and optimization.
14. **Understanding encoding** helps write efficient assembly and debug machine code issues.

### 2.2.7 Summary

Number representation and instruction encoding form the foundation of low-level programming. Two's complement enables efficient signed arithmetic with simple hardware, while sign extension preserves values across different data sizes. ARM's fixed 32-bit instruction format provides regularity but imposes constraints on immediate values, solved through clever encoding schemes. Logical operations—AND, OR, XOR, and NOT—combined with shift operations, provide powerful tools for bit manipulation essential in systems programming, embedded development, and performance optimization. Mastering these concepts enables efficient assembly programming and deeper understanding of how high-level operations translate to machine instructions. These fundamentals prepare us for more complex topics including branching, function calls, and memory management.

## 2.3 Lecture 6: Branching and Control Flow

*By Dr. Kisaru Liyanage*

### 2.3.1 Introduction

Control flow is what distinguishes computers from simple calculators—the ability to make decisions and alter execution based on conditions. This lecture explores conditional operations and branching in ARM assembly, covering comparison instructions, conditional branches, loop implementation, and PC-relative addressing. Understanding these mechanisms is essential for translating high-level control structures (if statements, loops) into assembly code and for comprehending how processors implement dynamic program behavior.

### 2.3.2 Fundamentals of Conditional Execution

#### Decision-Making in Computers

#### What Makes Computers Powerful

- Ability to make decisions based on data
- Execute different instructions depending on conditions
- Implement if statements, loops, and function calls
- Respond dynamically to input and computed values

### Control Flow Concepts

- **Sequential execution:** Default behavior (PC += 4)
- **Conditional branching:** Jump if condition is true
- **Unconditional branching:** Always jump
- **Function calls:** Branch with return address saving

### Program Status Register (PSR)

#### Status Flags

- **N (Negative):** Set if result is negative (bit 31 = 1)
- **Z (Zero):** Set if result is zero
- **C (Carry):** Set if unsigned overflow occurred
- **V (oVerflow):** Set if signed overflow occurred

#### How Flags Are Set

- Comparison instructions (CMP, CMN, TST, TEQ)
- Arithmetic/logic instructions with S suffix (ADDS, SUBS)
- Flags reflect the result of the operation
- Used by subsequent conditional branches

#### Example

```
CMP R1, R2          ; Compare R1 and R2 (computes R1 - R2)
                   ; Sets flags based on result
```

If R1 = 5, R2 = 3:

- Result of R1 - R2 = 2 (positive, non-zero)
- N = 0 (not negative)
- Z = 0 (not zero)
- C = 1 (no borrow needed)
- V = 0 (no overflow)

### 2.3.3 Comparison Instructions

#### Compare (CMP)

##### Syntax

```
CMP Rn, Rm          ; Compare Rn with Rm
CMP Rn, #imm        ; Compare Rn with immediate
```

##### Operation

- Performs  $Rn - Rm$  (subtraction)
- Updates PSR flags based on result
- Does NOT store the result
- Does NOT modify any register

##### Example Usage

```
MOV R1, #10
MOV R2, #5
CMP R1, R2          ; Compares 10 with 5
                   ; Result:  $10 - 5 = 5$  (positive, non-zero)
                   ; Z = 0, N = 0
```

#### Compare Negative (CMN)

##### Syntax

```
CMN Rn, Rm          ; Compare Negative
CMN Rn, #imm
```

##### Operation

- Performs  $Rn + Rm$  (addition)
- Updates PSR flags
- Equivalent to `CMP Rn, -Rm`
- Useful for checking if sum equals zero

#### Test (TST)

##### Syntax

```
TST Rn, Rm          ; Test bits
TST Rn, #imm
```

##### Operation

- Performs  $Rn \text{ AND } Rm$  (bitwise AND)
- Updates PSR flags
- Result not stored
- Used to test if specific bits are set

##### Example: Check if bit 5 is set

```
TST R1, #0x20        ; Test bit 5
BEQ bit_clear        ; Branch if bit was clear (Z=1)
```

**Test Equivalence (TEQ)****Syntax**

```
TEQ Rn, Rm          ; Test Equivalence
TEQ Rn, #imm
```

**Operation**

- Performs Rn XOR Rm (exclusive OR)
- Updates PSR flags
- Z=1 if values are equal
- Used to compare values without affecting C or V flags

**2.3.4 Conditional Branch Instructions****Branch if Equal (BEQ)****Syntax**

```
BEQ label          ; Branch if equal (Z=1)
```

**Condition**

- Branches if Zero flag is set ( $Z = 1$ )
- Typically used after CMP to check equality

**Example**

```
CMP R1, R2          ; Compare R1 and R2
BEQ equal_label     ; Jump to equal_label if R1 == R2
; Code if not equal
equal_label:
; Code if equal
```

**Branch if Not Equal (BNE)****Syntax**

```
BNE label          ; Branch if not equal (Z=0)
```

**Condition**

- Branches if Zero flag is clear ( $Z = 0$ )
- Opposite of BEQ

**Example**

```
CMP R3, #0
BNE not_zero       ; Jump if R3 != 0
; Code if R3 is zero
not_zero:
; Code if R3 is non-zero
```

## Signed Comparison Branches

### Branch if Greater or Equal (BGE)

```
BGE label          ; Branch if Rn >= Rm (signed)
                   ; Condition: N == V
```

### Branch if Less Than (BLT)

```
BLT label          ; Branch if Rn < Rm (signed)
                   ; Condition: N != V
```

### Branch if Greater Than (BGT)

```
BGT label          ; Branch if Rn > Rm (signed)
                   ; Condition: Z==0 AND N==V
```

### Branch if Less or Equal (BLE)

```
BLE label          ; Branch if Rn <= Rm (signed)
                   ; Condition: Z==1 OR N!=V
```

### Example

```
CMP R1, R2
BGE greater_equal  ; Branch if R1 >= R2 (signed)
; Code if R1 < R2
greater_equal:
; Code if R1 >= R2
```

## Unsigned Comparison Branches

### Branch if Higher or Same (BHS) (also called BCS - Branch if Carry Set)

```
BHS label          ; Branch if Rn >= Rm (unsigned)
                   ; Condition: C == 1
```

### Branch if Lower (BLO) (also called BCC - Branch if Carry Clear)

```
BLO label          ; Branch if Rn < Rm (unsigned)
                   ; Condition: C == 0
```

### Branch if Higher (BHI)

```
BHI label          ; Branch if Rn > Rm (unsigned)
                   ; Condition: C==1 AND Z==0
```

### Branch if Lower or Same (BLS)

```
BLS label          ; Branch if Rn <= Rm (unsigned)
                   ; Condition: C==0 OR Z==1
```

**Signed vs. Unsigned Example****Key Difference**

```

MOV R0, #0xFFFFFFFF ; R0 = -1 (signed) or 4,294,967,295 (unsigned)
MOV R1, #1           ; R1 = 1
CMP R0, R1

BLO lower_unsigned   ; BRANCH NOT TAKEN
                     ; Unsigned: 4,294,967,295 > 1

BLT less_signed      ; BRANCH TAKEN
                     ; Signed: -1 < 1

```

**When to Use Each**

- **Signed:** Comparing integers that can be negative (temperatures, offsets, differences)
- **Unsigned:** Comparing addresses, array indices, sizes, counts

**Unconditional Branch****Syntax**

```
B label ; Branch always
```

**Purpose**

- Jump without checking any condition
- Skip code sections
- Implement infinite loops
- Return to loop start

**Example**

```

B end ; Skip this section
; Code to skip
end:
; Continue execution here

```

**2.3.5 Labels in Assembly****Label Definition****Purpose**

- Mark specific instruction locations
- Provide symbolic names for addresses
- Enable branches and data references

**Syntax**

```
label:                ; Label definition (note colon)
    MOV R0, #1        ; Instruction at this label
```

### Naming Rules

- Can be almost any identifier
- Common conventions: loop, exit, done, L1, L2
- Cannot conflict with instruction mnemonics
- Case-sensitive

### Example

```
start:
    MOV R0, #0
loop:
    ADD R0, R0, #1
    CMP R0, #10
    BLT loop          ; Branch to loop label
    B start           ; Branch to start label
```

### Label Resolution

#### Assembly Process

1. First pass: Record label addresses
2. Second pass: Replace labels with addresses
3. Calculate offsets for PC-relative branches

### Virtual Addresses

- Assembler assigns virtual addresses from 0
- First instruction: address 0
- Second instruction: address 4
- Third instruction: address 8
- Physical addresses determined at load time

## 2.3.6 Implementing Control Structures

### If Statement

#### C Code

```
if (i == j)
    f = g + h;
else
    f = g - h;
```

#### ARM Assembly (Method 1: Branch on False)



```

    CMP R3, R4      ; Compare i (R3) and j (R4)
    BNE else       ; Branch to else if not equal
    ADD R0, R1, R2  ; f = g + h (then clause)
    B exit         ; Skip else clause

```

else:

```

    SUB R0, R1, R2  ; f = g - h (else clause)

```

exit:

```

    ; Continue...

```

### ARM Assembly (Method 2: Conditional Execution)

```

    CMP R3, R4      ; Compare i and j
    ADDEQ R0, R1, R2 ; f = g + h (executed only if equal)
    SUBNE R0, R1, R2 ; f = g - h (executed only if not equal)

```

### If-Else Ladder

#### C Code

```

if (x < 0)
    result = -1;
else if (x == 0)
    result = 0;
else
    result = 1;

```

#### ARM Assembly

```

    CMP R1, #0      ; Compare x with 0
    BLT negative    ; Branch if x < 0
    BEQ zero        ; Branch if x == 0
    ; x > 0
    MOV R0, #1
    B done

```

negative:

```

    MOV R0, #-1
    B done

```

zero:

```

    MOV R0, #0

```

done:

```

    ; Continue...

```

### While Loop

#### C Code

```

while (i < n) {
    sum += i;
    i++;
}

```

#### ARM Assembly

```

loop:
    CMP R1, R2        ; Compare i (R1) with n (R2)
    BGE end_loop      ; Exit if i >= n
    ADD R0, R0, R1     ; sum = sum + i
    ADD R1, R1, #1     ; i++
    B loop            ; Branch back to loop start
end_loop:
    ; Continue...

```

## For Loop

### C Code

```

for (i = 0; i < 10; i++) {
    sum += i;
}

```

### ARM Assembly

```

    MOV R1, #0        ; i = 0 (initialization)

for_loop:
    CMP R1, #10       ; Compare i with 10
    BGE end_for       ; Exit if i >= 10
    ADD R0, R0, R1     ; sum = sum + i (loop body)
    ADD R1, R1, #1     ; i++ (increment)
    B for_loop        ; Branch back to loop start
end_for:
    ; Continue...

```

## Do-While Loop

### C Code

```

do {
    sum += i;
    i++;
} while (i < n);

```

### ARM Assembly

```

do_loop:
    ADD R0, R0, R1     ; sum = sum + i (loop body first)
    ADD R1, R1, #1     ; i++
    CMP R1, R2        ; Compare i with n
    BLT do_loop       ; Branch back if i < n
    ; Continue...

```

### Key Difference from While

- Body executes at least once
- Condition checked at end, not beginning

### 2.3.7 Array Access in Loops

#### Static Array Indexing

##### C Code

```
while (save[i] == k)
    i++;
```

##### ARM Assembly

```
; R6 = base address of save array
; R3 = i (index)
; R5 = k (comparison value)
```

loop:

```
ADD R12, R6, R3, LSL #2 ; address = base + (i * 4)
LDR R0, [R12, #0]      ; R0 = save[i]
CMP R0, R5              ; Compare save[i] with k
BNE exit                ; Exit if not equal
ADD R3, R3, #1          ; i++
B loop                  ; Continue loop
```

exit:

```
; Continue...
```

##### Dynamic Offset Calculation

- R3, LSL #2 means  $R3 \times 4$  (shift left 2 = multiply by 4)
- Words are 4 bytes, so array element i is at  $\text{base} + (i \times 4)$
- Efficient: shift is faster than multiplication

#### Array Traversal

##### C Code

```
int sum = 0;
for (int i = 0; i < 10; i++) {
    sum += arr[i];
}
```

##### ARM Assembly

```
LDR R6, =arr ; R6 = base address of array
MOV R0, #0   ; sum = 0
MOV R1, #0   ; i = 0
```

loop:

```
CMP R1, #10
BGE done
ADD R12, R6, R1, LSL #2 ; address = base + i*4
LDR R2, [R12]           ; R2 = arr[i]
ADD R0, R0, R2          ; sum += arr[i]
ADD R1, R1, #1          ; i++
B loop
```

done:

```
; R0 contains sum
```

### 2.3.8 PC-Relative Addressing

#### Branch Instruction Encoding

##### 32-Bit Format

[Cond] [1010] [Offset]  
 4-bit 4-bit 24-bit

##### Fields

- **Cond:** Condition code (EQ, NE, LT, etc.)
- **1010:** Fixed format field for branch
- **Offset:** 24-bit signed offset

#### Address Calculation

##### Problem with Absolute Addressing

- 24 bits can address  $2^{24} = 16$  MB
- Limits program size to 16 MB
- Fixed addresses complicate relocation

##### PC-Relative Solution

- Store offset from current PC, not absolute address
- Target = PC + offset
- Can branch  $\pm 16$  MB from current instruction
- Total program can exceed 16 MB

##### Offset Calculation

Offset = (Target Address - PC) / 4

##### Why Divide by 4?

- All instructions are 4-byte aligned
- Least significant 2 bits always 00
- Omit these bits in encoding
- Effective range:  $\pm 64$  MB (24-bit offset  $\times 4$ )

##### Example

Current PC: 0x1000

Target: 0x1020

Offset = (0x1020 - 0x1000) / 4 = 0x20 / 4 = 8 instructions

Encoded offset in branch instruction: 8

At execution: PC = 0x1000 + (8  $\times$  4) = 0x1020

**Advantages of PC-Relative****Position-Independent Code**

- Code can load at any address
- Branches remain correct regardless of location
- Essential for libraries and shared code

**Simplified Linking**

- Linker doesn't need to patch all branches
- Only external function calls need adjustment

**Branch Locality**

- Most branches are to nearby instructions
- PC-relative naturally handles this case
- Absolute addressing wastes bits for nearby targets

**2.3.9 Conditional Execution (Alternative to Branching)****Conditional Instruction Suffixes****Concept**

- Add condition code to instruction mnemonic
- Instruction executes only if condition is true
- Otherwise, instruction is skipped (NOP)

**Available Suffixes**

- EQ (equal), NE (not equal)
- GT, LT, GE, LE (signed comparisons)
- HI, LO, HS, LS (unsigned comparisons)
- Many others (see ARM documentation)

**Examples**

```
CMP R1, R2
ADDEQ R0, R3, R4    ; Execute ADD only if R1 == R2
SUBNE R0, R3, R4    ; Execute SUB only if R1 != R2
MOVGT R5, #10       ; Execute MOV only if R1 > R2
```

## Conditional Execution Example

### C Code

```
if (a == b)
    max = a;
else
    max = b;
```

### Method 1: Branching

```
    CMP R1, R2        ; Compare a and b
    BNE else
    MOV R0, R1        ; max = a
    B done

else:
    MOV R0, R2        ; max = b
done:
```

### Method 2: Conditional Execution

```
    CMP R1, R2        ; Compare a and b
    MOVEQ R0, R1      ; max = a (if equal)
    MOVNE R0, R2      ; max = b (if not equal)
```

## Advantages and Limitations

### Advantages

- More compact code (fewer instructions)
- No branch misprediction penalty
- Faster for simple conditions
- Clearer intent in some cases

### Limitations

- Only works for simple, short sequences
- Cannot conditionally execute blocks of code
- All conditional instructions must fit in pipeline
- May execute both paths (but discard one result)

### When to Use

- Simple assignments
- Min/max operations
- Short computations with single result
- Performance-critical paths where branches hurt

### 2.3.10 Basic Blocks

#### Definition

##### Basic Block Characteristics

- Sequence of instructions with:
  - No embedded branches (except possibly at end)
  - No branch targets (except possibly at beginning)
- Executed atomically: all or nothing
- Single entry point, single exit point

#### Example

```
; Basic Block 1 (entry point)
MOV R0, #0
MOV R1, #10
CMP R1, #10
BNE block2          ; Exit point of block 1

; Basic Block 2 (entry and exit point)
block2:
    ADD R0, R0, #1
    CMP R0, R1
    BLT block2      ; Exit point of block 2
```

#### Importance in Compilation

##### Compiler Optimizations

- Identify basic blocks for analysis
- Optimize within blocks (register allocation, scheduling)
- Build control flow graph from blocks
- Apply inter-block optimizations

##### Processor Optimizations

- Predict block execution
- Prefetch instructions in block
- Schedule instructions more aggressively
- Reduce branch overhead

### 2.3.11 Key Takeaways

1. **Conditional execution** distinguishes computers from calculators, enabling decision-making and dynamic behavior.
2. **CMP instruction** sets PSR flags by performing subtraction without storing the result.
3. **Conditional branches** (BEQ, BNE, BGE, BLT, etc.) check PSR flags to decide whether to jump.
4. **Signed vs. unsigned branches** interpret the same bit patterns differently based on context.
5. **Labels** provide symbolic names for addresses, enabling readable branch targets.
6. **If statements** translate to compare + conditional branch + unconditional branch to skip alternate path.
7. **Loops** use compare + conditional branch (to exit) + unconditional branch (to continue).
8. **Array access** in loops uses dynamic offset calculation with shifts (LSL #2 for word arrays).
9. **PC-relative addressing** stores branch offset from current PC, enabling position-independent code and large programs.
10. **Word-based offsets** effectively quadruple branch range by encoding instruction count instead of byte offset.
11. **Conditional execution** provides alternative to branching for simple cases, improving performance and code density.
12. **Basic blocks** are atomic instruction sequences used by compilers and processors for optimization.
13. **Branch locality** means most branches target nearby instructions, making PC-relative addressing natural and efficient.

### 2.3.12 Summary

Branching and conditional execution form the foundation of program control flow, translating high-level constructs like if statements and loops into machine instructions. The ARM architecture provides a rich set of conditional branches for both signed and unsigned comparisons, enabling efficient implementation of diverse control structures. Understanding the distinction between comparison (which sets flags) and branching (which checks flags) is essential for correct assembly programming. PC-relative addressing solves program size limitations while enabling position-independent code, and conditional execution offers a performant alternative to branching for simple cases. Mastering these concepts is crucial for translating algorithms into assembly code, optimizing performance-critical sections, and understanding how processors implement dynamic program behavior. These fundamentals prepare us for more advanced topics including function calls, stack management, and processor pipelining.

## 2.4 Lecture 7: Function Call and Return

*By Dr. Kisaru Liyanage*



### 2.4.1 Introduction

Function calling is a fundamental mechanism that enables modular programming and code reuse. This lecture explores how ARM assembly implements function calls, covering parameter passing, return value handling, the call stack, register preservation conventions, and recursion. Understanding these mechanisms is essential for translating high-level function-based programs into assembly and for comprehending how processors manage execution context across function boundaries.

### 2.4.2 Function Calling Fundamentals

#### Function Calling Steps

##### Complete Call Sequence

1. **Place parameters** in argument registers (R0-R3)
2. **Transfer control** to callee function using BL
3. **Acquire stack storage** for temporary values
4. **Back up registers** that need preservation (R4-R11)
5. **Perform function operations** (the actual work)
6. **Place result** in return register (R0)
7. **Restore backed-up registers** from stack
8. **Return to caller** using MOV PC, LR

##### Why This Complexity?

- Enables nested and recursive function calls
- Protects caller's data in registers
- Provides local storage for function variables
- Supports arbitrary call depth

#### Why Use Functions?

##### Benefits

- **Code reuse:** Write once, call many times
- **Modularity:** Break complex problems into manageable pieces
- **Abstraction:** Hide implementation details
- **Maintainability:** Easier to debug and modify

##### Example

```
int add(int a, int b) {  
    return a + b;  
}  
  
int main() {  
    int result = add(5, 3);  // Function call  
}
```

### 2.4.3 ARM Register Conventions

#### Register Usage Rules

#### Register Classification

- R0-R1: Arguments and return results
- Caller does NOT expect these preserved
  - Scratch registers
- R2-R3: Additional arguments
- Also scratch registers
  - Caller does NOT expect preservation
- R4-R11: Local variables
- MUST be preserved across function calls
  - Callee saves if it uses these registers
- R12: Intra-procedure-call scratch register
- Can be corrupted by function calls
  - Not preserved
- R13 (SP): Stack Pointer
- Points to top of stack
  - MUST always be valid
- R14 (LR): Link Register
- Stores return address
  - Set by BL instruction
- R15 (PC): Program Counter
- Next instruction address
  - Modified to return from function

#### Shared Register File

#### Key Concept

- ALL functions share the SAME 16 registers
- No separate register sets per function
- Registers are a shared resource requiring careful management

#### Implications

- Functions must coordinate register usage
- Conventions prevent conflicts
- Callee must preserve certain registers (R4-R11)
- Caller can assume R4-R11 unchanged after call

#### Example Scenario

```

main:
    MOV R4, #10      ; main uses R4
    MOV R0, #5       ; Pass argument
    BL function      ; Call function
    ; R4 still contains 10 (guaranteed)
    ADD R5, R4, R0    ; Use preserved R4 and return value

function:
    ; Must preserve R4 if we use it
    ; Can freely modify R0-R3, R12
    MOV R0, #20       ; Return value
    MOV PC, LR        ; Return

```

### 2.4.4 Function Call Instructions

#### Branch and Link (BL)

##### Syntax

```
BL function_label    ; Branch and Link
```

##### Operation

1. **Save return address:** LR = address of next instruction
2. **Jump to function:** PC = function\_label address

##### Example

```

MOV R0, #10          ; Address: 0x1000
BL fun               ; Address: 0x1004
ADD R1, R0, #5       ; Address: 0x1008 (return point)

```

```

fun:
    ; LR contains 0x1008 (address after BL)
    ; Function code here
    MOV PC, LR       ; Return to 0x1008

```

##### Why "Link"?

- Creates a "link" back to caller
- LR provides the connection
- Enables function to return

#### Return from Function

##### Basic Return

```
MOV PC, LR          ; Copy LR to PC
```

##### Operation

- PC = LR (jump to return address)
- Execution continues at instruction after BL
- Simple and fast

#### Alternative (older ARM)

`BX LR` ; Branch and Exchange

### 2.4.5 Parameter Passing

#### Using R0-R3

##### Convention

- First 4 arguments in R0-R3
- Arguments loaded before BL instruction
- Callee reads R0-R3 to get parameters

#### Example: Two Parameters

```
int multiply(int a, int b) {
    return a * b;
}
```

```
int result = multiply(6, 7);
```

#### ARM Assembly

```
MOV R0, #6      ; First argument (a)
MOV R1, #7      ; Second argument (b)
BL multiply     ; Call function
                ; R0 now contains result (42)
```

```
multiply:
    MUL R0, R0, R1    ; R0 = R0 * R1
    MOV PC, LR       ; Return
```

#### More Than 4 Arguments

##### Solution: Use Stack

- Arguments 1-4 in R0-R3
- Additional arguments pushed to stack
- Callee reads from stack

#### Example: 6 Arguments

```
int sum6(int a, int b, int c, int d, int e, int f) {
    return a + b + c + d + e + f;
}
```

**ARM Assembly**

```

MOV R0, #1      ; arg1
MOV R1, #2      ; arg2
MOV R2, #3      ; arg3
MOV R3, #4      ; arg4
MOV R4, #5
MOV R5, #6
SUB SP, SP, #8  ; Space for 2 more args
STR R4, [SP, #0] ; arg5 on stack
STR R5, [SP, #4] ; arg6 on stack
BL sum6
ADD SP, SP, #8  ; Clean up stack

```

```

sum6:
    ; R0-R3 have first 4 args
    ; Load arg5 and arg6 from stack
    LDR R4, [SP, #0] ; arg5
    LDR R5, [SP, #4] ; arg6
    ADD R0, R0, R1
    ADD R0, R0, R2
    ADD R0, R0, R3
    ADD R0, R0, R4
    ADD R0, R0, R5
    MOV PC, LR

```

**2.4.6 Return Values****Primary Return Register (R0)****Convention**

- Result placed in R0
- Caller reads R0 after function returns
- Works for 32-bit values

**Example**

```

add:
    ADD R0, R0, R1    ; R0 = R0 + R1
    MOV PC, LR        ; Return with result in R0

main:
    MOV R0, #10
    MOV R1, #20
    BL add             ; Call function
    ; R0 now contains 30

```

## 64-Bit Return Values

### Convention

- Lower 32 bits in R0
- Upper 32 bits in R1
- Example: 64-bit integer or two 32-bit values

### Example

```
long long multiply64(int a, int b) {
    return (long long)a * b;
}
```

### ARM Assembly

```
multiply64:
    SMULL R0, R1, R0, R1    ; Signed multiply long
    ; R0 = lower 32 bits
    ; R1 = upper 32 bits
    MOV PC, LR
```

## 2.4.7 The Stack

### Stack Structure

#### Definition

- Last In, First Out (LIFO) data structure
- Part of main memory
- Used for temporary storage

#### Characteristics

- **Starts at high address:** Top of memory
- **Grows downward:** Toward lower addresses
- **Stack Pointer (SP/R13):** Points to top of stack
- **Dynamic size:** Grows and shrinks as needed

#### Memory Layout

High Address

```
+-----+
| Stack | SP points here
|       | (grows downward)
|       |
+-----+
| Heap  |
|       | (grows upward)
+-----+
| Data  | (static variables)
+-----+
| Text  | (instructions)
+-----+
```

Low Address

**Stack Uses****Primary Purposes**

1. **Saving register values** (preserve R4-R11)
2. **Storing local variables** (arrays, structures)
3. **Preserving return addresses** (nested calls)
4. **Extra function arguments** (beyond R0-R3)
5. **Storing local arrays** that don't fit in registers

**2.4.8 Stack Operations****Allocating Stack Space (Pushing)****Decrement Stack Pointer**

```
SUB SP, SP, #4      ; Allocate 4 bytes (1 register)
SUB SP, SP, #12     ; Allocate 12 bytes (3 registers)
```

**Why Subtract?**

- Stack grows toward lower addresses
- Allocating space moves SP downward
- Each 32-bit register needs 4 bytes

**Storing Values to Stack****Single Register**

```
SUB SP, SP, #4      ; Allocate space
STR R4, [SP, #0]    ; Store R4 at top of stack
```

**Multiple Registers**

```
SUB SP, SP, #12     ; Space for 3 registers
STR R4, [SP, #0]    ; Store R4
STR R5, [SP, #4]    ; Store R5
STR R6, [SP, #8]    ; Store R6
```

**Push Multiple (Convenient)**

```
PUSH {R4-R6}       ; Allocate and store in one instruction
```

## Loading Values from Stack

### Single Register

```
LDR R4, [SP, #0]      ; Load R4 from stack
ADD SP, SP, #4         ; Release space
```

### Multiple Registers

```
LDR R4, [SP, #0]      ; Restore R4
LDR R5, [SP, #4]      ; Restore R5
LDR R6, [SP, #8]      ; Restore R6
ADD SP, SP, #12       ; Release space
```

### Pop Multiple

```
POP {R4-R6}           ; Restore and release in one instruction
```

## Stack Space Lifecycle

### Pattern

1. **Allocate:** SUB SP, SP, #n
2. **Use:** STR/LDR with [SP, offset]
3. **Release:** ADD SP, SP, #n

### Important: Balance

- Every SUB must have corresponding ADD
- Unbalanced stack causes bugs and crashes
- SP must be restored before return

## 2.4.9 Register Preservation

### Why Preserve R4-R11?

#### Problem

- All functions share same registers
- Main function may be using R4-R11
- Called function needs registers for its work
- Must not corrupt caller's data

#### Solution

- Callee saves R4-R11 to stack at function start
- Uses registers freely during execution
- Restores R4-R11 from stack before return
- Caller expects R4-R11 unchanged



**Preservation Pattern****Function Template**

```

function:
    ; Prologue: Save registers
    SUB SP, SP, #12      ; Allocate space
    STR R4, [SP, #0]     ; Save R4
    STR R5, [SP, #4]     ; Save R5
    STR R6, [SP, #8]     ; Save R6

    ; Function body: Use R4-R6 freely
    ; ...

    ; Epilogue: Restore registers
    LDR R4, [SP, #0]     ; Restore R4
    LDR R5, [SP, #4]     ; Restore R5
    LDR R6, [SP, #8]     ; Restore R6
    ADD SP, SP, #12      ; Release space
    MOV PC, LR           ; Return

```

**Optimization**

- Only preserve registers actually used
- If function doesn't use R5, don't save/restore it
- Saves stack space and execution time

**2.4.10 Nested Function Calls (Non-Leaf Functions)****The Problem****Leaf Function**

- Doesn't call other functions
- LR preserved automatically (not overwritten)
- Simple return: MOV PC, LR

**Non-Leaf Function**

- Calls other functions
- BL overwrites LR with new return address
- Original LR lost!
- Cannot return to original caller

**Example Problem**

```

main:
    BL funcA          ; LR = address after this BL

funcA:
    ; LR contains return address to main
    BL funcB          ; LR OVERWRITTEN with return to funcA!
    MOV PC, LR        ; Returns to funcA, not main (WRONG!)

funcB:
    MOV PC, LR        ; Correctly returns to funcA

```

### Solution: Save LR to Stack

#### Pattern

```

function:
    ; Save LR first!
    SUB SP, SP, #4
    STR LR, [SP, #0]

    ; Now safe to call other functions
    BL other_function

    ; Restore LR before return
    LDR LR, [SP, #0]
    ADD SP, SP, #4
    MOV PC, LR

```

### Complete Example

```

main:
    MOV R0, #5
    BL outer          ; LR = return_to_main
    ; Execution returns here

outer:
    SUB SP, SP, #4
    STR LR, [SP, #0] ; Save LR (return_to_main)

    MOV R1, R0
    ADD R0, R0, #10
    BL inner          ; LR = return_to_outer (overwrites!)

    ADD R0, R0, R1
    LDR LR, [SP, #0] ; Restore LR (return_to_main)
    ADD SP, SP, #4
    MOV PC, LR        ; Returns to main

inner:
    MUL R0, R0, R0
    MOV PC, LR        ; Returns to outer

```

### 2.4.11 Recursion Example: Factorial

#### Factorial Function

##### C Code

```
int fact(int n) {
    if (n <= 1)
        return 1;
    else
        return n * fact(n-1);
}
```

##### Key Points

- Base case:  $n \leq 1$ , return 1
- Recursive case: return  $n \times \text{fact}(n - 1)$
- Each call creates new stack frame
- Stack unwinds as recursion returns

#### ARM Assembly Implementation

```
fact:
    ; Save LR and n
    SUB SP, SP, #8
    STR LR, [SP, #4]    ; Save return address
    STR R0, [SP, #0]    ; Save n

    ; Base case: if (n <= 1) return 1
    CMP R0, #1
    BGT recursive
    MOV R0, #1          ; Return 1
    B fact_end

recursive:
    ; Recursive case: n * fact(n-1)
    SUB R0, R0, #1      ; n-1
    BL fact              ; fact(n-1)
    LDR R1, [SP, #0]    ; Restore original n
    MUL R0, R0, R1      ; n * fact(n-1)

fact_end:
    ; Restore and return
    LDR LR, [SP, #4]
    ADD SP, SP, #8
    MOV PC, LR
```

#### Stack Growth During Recursion

Call: fact(3)

Initial: SP = 0x1000

fact(3) call:

SP = 0xFF8: [LR\_main, 3]

fact(2) call:

SP = 0xFF0: [LR\_fact3, 2]

fact(1) call:

SP = 0xFE8: [LR\_fact2, 1]

Base case returns 1

Unwinds to fact(2): returns  $1 \times 2 = 2$

Unwinds to fact(3): returns  $2 \times 3 = 6$

Returns to main with result 6

Final: SP = 0x1000 (restored)

### Stack Space Per Call

- 8 bytes (LR + n)
- fact(5) needs  $5 \times 8 = 40$  bytes
- fact(10) needs 80 bytes
- Deep recursion can overflow stack!

## 2.4.12 Memory Layout and Stack vs. Heap

### Complete Memory Layout

High Address (0xFFFFFFFF)

+-----+	
Reserved	OS and system
+-----+	
Stack	<- SP (grows down)
v	Automatic storage
	Function call data
	Local variables
(unused)	
^	
Heap	Dynamic allocation
	malloc/free, new/delete
+-----+	
Static Data	Global variables
	String constants
+-----+	
Text	Program instructions
(Code)	Read-only
+-----+	

Low Address (0x00000000)

**Stack Characteristics****Automatic Storage**

- Allocated when function called
- Released when function returns
- Managed automatically by compiler/runtime

**Fast Access**

- Fixed addressing pattern
- SP always points to top
- Simple offset calculations

**Limited Size**

- Typically 1-8 MB
- Stack overflow if exceeded
- Recursion depth limited

**Scope**

- Local to function
- Not accessible after return
- Perfect for temporary data

**Heap Characteristics****Dynamic Allocation**

- malloc/free in C
- new/delete in C++
- Programmer controls lifetime

**Flexible Size**

- Can grow large (limited by available memory)
- Variable-sized allocations

**Manual Management**

- Must explicitly free memory
- Memory leaks if not freed
- Fragmentation possible

**Global Scope**

- Persists until explicitly freed
- Can pass pointers across functions
- Suitable for data structures

### 2.4.13 Key Takeaways

1. **Function calling requires** parameter passing, return value handling, and register preservation.
2. **R0-R3 for arguments and returns** - caller doesn't expect preservation.
3. **R4-R11 must be preserved** by callee if used, protecting caller's data.
4. **BL instruction** saves return address in LR and jumps to function.
5. **Return via MOV PC, LR** copies link register to program counter.
6. **Stack is LIFO structure** growing downward from high addresses, pointed to by SP.
7. **Stack usage** includes saving registers, local variables, return addresses, and extra arguments.
8. **Allocate with SUB SP, release with ADD SP** - must balance allocations and releases.
9. **Non-leaf functions** must save LR to stack before making nested calls.
10. **Recursion** creates multiple stack frames, one per call, unwinding as calls return.
11. **Stack vs. Heap** - stack is automatic/local/fast/limited, heap is manual/global/flexible/larger.
12. **Register conventions** enable modularity and prevent conflicts in shared register file.

### 2.4.14 Summary

Function calling mechanisms enable modular programming by providing structured ways to pass control, data, and return values between code sections. ARM's register conventions balance efficiency (passing arguments in registers) with safety (preserving callee-saved registers). The stack provides essential temporary storage for register preservation, local variables, and handling nested calls including recursion. Understanding these mechanisms is crucial for translating high-level function-based code to assembly, optimizing performance, and debugging stack-related issues. The interplay between registers, stack, and calling conventions forms the foundation for understanding how real programs execute, preparing us for more advanced topics like exception handling, operating systems, and compiler optimization.

## 2.5 Lecture 8: Memory Access and String Operations

*By Dr. Kisaru Liyanage*

### 2.5.1 Introduction

This lecture explores character data handling, string operations, and the compilation/linking/loading process. We examine byte and half-word memory operations, implement string manipulation functions, use library functions like `scanf` and `printf`, and understand how programs transform from source code to executable binaries. These topics bridge high-level programming concepts and low-level assembly implementation, essential for systems programming and understanding program execution.

## 2.5.2 Character Data and Encoding

### ASCII Encoding

#### Basic 7-Bit Standard

- Represents 128 characters using 7 bits ( $2^7 = 128$ )
- 95 graphic symbols (printable): A-Z, a-z, 0-9, punctuation
- 33 control symbols: newline ('`\n`'), tab ('`\t`'), null ('`\0`')
- Most basic and widely used encoding

#### ASCII Examples

'A' = 65 (0x41)  
'a' = 97 (0x61)  
'0' = 48 (0x30)  
'\n' = 10 (0x0A)  
'\0' = 0 (0x00) - null terminator

### Latin-1 Encoding

#### Extended 8-Bit Standard

- Supports 256 characters using 8 bits ( $2^8 = 256$ )
- Includes all ASCII characters (first 128)
- Adds 96 additional graphic characters
- European language support (accented characters)

### Unicode Encoding

#### Modern Universal Standard

- Uses 32-bit character set ( $2^{32}$  possible characters)
- Can represent most world alphabets and symbols
- Used in modern languages (Java, C++, Python 3)
- Variable-length encodings: UTF-8, UTF-16
- UTF-8: 1-4 bytes per character (backward compatible with ASCII)

#### Why Unicode?

- Global language support
- Emoji and special symbols
- Mathematical and technical symbols
- Historical scripts and languages

### 2.5.3 Byte Load/Store Operations

#### Load Register Byte (LDRB)

##### Syntax

LDRB Rd, [Rn, #offset] ; Load byte from memory

##### Operation

- Reads 8 bits (1 byte) from memory
- Fills upper 24 bits of register with zeros (zero-extension)
- Lower 8 bits contain the loaded byte

##### Example

```
; Memory[0x1000] = 0x42 ('B')
LDR R1, =0x1000
LDRB R0, [R1]
; R0 = 0x00000042
```

##### Use Cases

- Loading single characters
- Reading byte arrays
- Accessing packed data structures
- I/O port access

#### Store Register Byte (STRB)

##### Syntax

STRB Rd, [Rn, #offset] ; Store byte to memory

##### Operation

- Writes lower 8 bits of register to memory
- Upper 24 bits of register ignored
- Only affects 1 byte in memory

##### Example

```
MOV R0, #0x41 ; 'A'
LDR R1, =0x2000
STRB R0, [R1] ; Memory[0x2000] = 0x41
```



**Load Register Signed Byte (LDRSB)****Syntax**

LDRSB Rd, [Rn, #offset] ; Load signed byte

**Operation**

- Loads 8 bits from memory
- Replicates sign bit (bit 7) to fill upper 24 bits
- Sign-extension preserves signed value

**Example**

```
; Memory[0x1000] = 0xFE (-2 in signed byte)
LDR R1, =0x1000
LDRSB R0, [R1]
; R0 = 0xFFFFFFF0 (-2 in 32-bit signed)

; Memory[0x1001] = 0x7F (+127)
LDRSB R0, [R1, #1]
; R0 = 0x0000007F (+127)
```

**When to Use**

- Loading signed characters (int8\_t)
- Temperature values
- Signed offsets or deltas

**Memory Alignment****LDRB Advantages**

- Can access ANY byte address
- No alignment requirement
- Example: addresses 0, 1, 2, 3, 4, 5...

**LDR Requirement**

- Must use word-aligned addresses (multiples of 4)
- Valid addresses: 0, 4, 8, 12, 16...
- Invalid: 1, 2, 3, 5, 6, 7, 9...
- Unaligned access causes errors or performance penalties

### 2.5.4 Half-Word Load/Store Operations

#### Load Register Half-word (LDRH)

##### Syntax

LDRH Rd, [Rn, #offset] ; Load 16 bits

##### Operation

- Loads 16 bits (2 bytes) from memory
- Fills upper 16 bits with zeros (zero-extension)

##### Example

```
; Memory[0x1000-0x1001] = 0xABCD
LDR R1, =0x1000
LDRH R0, [R1]
; R0 = 0x0000ABCD
```

##### Use Cases

- Loading 16-bit integers (short)
- Unicode characters (UTF-16)
- 16-bit data types

#### Store Register Half-word (STRH)

##### Syntax

STRH Rd, [Rn, #offset] ; Store 16 bits

##### Operation

- Writes lower 16 bits of register to memory
- Upper 16 bits ignored

##### Example

```
MOV R0, #0x1234
LDR R1, =0x2000
STRH R0, [R1]
; Memory[0x2000-0x2001] = 0x1234
```

**Load Register Signed Half-word (LDRSH)****Syntax**

LDRSH Rd, [Rn, #offset] ; Load signed 16-bit

**Operation**

- Loads 16 bits from memory
- Replicates sign bit (bit 15) to upper 16 bits
- Sign-extension

**Example**

```
; Memory = 0x8000 (-32768 as signed 16-bit)
LDRSH R0, [R1]
; R0 = 0xFFFF8000 (-32768 as signed 32-bit)
```

**2.5.5 String Copy Example (strcpy)****C Implementation****Code**

```
void strcpy(char x[], char y[]) {
    int i = 0;
    while ((x[i] = y[i]) != '\0') {
        i++;
    }
}
```

**Algorithm**

1. Copy characters from y to x one at a time
2. Stop when null terminator ('\0') encountered
3. Null terminator also copied

**ARM Assembly Implementation****Register Allocation**

```
R0: Base address of x (destination)
R1: Base address of y (source)
R4: Loop counter i
R2: Address of y[i]
R3: Value of y[i]
R12: Address of x[i]
```

**Complete Assembly**

```

strcpy:
    ; Prologue: Save R4 (must preserve)
    SUB SP, SP, #4
    STR R4, [SP, #0]

    ; Initialize counter
    MOV R4, #0          ; i = 0

loop:
    ; Calculate address of y[i]
    ADD R2, R4, R1      ; R2 = y + i

    ; Load y[i]
    LDRB R3, [R2, #0]   ; R3 = y[i]

    ; Calculate address of x[i]
    ADD R12, R4, R0     ; R12 = x + i

    ; Store to x[i]
    STRB R3, [R12, #0]  ; x[i] = y[i]

    ; Check for null terminator
    CMP R3, #0          ; Is y[i] == '\0'?
    BEQ done            ; If yes, exit loop

    ; Increment counter
    ADD R4, R4, #1      ; i++
    B loop              ; Continue loop

done:
    ; Epilogue: Restore R4
    LDR R4, [SP, #0]
    ADD SP, SP, #4
    MOV PC, LR          ; Return

```

## Key Points

### Why LDRB/STRB?

- Strings are char arrays (8-bit elements)
- Must use byte operations

### Register Preservation

- R4 must be saved/restored (callee-saved)
- R12 doesn't need preservation (scratch register)

### Offsets Are Immediate

- [R2, #0] uses immediate offset (hash symbol)
- Cannot use [R2, R3] directly without proper syntax

### 2.5.6 Library Functions: scanf and printf

#### scanf Function

##### Purpose

- Read input from standard input (keyboard)
- Parse formatted input

##### C Signature

```
int scanf(const char *format, ...);
```

##### Arguments

- R0: Address of format string ("%d", "%c", "%s", etc.)
- R1: Address where to store input (NOT the value!)
- R2, R3: Additional addresses for more inputs

##### Example: Read Integer

##### C Code

```
int x;
scanf("%d", &x); // Note: &x (address of x)
```

##### ARM Assembly

```
.data
formatS: .asciz "%d"

.text
; Allocate space for variable
SUB SP, SP, #4 ; Space for x

; Load format string address
LDR R0, =formatS ; R0 = address of "%d"

; Load stack address
MOV R1, SP ; R1 = address where to store

; Call scanf
BL scanf

; Value now stored at [SP]
LDR R2, [SP, #0] ; R2 = x
```

#### printf Function

##### Purpose

- Print output to standard output (screen)
- Format and display data

**C Signature**

```
int printf(const char *format, ...);
```

**Arguments**

- R0: Address of format string
- R1, R2, R3: VALUES to print (not addresses!)

**Example: Print Integer****C Code**

```
printf("Result: %d\n", result);
```

**ARM Assembly**

```
.data
formatP: .asciz "Result: %d\n"

.text
    ; Load value to print
    LDR R1, [SP, #0]      ; R1 = result (value, not address)

    ; Release stack space (before printf)
    ADD SP, SP, #4

    ; Load format string
    LDR R0, =formatP

    ; Call printf
    BL printf
```

**Data Section and Format Strings****Data Section**

```
.data
formatS: .asciz "%d"      ; Input format
formatP: .asciz "Result: %d\n" ; Output format
array: .word 1, 2, 3, 4   ; Array
message: .asciz "Hello"   ; String
```

**.asciz Directive**

- Defines null-terminated string
- Automatically adds '\0' at end
- Stored in data section (separate from code)

**Pseudo-Operation: LDR Rd, =label**

```
LDR R0, =formatS      ; Loads ADDRESS of formatS into R0
```

- Not actual LDR instruction
- Assembler converts to appropriate instruction(s)
- Loads memory address (pointer), not content

**scanf vs printf Argument Differences****scanf: Needs Addresses**

```
SUB SP, SP, #4
MOV R1, SP          ; R1 = address (where to store)
BL scanf
```

**printf: Needs Values**

```
LDR R1, [SP]        ; R1 = value (what to print)
BL printf
```

**Why This Difference?**

- scanf modifies variables (needs addresses to write to)
- printf only reads values (copies values)

**Calling Convention Rules****Follow Exact Order**

- R0 first, R1 second, R2 third, R3 fourth
- Library functions expect specific argument positions
- Assembly won't check violations
- Mistakes cause wrong behavior or crashes

**Know Function Signatures**

- Read documentation
- Understand parameter types and order
- Match assembly to C function prototype

**2.5.7 Compilation, Linking, and Loading****Translation Overview****Complete Process**

```
C Program (.c)
  | [Compiler]
  v
Assembly (.s)
  | [Assembler]
  v
Object Module (.o)
  | [Linker]
  v
Executable (a.out)
  | [Loader]
  v
Memory (running program)
```

## Compiler

### Function

- Converts high-level C code to assembly language
- Complex task requiring sophisticated algorithms
- Performs optimizations

### Optimizations

- Register allocation
- Instruction selection
- Loop unrolling
- Dead code elimination
- Function inlining

### Example

```
int add(int a, int b) {  
    return a + b;  
}
```

↓ Compiler

```
add:  
    ADD R0, R0, R1  
    MOV PC, LR
```

## Assembler

### Function

- Converts assembly language to machine code (binary)
- Simpler than compilation (mostly 1-to-1 mapping)
- Produces object modules

### Tasks

1. Translate instructions to binary opcodes
2. Resolve local labels to addresses
3. Generate symbol table
4. Create relocation information

### Object Module Structure Header

- Describes contents and sizes



**Text Segment**

- Machine instructions (binary code)

**Static Data Segment**

- Initialized global variables
- String constants (format strings)

**Relocation Info**

- Instructions/data depending on absolute addresses
- Needed when program loaded at different address

**Symbol Table**

- Global definitions: functions, variables defined here
- External references: functions/variables from other modules
- Enables linking

**Debug Info**

- Maps machine code to source code lines
- Used by debuggers (gdb)

**Linker****Function**

- Combines multiple object modules into executable
- Links program code with library code

**Tasks****1. Merge Segments**

program.o:	lib.o:	Result:
[Text1]	[Text2]	-> [Text1+Text2]
[Data1]	[Data2]	-> [Data1+Data2]

**2. Resolve Labels**

- Convert symbolic names to actual addresses
- Example: "printf" → 0x80481234
- Processor only understands addresses

**3. Patch References**

- Update function calls to correct addresses
- Fix relocatable addresses
- May leave some for loader

## Static vs Dynamic Linking

### Static Linking

- Library code copied into executable at compile time
- Larger executable files
- Self-contained (no external dependencies)
- All code in one file

#### Advantages

- No runtime dependencies
- Faster load time
- Predictable behavior

#### Disadvantages

- Large file sizes
- No benefit from library updates
- Memory duplication across programs

### Dynamic Linking

- Library code loaded at runtime when called
- Smaller executables
- Shared libraries on system

#### Advantages

- Smaller executables
- Shared libraries (less memory usage)
- Automatic library updates
- Less disk space

#### Disadvantages

- Requires libraries installed on system
- "DLL not found" errors
- Slightly slower initial load

### DLL (Dynamic Link Library) - Windows

- File extension: .dll
- Shared by multiple programs
- Must be present on system
- Example: msvcrt.dll (C runtime library)

**Loader****Function**

- Loads executable from disk into memory
- Prepares program for execution
- Initializes execution environment

**Loading Steps****1. Read Header**

- Determine segment sizes
- Text segment size
- Data segment size
- Other metadata

**2. Create Virtual Address Space**

- Allocate memory for program
- Set up page tables (virtual memory)
- Map segments to physical memory

**3. Copy Segments to Memory**

- Text segment (instructions)
- Initialized data
- Set up page table entries
- Mark text as read-only, data as read-write

**4. Set Up Arguments on Stack**

- Command-line arguments: argc, argv
- Environment variables
- Initial stack frame

**Example**

```
./program arg1 arg2
```

- argc = 3
- argv[0] = "./program"
- argv[1] = "arg1"
- argv[2] = "arg2"

**5. Initialize Registers**

- Set up register file
- PC points to entry point (`_start`)
- SP points to top of stack
- Other registers to initial values

## 6. Jump to Startup Routine

- Calls C runtime initialization
- Sets up standard library
- Calls `main()` function
- When main returns, calls `exit()`

### 2.5.8 Exercises

#### Common String Operations

##### String Length

```
int strlen(char *s) {
    int len = 0;
    while (s[len] != '\0')
        len++;
    return len;
}
```

##### String Reverse

```
void strrev(char *s) {
    int len = strlen(s);
    for (int i = 0; i < len/2; i++) {
        char temp = s[i];
        s[i] = s[len-1-i];
        s[len-1-i] = temp;
    }
}
```

#### Integer I/O

##### Read Two Integers, Print Sum

```
; Read x and y
; Print x + y
```

##### Read n, Print 1 to n

```
; Read n
; Loop from 1 to n, print each
```

### Skills Required

- Character data handling (LDRB/STRB)
- String manipulation
- scanf for input
- printf for output
- Stack management
- Function calling conventions
- Loop implementation
- Array indexing

### 2.5.9 Key Takeaways

1. **ASCII (7-bit), Latin-1 (8-bit), Unicode (32-bit)** represent character data with increasing capacity.
2. **LDRB/STRB for byte operations**, LDRH/STRH for half-words - smaller than word operations.
3. **Byte operations don't require alignment** unlike word operations (LDR/STR).
4. **Sign extension (LDRSB/LDRSH)** replicates sign bit to preserve signed values.
5. **Strings in C are char arrays** terminated with null character (`'\0' = 0`).
6. **scanf and printf are library functions** called via BL instruction.
7. **scanf needs addresses (where to store)**, printf needs values (what to print).
8. **Format strings stored in .data section** using `.asciz` directive.
9. **Arguments passed in R0-R3** following ARM calling convention.
10. **Compilation chain: Compile → Assemble → Link → Load → Execute.**
11. **Static linking includes libraries in executable**, dynamic linking loads at runtime.
12. **Loader sets up virtual memory, copies segments, initializes stack** with arguments.

### 2.5.10 Summary

Character data handling and library function usage bridge high-level programming concepts and assembly implementation. Understanding byte/half-word operations enables efficient string manipulation and compact data storage. The scanf/printf functions demonstrate how assembly code interfaces with system libraries, requiring careful attention to calling conventions and argument types. The compilation, linking, and loading process reveals how source code transforms into running programs, involving multiple stages with distinct responsibilities. Static and dynamic linking represent different trade-offs between self-containment and flexibility. These concepts are essential for systems programming, understanding program structure, and debugging low-level issues. This knowledge prepares us for advanced topics including operating systems, compilers, and system-level optimization.



# Chapter 3

## Processor Design

### 3.1 Lecture 9: Microarchitecture and Datapath Design

*By Dr. Isuru Nawinne*

#### 3.1.1 Introduction

This lecture transitions from instruction set architecture (ISA) to microarchitecture—the hardware implementation of the ISA. We explore how to build a processor that executes MIPS instructions, covering instruction formats, digital logic fundamentals, datapath construction, and single-cycle processor design. Understanding microarchitecture reveals how software instructions translate to hardware operations and provides the foundation for studying advanced processor designs including pipelining and superscalar execution.

#### 3.1.2 Course Context and MIPS ISA

##### Transition to Hardware Implementation

**Previous Focus:** ARM ISA

- Instruction set
- Assembly programming
- Software perspective

**Current Focus:** MIPS Microarchitecture

- Hardware implementation
- Processor design
- Hardware perspective

**Why MIPS for Hardware Study?**

- Simpler than ARM (educational clarity)
- Clean RISC design
- Well-documented architecture
- Concepts apply to all processors

## MIPS Instruction Categories

### Three Instruction Types (based on encoding)

#### I-Type (Immediate)

- Contains one immediate operand
- Covers data processing, data transfer, control flow
- Examples: ADDI, LW, SW, BEQ
- Most common type

#### R-Type (Register)

- All operands are registers
- Primarily arithmetic and logic
- Examples: ADD, SUB, AND, OR
- Opcode always 0, funct field specifies operation

#### J-Type (Jump)

- Jump instructions
- Examples: J, JAL
- 26-bit address field

### Contrast with ARM

- ARM: Data processing, data transfer, flow control
- MIPS: I-type, R-type, J-type
- Different classification philosophy

## MIPS Instruction Encoding

### Fixed 32-Bit Length

- Every instruction exactly 32 bits
- Simplifies fetch and decode
- Enables efficient pipelining

### R-Type Format

[Opcode] [RS] [RT] [RD] [SHAMT] [Funct]  
 6 bits   5   5   5   5   6 bits

Fields:

- **Opcode:** Always 0 for R-type
- **RS:** Source register 1 (5 bits for 32 registers)



- **RT**: Source register 2
- **RD**: Destination register
- **SHAMT**: Shift amount (for shift instructions)
- **Funct**: Function code (actual operation)

#### **I-Type Format**

[Opcode] [RS] [RT] [Immediate]  
6 bits 5 5 16 bits

Fields:

- **Opcode**: Varies by instruction
- **RS**: Source/base register
- **RT**: Source/destination register
- **Immediate**: 16-bit immediate value or offset

#### **J-Type Format**

[Opcode] [Address]  
6 bits 26 bits

Fields:

- **Opcode**: 2 for J, 3 for JAL
- **Address**: 26-bit jump target (word address)

### **3.1.3 Digital Logic Review**

#### **Information Encoding**

##### **Binary Representation**

- Low voltage = Logic 0
- High voltage = Logic 1
- Digital signals immune to analog noise

##### **Multi-Bit Signals**

- One wire per bit
- 32-bit instruction needs 32 wires
- Parallel transmission within CPU

## Combinational Elements

### Definition

- Output is function of inputs ONLY
- No internal state or memory
- Purely functional relationship

### Examples

- AND, OR, NOT gates
- Multiplexers:  $Y = (S == 0)?I0 : I1$
- Adders:  $Y = A + B$
- ALU:  $Y = \text{function}(A, B, \text{operation})$

### Characteristics

- Output changes immediately with input (plus propagation delay)
- Can draw complete truth table
- Asynchronous operation (no clock needed)

## Sequential Elements (State Elements)

### Definition

- Output is function of inputs AND internal state
- Has memory—stores information over time
- State persists between clock cycles

### Examples

- Registers
- Flip-flops
- Register files
- Memory units

### Characteristics

- Store information
- Synchronized to clock signal
- Output depends on history

## Clocking and Timing

### Clock Signal

- Periodic alternating signal: Low  $\rightarrow$  High  $\rightarrow$  Low  $\rightarrow$  High...
- Synchronizes all sequential operations

### Edge-Triggered

- Rising edge: Transition 0  $\rightarrow$  1
- Falling edge: Transition 1  $\rightarrow$  0
- Most processors use rising edge

### Clock Period and Frequency

Clock Period (T): Duration of one cycle

Clock Rate (f): Cycles per second

Relationship:  $f = 1/T$

Example:

$T = 250 \text{ ps} = 0.25 \text{ ns}$

$f = 1/(250 \times 10^{-12}) = 4 \text{ GHz}$

## Register Operations

### Basic Register

- Stores multi-bit value (e.g., 32 bits)
- Updates on clock edge: D (input)  $\rightarrow$  Q (output state)

### Timing Example




Figure 3.1: Register Timing Diagram

### Register with Write Control

- Additional Write Enable signal
- Updates ONLY when clock edge AND Write Enable = 1
- Otherwise holds previous value

### Timing Example




Figure 3.2: Register with Write Enable Timing Diagram

### Critical Path and Clock Period

#### Combinational Logic Delay

- All combinational elements have propagation delay
- Different elements, different delays

#### Clock Period Constraint

Clock Period  $\geq$  Longest Path Delay

Path: Register  $\rightarrow$  Combinational Logic  $\rightarrow$  Register

Must allow time for:

1. Register output stabilization
2. Combinational logic computation
3. Result reaching next register input
4. Setup time before next clock edge

#### Critical Path

- Longest delay path from register to register
- Determines minimum clock period
- Limits maximum clock frequency

#### Single-Cycle Constraint

- Complete one instruction per clock cycle
- Clock period must accommodate slowest instruction
- All instructions take same time (inefficient!)

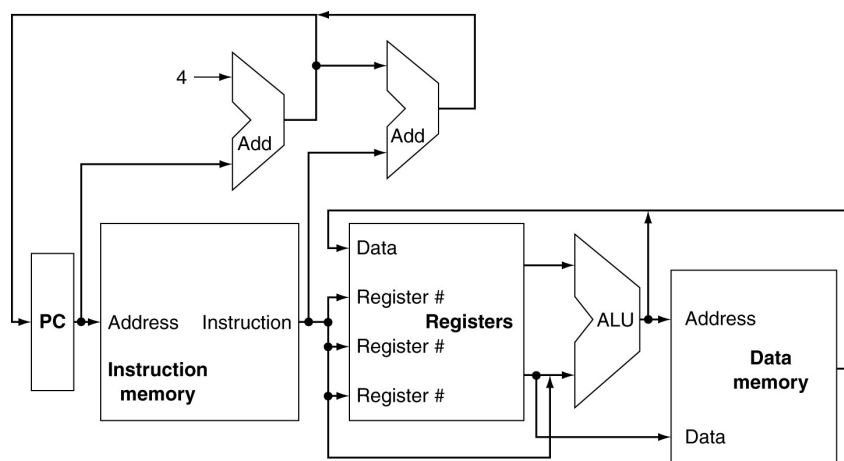


Figure 3.3: CPU Execution Stages Overview

### 3.1.4 CPU Execution Stages

#### Instruction Fetch (IF)

**Purpose:** Retrieve next instruction from memory

**Steps:**

1. Use Program Counter (PC) for instruction address
2. Access Instruction Memory with PC
3. Retrieve 32-bit instruction word
4. Instruction now in CPU for processing

**Hardware:**

- Program Counter (32-bit register)
- Instruction Memory (read-only during execution)
- Address bus from PC to memory
- Data bus from memory to CPU

#### Instruction Decode (ID)

**Purpose:** Interpret instruction and extract fields

**Decode Operations:**

1. **Examine Opcode** (bits 26-31):
  - If opcode = 0: R-type
  - If opcode = 2 or 3: J-type
  - Otherwise: I-type
2. **Extract Register Numbers:**
  - R-type: RS, RT, RD (three 5-bit fields)
  - I-type: RS, RT (two 5-bit fields)

- J-type: No registers

### 3. Extract Immediate/Address:

- I-type: 16-bit immediate
- J-type: 26-bit address

### 4. Extract Function/Shift (R-type only):

- Funct: bits 0-5 (ALU operation)
- SHAMT: bits 6-10 (shift amount)

### Control Unit Role:

- Decodes opcode
- Generates control signals
- Determines datapath activation

## Execute (EX)

**Purpose:** Perform operation or calculate address

### Operations by Type:

#### Arithmetic/Logic (R-type, I-type arithmetic):

- Send operands to ALU
- ALU performs operation
- Operation from funct field (R-type) or opcode (I-type)

### Memory Access (Load/Store):

- ALU calculates address:  $\text{Base} + \text{Offset}$
- Always performs addition
- Result is memory address

### Branch:

- ALU compares registers:  $\text{RS} - \text{RT}$
- Zero flag indicates equality
- Result determines branch decision

**Memory Access (MEM)**

**Purpose:** Read or write data memory

**Applies To:**

- Load instructions: Read from memory
- Store instructions: Write to memory
- NOT arithmetic/logic (skip this stage)

**Load Operation:**

1. Use address from ALU
2. Read data from memory
3. Data will be written to register

**Store Operation:**

1. Use address from ALU
2. Get data from RT register
3. Write data to memory

**Register Write-Back (WB)**

**Purpose:** Write result to destination register

**Applies To:**

- Arithmetic/Logic: Write ALU result
- Load: Write memory data
- NOT store or branch

**Source Selection:**

- Arithmetic/Logic: Data from ALU
- Load: Data from memory
- Multiplexer selects appropriate source

**PC Update**

**Purpose:** Determine next instruction address

**Default:**  $PC = PC + 4$  (sequential)

**Branch/Jump:**  $PC = \text{calculated target address}$

**Control Flow:**

- Multiplexer selects next PC value
- Sequential or branch/jump target
- Update happens at clock edge

### 3.1.5 R-Type Instruction Datapath

#### Register File

##### Structure:

- 32 registers (R0-R31), 32 bits each
- Three ports: 2 read, 1 write

##### Read Ports:

- Read Address 1: RS (5 bits)
- Read Address 2: RT (5 bits)
- Read Data 1: 32-bit output
- Read Data 2: 32-bit output
- Combinational (no clock)

##### Write Port:

- Write Address: RD (5 bits)
- Write Data: 32-bit input
- Write Enable: Control signal
- Synchronized (clock edge)

#### R-Type Execution Flow

**Instruction:** ADD \$t0, \$t1, \$t2 ( $R0 = R1 + R2$ )

##### Step 1: Register Read

- Extract RS (R1) and RT (R2) fields
- Register file outputs two 32-bit values

##### Step 2: ALU Operation

- Inputs: Two register values
- Funct field (6 bits)  $\rightarrow$  ALU control (4 bits)
- ALU performs specified operation
- Examples: ADD, SUB, AND, OR, SLT

##### Step 3: Write-Back

- ALU result  $\rightarrow$  Register file write data
- RD field specifies destination
- Write Enable = 1
- At clock edge: Result written



**ALU Control****Function Field Encoding:**

Funct	Operation	ALU Control
-----	-----	-----
0x20	ADD	0010
0x22	SUB	0110
0x24	AND	0000
0x25	OR	0001
0x2A	SLT	0111

**ALU Control Logic:**

- Input: 6-bit funct field
- Output: 4-bit ALU operation
- Combinational logic (lookup table)

**3.1.6 I-Type Instruction Datapath****Differences from R-Type****Operand Sources:**

- R-type: Both from registers
- I-type: One register, one immediate

**Register Usage:**

- RS: Source register
- RT: Destination register (NOT source!)
- Immediate: 16-bit operand

**Sign Extension****Problem:** 16-bit immediate, 32-bit ALU**Process:**

1. Take 16-bit immediate
2. Examine bit 15 (sign bit)
3. Replicate sign bit to bits 16-31
4. Result: 32-bit signed value

**Examples:**

16-bit: 0x0005 -&gt; 32-bit: 0x00000005 (+5)

16-bit: 0xFFFFB -&gt; 32-bit: 0xFFFFFFF B (-5)

**Hardware:** Simple wire replication (fast)

### Multiplexer for ALU Input

#### ALU Input B Selection:

- Input 0: Register data (RT) for R-type
- Input 1: Sign-extended immediate for I-type
- Select: ALUSrc control signal

#### ALUSrc Signal:

ALUSrc = 0: Use register (R-type, branch)

ALUSrc = 1: Use immediate (I-type)

### 3.1.7 Load/Store Instruction Datapath

#### Address Calculation

**Formula:** Address = Base + Offset

#### Components:

- Base: RS register (32-bit pointer)
- Offset: 16-bit signed immediate (sign-extended)
- ALU: Always performs addition

#### Examples:

```
LW $t1, 8($t0)    # Load from $t0 + 8
SW $t2, -4($sp)   # Store to $sp - 4
```

#### Load Word (LW)

#### Instruction Format:

- RS: Base register
- RT: Destination register
- Immediate: Offset

#### Execution:

1. Read RS (base address)
2. Sign-extend immediate (offset)
3. ALU adds: Address = RS + offset
4. Read data from memory at address
5. Write data to RT register

**Critical Path:** Longest in single-cycle design

- Fetch → Register Read → ALU → Memory → Register Write

**Store Word (SW)****Instruction Format:**

- RS: Base register
- RT: Source register (data to store)
- Immediate: Offset

**Execution:**

1. Read RS (base) and RT (data)
2. ALU calculates address
3. Write RT data to memory at address
4. NO register write-back

**Key Difference:**

- Reads TWO registers (RS and RT)
- Memory write instead of read
- No register write stage

**Data Memory****Interface:**

- Address: From ALU (32 bits)
- Write Data: From RT register
- Read Data: To register file (for loads)

**Control Signals:**

- MemRead: Enable read (LW)
- MemWrite: Enable write (SW)

**Multiplexer for Write-Back:**

- Input 0: ALU result (arithmetic/logic)
- Input 1: Memory data (load)
- Select: MemtoReg signal

### 3.1.8 Branch Instruction Datapath

#### Branch Types

##### BEQ (Branch if Equal):

- Compare RS and RT
- Branch if  $RS == RT$

##### BNE (Branch if Not Equal):

- Compare RS and RT
- Branch if  $RS \neq RT$

#### Branch Target Calculation

##### Components:

1.  $PC + 4$  (next sequential instruction)
2. Offset from immediate (in instructions)
3.  $Target = (PC + 4) + (Offset \times 4)$

##### Why $PC + 4$ ?

- Offset relative to NEXT instruction
- PC already incremented

##### Word to Byte Conversion:

- Immediate: Number of instructions
- Multiply by 4: Byte offset
- Shift left 2 (wire routing, no hardware!)

#### Branch Execution

##### Step 1: Register Comparison

- Read RS and RT
- ALU subtracts:  $RS - RT$
- Generate Zero flag

##### Step 2: Zero Flag Evaluation

- Zero = 1: Values equal
- Zero = 0: Values different

##### Step 3: Target Calculation (parallel)

- Sign-extend immediate

- Shift left 2
- Add to PC + 4

**Step 4: PC Update Decision**

BEQ: PCSrc = Branch AND Zero

BNE: PCSrc = Branch AND NOT(Zero)

**Multiplexer:**

- Input 0: PC + 4 (sequential)
- Input 1: Branch target
- Select: PCSrc

**Sign Extension and Shifting**

**Sign Extension:** Preserves signed offset

- Forward branch: Positive offset
- Backward branch: Negative offset

**Shift Left 2:** Wire routing trick

- Take bits 0-29 of sign-extended value
- Connect to bits 2-31 of result
- Append two zero wires at bits 0-1
- NO actual shifter hardware!

**3.1.9 Complete Single-Cycle Datapath****Integrated Components****Instruction Fetch:**

- PC register
- Instruction memory
- PC + 4 adder

**Register File:**

- 32 registers with 3 ports
- Two read, one write

**ALU:**

- Two 32-bit inputs
- Operation control

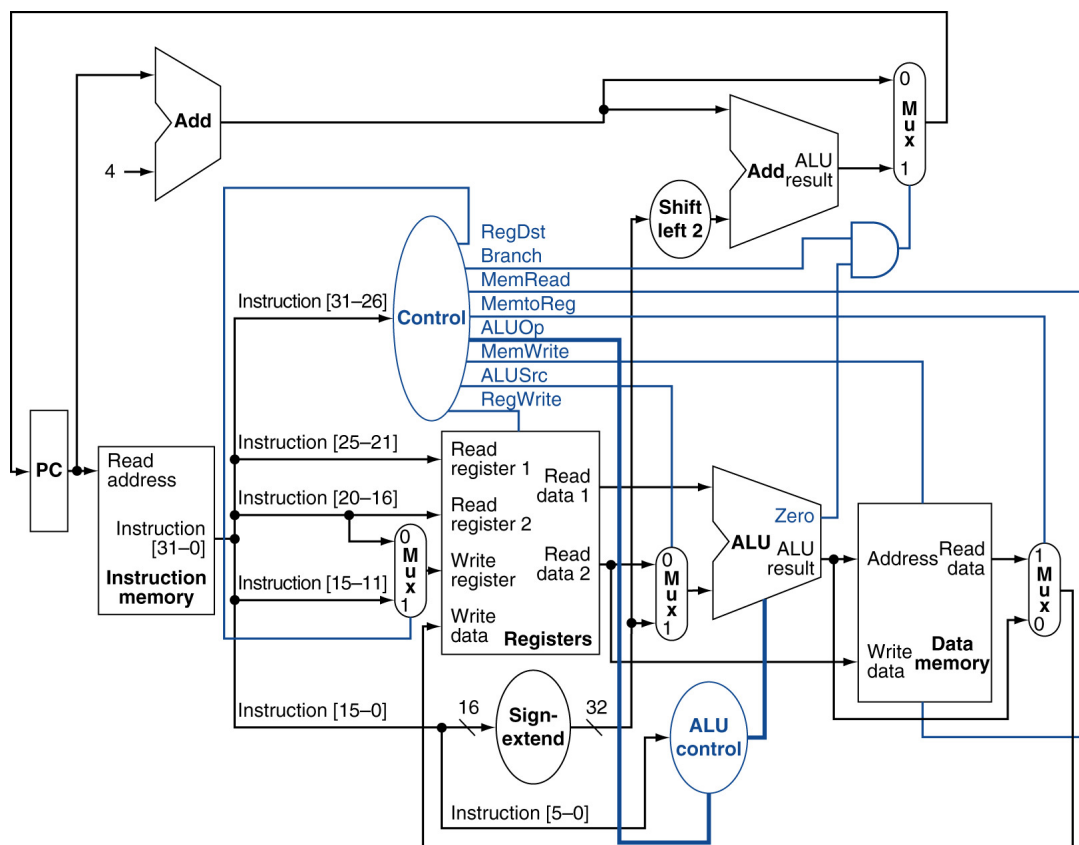


Figure 3.4: Complete Single-Cycle CPU Control and Datapath

- Result output
- Zero flag

**Data Memory:**

- Address from ALU
- Write data from register
- Read data to register

**Sign Extender:**

- 16-bit input
- 32-bit output

**Branch Logic:**

- Target adder
- PC multiplexer

**Multiplexers:**

- ALU input B (register vs immediate)
- Register write data (ALU vs memory)
- Next PC (PC+4 vs branch target)

**Control Signals****Generated by Control Unit:**

1. RegDst: Register destination select
2. Branch: Branch instruction indicator
3. MemRead: Memory read enable
4. MemtoReg: Memory to register select
5. MemWrite: Memory write enable
6. ALUSrc: ALU source select
7. RegWrite: Register write enable
8. ALUOp: ALU operation type

**Parallel Operations**

**Key Insight:** Hardware operates in PARALLEL

- All datapath elements active simultaneously
- Some produce meaningless results
- Control signals select valid paths

**Example:** R-type instruction

- Sign extender operates on bits 0-15
- Produces meaningless output (no immediate in R-type)
- Multiplexer doesn't select it (ALUSrc = 0)

**Critical Path Analysis**

**Path for Load Word** (longest):

1. Instruction fetch:	200 ps
2. Register read:	150 ps
3. Sign extend:	50 ps
4. Multiplexer:	25 ps
5. ALU address calc:	200 ps
6. Data memory access:	200 ps
7. Multiplexer:	25 ps
8. Register write setup:	100 ps
<b>Total:</b>	<b>950 ps</b>

**Clock Period:** Must be  $\geq 950$  ps

**Max Frequency:**  $1/950$  ps  $\approx 1.05$  GHz

**Inefficiency:**

- ALL instructions take 950 ps
- Fast R-type (650 ps) waits
- Wasted time per fast instruction

### Single-Cycle Disadvantages

#### Inefficiency:

- Fast instructions wait for slow ones
- Clock period by worst case
- Cannot optimize common case

#### Hardware Duplication:

- Separate instruction/data memories
- Multiple adders
- Cannot reuse hardware in same cycle

#### No Parallelism:

- One instruction at a time
- Hardware mostly idle
- Poor resource utilization

#### Advantages:

- Simple design
- Simple control
- One instruction per cycle (conceptually)
- Good for learning

### 3.1.10 Key Takeaways

1. **Microarchitecture is hardware implementation of ISA** - translating instruction semantics to hardware operations.
2. **MIPS uses three instruction types:** R-type (registers), I-type (immediate), J-type (jump).
3. **Fixed 32-bit instructions** simplify fetch/decode and enable efficient pipelining.
4. **Combinational elements** have output as function of inputs only; sequential elements have state.
5. **Clock period must exceed longest combinational path** between sequential elements.
6. **Six execution stages:** Fetch, Decode, Execute, Memory, Write-back, PC Update.
7. **Register file has three ports:** two read (combinational), one write (clocked).
8. **Sign extension** converts 16-bit immediate to 32-bit preserving signed value.
9. **Multiplexers select between data sources** based on control signals.



10. **ALU operations vary by instruction:** addition (load/store), subtraction (branch), varies (R-type).
11. **Critical path determines clock period** - load word is longest in single-cycle design.
12. **Single-cycle processor completes one instruction per cycle** but inefficiently (all take same time).
13. **Separate instruction and data memories** required for single-cycle (both accessed same cycle).
14. **Control signals orchestrate datapath** - generated by control unit from opcode.
15. **All hardware operates in parallel** - control signals select valid results, ignore others.

### 3.1.11 Summary

Microarchitecture bridges the gap between software instructions and hardware implementation, revealing how processors execute programs. Building a single-cycle MIPS processor requires understanding digital logic fundamentals, datapath component design, and control signal generation. While conceptually simple (one instruction per cycle), the single-cycle design is inefficient because all instructions must complete within the time required by the slowest instruction. The critical path—typically the load word instruction—determines the maximum clock frequency. Understanding this foundation prepares us for more sophisticated designs including multi-cycle processors (which break execution into multiple stages) and pipelined processors (which overlap instruction execution for higher throughput). These microarchitecture concepts apply broadly across processor design, from embedded systems to high-performance superscalar processors.

## 3.2 Lecture 10: Processor Control

*By Dr. Isuru Nawinne*

### 3.2.1 Introduction

This lecture completes the single-cycle MIPS processor design by exploring the control unit—the component that generates control signals based on instruction opcodes. We examine ALU control generation using a two-stage approach, design the main control unit, analyze control signal purposes, and create truth tables mapping instructions to control patterns. Understanding control unit design reveals how hardware interprets instructions and orchestrates datapath operations, completing our understanding of processor implementation.

### 3.2.2 Control Unit Overview

#### Recap of Datapath Components

#### Previously Covered:

- Register File (32 registers, 3 ports)
- ALU (arithmetic/logic operations)
- Instruction Memory (stores program)
- Data Memory (stores data)
- Adders (PC+4, branch target)

- Multiplexers (data source selection)
- Sign Extender (16-bit to 32-bit)
- Shifter (branch offset left 2)

### Control Unit Purpose

**Function:** Generate control signals based on instruction

**Inputs:**

- Opcode (bits 26-31, 6 bits)
- Funct field (bits 0-5, 6 bits) for R-type

**Outputs:** Control signals for datapath

- Multiplexer selections
- Register write enable
- Memory read/write
- ALU operation
- Branch decision

### Instruction Subset for Study

**Selected Instructions:**

- **Load Word (LW):** Memory read
- **Store Word (SW):** Memory write
- **Branch if Equal (BEQ):** Conditional branch
- **R-type:** Arithmetic, logic, shift

**Coverage:**

- Uses almost all datapath hardware
- Representative of most control signals
- Excludes: Jump instructions, I-type arithmetic

### 3.2.3 ALU Operations for Different Instructions

#### Load/Store Instructions

**Address Calculation:**

Address = Base Register + Immediate Offset  
           = RS + Sign\_Extend(Immediate)

**ALU Function:** ADDITION (always)

- Input A: RS register value

- Input B: Sign-extended immediate
- Operation: ADD
- ALU Control: 0010 (binary)
- Result: Memory address

**Example:**

```
LW $t1, 8($t0)    # Address = $t0 + 8
SW $t2, -4($sp)   # Address = $sp + (-4)
```

**Branch Instructions****Comparison Operation:**

Compare RS and RT for equality  
 Method: Subtract RT from RS

**ALU Function: SUBTRACTION**

- Input A: RS register value
- Input B: RT register value
- Operation: SUB
- ALU Control: 0110 (binary)
- Result: RS - RT
- Zero Flag: Indicates if result is zero (equal)

**Branch Decision:**

Zero = 1: RS == RT, take branch  
 Zero = 0: RS != RT, don't take branch

**R-Type Instructions**

**Variable Operations:** Determined by funct field

**ALU Function: DEPENDS ON FUNCT**

- Input A: RS register value
- Input B: RT register value
- Operation: From funct field
- ALU Control: Varies
- Result: Written to RD register

**Funct Field Mapping:**

Funct	Operation	ALU Control
----- ----- -----		
0x20	ADD	0010
0x22	SUB	0110
0x24	AND	0000
0x25	OR	0001
0x2A	SLT	0111

### 3.2.4 ALU Control Signal

#### Signal Format

**4-Bit Signal:** Specifies ALU operation

**Possible Operations** ( $2^4 = 16$ ):

0000: AND  
0001: OR  
0010: ADD  
0110: SUBTRACT  
0111: Set on Less Than (SLT)  
1100: NOR

#### Usage:

- Not all 16 combinations used
- Could use 3 bits for 8 operations
- 4-bit standard allows expansion

#### Control Signal Usage by Instruction

##### Load/Store:

- ALU Control = 0010 (ADD)
- Fixed operation
- Independent of instruction specifics

##### Branch:

- ALU Control = 0110 (SUBTRACT)
- Fixed operation
- Zero flag is critical output

##### R-Type:

- ALU Control = Varies
- Must decode funct field
- Different operations need different controls

### 3.2.5 Two-Stage ALU Control Generation

#### Design Rationale

##### Why Two Stages?

##### Efficiency:

- Some instructions don't need funct field
- Separates opcode-level from operation-level

- Faster for non-R-type instructions

#### Timing Optimization:

- Other control signals needed faster
- Examples: Register addressing, immediate routing
- ALU control can afford slight delay

#### Modularity:

- Stage 1: Main control (opcode-based)
- Stage 2: ALU control (operation-specific)
- Cleaner design separation

### Stage 1: Generate ALUOp

**Input:** Opcode (6 bits)

**Output:** ALUOp (2 bits)

**Encoding:**

Instruction	Opcode	ALUOp
-----	-----	-----
Load Word	100011	00
Store Word	101011	00
Branch Equal	000100	01
R-type	000000	10

#### ALUOp Meaning:

- **00:** Perform ADD (address calculation)
- **01:** Perform SUBTRACT (comparison)
- **10:** Operation from funct field

**Logic:** Purely combinational based on opcode

### Stage 2: Generate ALU Control

**Inputs:**

- ALUOp (2 bits from Stage 1)
- Funct field (6 bits from instruction)
- Total: 8 input bits

**Output:** ALU Control (4 bits)

**Truth Table:**

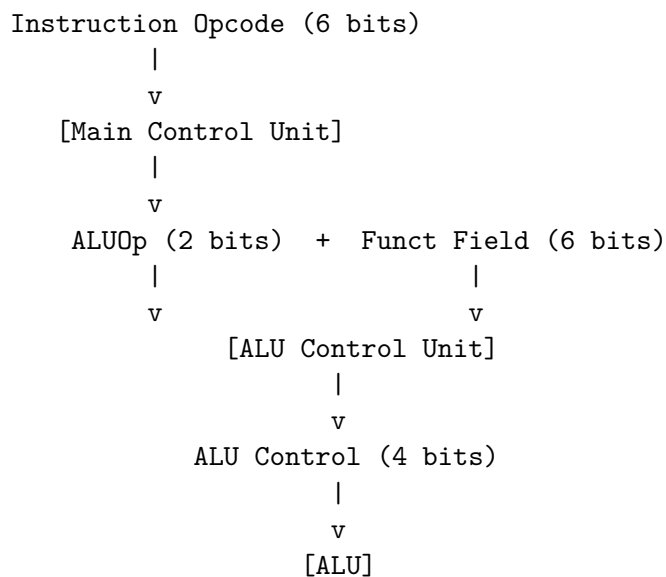
ALUOp	Func	ALU Control	Operation
-----	-----	-----	-----
00	XXXXXX	0010	ADD (LW/SW)
01	XXXXXX	0110	SUB (BEQ)
10	100000	0010	ADD (R-type)
10	100010	0110	SUB (R-type)
10	100100	0000	AND
10	100101	0001	OR
10	101010	0111	SLT

**"X" Notation:** Don't Care

- For ALUOp = 00 or 01, funct irrelevant
- Simplifies logic design
- Reduces gate count

### Complete ALU Control Path

**Flow Diagram:**



**Advantages:**

- Modular design
- Simplified main control
- Localized R-type complexity
- Easier to verify

### 3.2.6 Main Control Signals

**Complete Signal List**

**Signals Generated:**

1. **RegDst** (1 bit): Register destination select

2. **Branch** (1 bit): Branch instruction indicator
3. **MemRead** (1 bit): Memory read enable
4. **MemtoReg** (1 bit): Memory to register select
5. **MemWrite** (1 bit): Memory write enable
6. **ALUSrc** (1 bit): ALU source select
7. **RegWrite** (1 bit): Register write enable
8. **ALUOp** (2 bits): To ALU control unit

**Total:** 9 control bits from main control

### RegDst (Register Destination)

**Purpose:** Select which field specifies write destination

**Multiplexer Control:**

- Input 0: RT field (bits 16-20)
- Input 1: RD field (bits 11-15)
- Output: Register write address (5 bits)

**Settings:**

RegDst = 0: Write to RT (Load Word)

RegDst = 1: Write to RD (R-type)

**Rationale:**

- Load Word: RT is destination (I-type format)
- R-type: RD is destination (R-type format)
- Store/Branch: Don't care (no write)

**Examples:**

```
LW $t1, 8($t0)      # Write to $t1 (RT) -> RegDst = 0
ADD $t2, $t3, $t4    # Write to $t2 (RD) -> RegDst = 1
```

### Branch

**Purpose:** Indicate if instruction is branch

**Usage:** Combined with Zero flag for PC selection

**Settings:**

Branch = 0: Not a branch (LW, SW, R-type)

Branch = 1: Branch instruction (BEQ, BNE)

**PC Selection Logic:**

For BEQ:

PCSrc = Branch AND Zero

(Take branch if instruction is branch AND comparison equal)

For BNE:

PCSrc = Branch AND NOT(Zero)

(Take branch if instruction is branch AND comparison not equal)

**MemRead**

**Purpose:** Enable reading from data memory

**Settings:**

MemRead = 0: No memory read (R-type, SW, BEQ)

MemRead = 1: Read from memory (LW)

**Function:**

- Controls data memory read enable
- When high: Memory outputs data
- When low: Memory read inactive

**MemtoReg (Memory to Register)**

**Purpose:** Select source of register write data

**Multiplexer Control:**

- Input 0: ALU result
- Input 1: Data memory read data
- Output: Register write data (32 bits)

**Settings:**

MemtoReg = 0: Write ALU result (R-type)

MemtoReg = 1: Write memory data (LW)

**Examples:**

ADD \$t1, \$t2, \$t3 # \$t1 = ALU result -> MemtoReg = 0

LW \$t1, 8(\$t0) # \$t1 = memory data -> MemtoReg = 1

**MemWrite**

**Purpose:** Enable writing to data memory

**Settings:**

MemWrite = 0: No memory write (R-type, LW, BEQ)

MemWrite = 1: Write to memory (SW)

**Function:**

- Controls data memory write enable
- When high: Data written (on clock edge)
- When low: Memory write disabled



**ALUSrc (ALU Source)****Purpose:** Select second ALU operand source**Multiplexer Control:**

- Input 0: Register file Read Data 2 (RT value)
- Input 1: Sign-extended immediate
- Output: ALU Input B (32 bits)

**Settings:**

ALUSrc = 0: Use register (R-type, BEQ)

ALUSrc = 1: Use immediate (LW, SW)

**Examples:**

```
ADD $t1, $t2, $t3  # Use $t3 -> ALUSrc = 0
LW $t1, 8($t0)     # Use imm 8 -> ALUSrc = 1
```

**RegWrite****Purpose:** Enable writing to register file**Settings:**

RegWrite = 0: No register write (SW, BEQ)

RegWrite = 1: Write to register (R-type, LW)

**Usage by Instruction:**

```
R-type:      RegWrite = 1 (write ALU result)
Load Word:   RegWrite = 1 (write memory data)
Store Word:   RegWrite = 0 (no write)
Branch:      RegWrite = 0 (no write)
```

**3.2.7 Control Signal Truth Table****Complete Table**

Instruction	RegDst	ALUSrc	MemtoReg	RegWrite	MemRead	MemWrite	Branch	A
R-type	1	0	0	1	0	0	0	
Load Word	0	1	1	1	1	0	0	
Store Word	X	1	X	0	0	1	0	
Branch Eq	X	0	X	0	0	0	1	

Table 3.1: Control Signal Truth Table

**Legend:**

- **0**: Signal low/false/select input 0
- **1**: Signal high/true/select input 1
- **X**: Don't Care (not used, can be anything)

**R-Type Control****Settings:**

RegDst = 1:       Write to RD field  
 ALUSrc = 0:       Second operand from register (RT)  
 MemtoReg = 0:     Write ALU result  
 RegWrite = 1:     Enable register write  
 MemRead = 0:       No memory read  
 MemWrite = 0:      No memory write  
 Branch = 0:        Not a branch  
 ALUOp = 10:        Consult funct field

**Active Elements:**

- Instruction fetch
- Register file (read RS, RT; write RD)
- ALU (operation from funct)
- Register write from ALU
- PC updated to PC + 4

**Inactive Elements:**

- Data memory (not accessed)
- Branch target (computed but not used)
- Sign extender (operates but ignored)

**Load Word Control****Settings:**

RegDst = 0:       Write to RT field  
 ALUSrc = 1:       Second operand from immediate  
 MemtoReg = 1:     Write memory data  
 RegWrite = 1:     Enable register write  
 MemRead = 1:       Enable memory read  
 MemWrite = 0:      No memory write  
 Branch = 0:        Not a branch  
 ALUOp = 00:        ALU performs ADD

**Active Elements:**

- Instruction fetch
- Register file (read RS; write RT)
- Sign extender
- ALU (ADD for address)
- Data memory (read)

- Register write from memory
- PC updated to  $PC + 4$

**Critical Path:** Longest delay

- Fetch  $\rightarrow$  Reg Read  $\rightarrow$  Sign Extend  $\rightarrow$  ALU  $\rightarrow$  Memory  $\rightarrow$  Reg Write

### Store Word Control

**Settings:**

RegDst = X: Don't care (no register write)  
 ALUSrc = 1: Second operand from immediate  
 MemtoReg = X: Don't care (no register write)  
 RegWrite = 0: No register write  
 MemRead = 0: No memory read  
 MemWrite = 1: Enable memory write  
 Branch = 0: Not a branch  
 ALUOp = 00: ALU performs ADD

**Key Difference from Load:**

- Read TWO registers (RS for base, RT for data)
- Memory write instead of read
- No register write stage

### Branch if Equal Control

**Settings:**

RegDst = X: Don't care (no register write)  
 ALUSrc = 0: Second operand from register (RT)  
 MemtoReg = X: Don't care (no register write)  
 RegWrite = 0: No register write  
 MemRead = 0: No memory read  
 MemWrite = 0: No memory write  
 Branch = 1: This is a branch  
 ALUOp = 01: ALU performs SUBTRACT

**Active Elements:**

- Instruction fetch
- Register file (read RS, RT)
- ALU (SUBTRACT for comparison, Zero flag)
- Sign extender + shift (branch target)
- Branch target adder ( $PC + 4 + \text{offset}$ )
- PC multiplexer (select based on Branch AND Zero)

**Branch Decision Logic:**

```

Zero = (RS - RT == 0)
PCSrc = Branch AND Zero
If PCSrc:
    Next PC = PC + 4 + (SignExtend(Imm) << 2)
Else:
    Next PC = PC + 4

```

### 3.2.8 Control Unit Implementation

#### Input to Control Unit

**Primary Input:** Opcode (bits 26-31, 6 bits)

- Identifies instruction type
- Determines all control signal values

**Secondary Input:** Funct field (bits 0-5, 6 bits)

- Only for R-type (opcode = 000000)
- Specifies ALU operation

#### Combinational Logic Design

**Method:** Standard digital logic techniques

**Steps:**

1. Create truth table (opcode  $\rightarrow$  control signals)
2. List all control signals as outputs
3. Fill in values for each instruction
4. Use Karnaugh maps or Boolean algebra to minimize
5. Implement with logic gates

**Example for RegWrite:**

```

RegWrite = (R-type) OR (Load Word)
RegWrite = (opcode == 000000) OR (opcode == 100011)

```

#### Control Unit Structure

**ROM-Based Implementation:**

- Opcode as ROM address
- ROM location stores control pattern
- Simple but inflexible

**PLA (Programmable Logic Array):**

- Implements minimized logic equations
- More efficient than ROM

- Standard for simple processors

**Hardwired Logic:**

- Custom logic gates
- Fastest implementation
- Most common for high-performance

**Microcode** (not typical for RISC):

- Control signals stored in memory
- More flexible but slower
- Used in CISC (e.g., x86)

**Timing Considerations****Signal Generation Time:**

- Must complete early in clock cycle
- Before datapath elements need signals
- Critical for clock frequency

**Signal Stability:**

- Must remain stable throughout cycle
- Changes only between instructions
- Combinational logic ensures this

**Clock Period Impact:**

- Control logic adds delay
- Typically small vs. ALU/memory
- Well-designed control has minimal impact

**3.2.9 Why Separate MemRead and MemWrite?****Initial Observation**

**Question:** Seem mutually exclusive—why not one signal?

- Could use: 0 = Read, 1 = Write
- Appears redundant

**Answer: Yes, Separate Signals Needed**

**Timing Control:**

- Write Enable: Specifies WHEN to write
- Read Enable: Specifies WHEN valid data available
- Different timing requirements

**No Operation State:**

- Both = 0: No memory access
- Common for R-type and branch
- Single signal couldn't represent this

**Three States Required:**

MemRead=1, MemWrite=0: Read

MemRead=0, MemWrite=1: Write

MemRead=0, MemWrite=0: No access

(MemRead=1, MemWrite=1: Invalid)

**Future: Pipelined Processors**

**Concurrent Access:**

- Different pipeline stages access memory
- One stage reading, another writing
- Separate signals essential

**Memory Banking:**

- Separate read/write ports
- Enables simultaneous access
- Separate signals control independent ports

**Design Philosophy**

**Orthogonality:**

- Each signal controls independent function
- Easier to understand and verify
- Reduces design errors

**Flexibility:**

- Supports future enhancements
- Allows memory optimization
- Standard practice

### 3.2.10 Complete Datapath with Control

#### Integrated System

##### Components Connected:

- Control Unit (generates signals)
- Datapath (executes operations)
- Blue lines: Control signals
- Black lines: Data paths

##### Control Unit Connections:

- Input: Instruction opcode
- Outputs: All control signals
- Fan out to datapath elements

##### ALU Control Unit:

- Separate box near ALU
- Inputs: ALUOp, Funct
- Output: ALU Control (4 bits)

#### Example: Load Word Execution

**Instruction:** LW \$t1, 8(\$t0)

##### Step 1: Fetch

PC -> Instruction Memory  
Opcode = 100011 (LW)

##### Step 2: Control Signals

RegDst=0, ALUSrc=1, MemtoReg=1, RegWrite=1,  
MemRead=1, MemWrite=0, Branch=0, ALUOp=00

##### Step 3: Register Read

RS field (\$t0) -> Register file  
Read Data 1 = \$t0 value

##### Step 4: ALU

Immediate = 8  
Sign-extended to 32 bits  
ALUSrc=1: Selects immediate  
ALU performs ADD: \$t0 + 8 = address

##### Step 5: Memory

MemRead=1: Memory reads at address  
Data output from memory

### Step 6: Write-Back

MemtoReg=1: Selects memory data  
RegDst=0: Selects RT (\$t1)  
RegWrite=1: Enables write  
At clock edge: Memory data -> \$t1

### Step 7: PC Update

Branch=0: PCSrc=0  
PC updated to PC + 4

### 3.2.11 Key Takeaways

1. **Control unit generates signals based on instruction opcode**, orchestrating datapath operations.
2. **ALU control uses two-stage generation**: Opcode  $\rightarrow$  ALUOp (2 bits)  $\rightarrow$  ALU Control (4 bits).
3. **Stage 1 (Main Control)**: Opcode to ALUOp - identifies operation category.
4. **Stage 2 (ALU Control)**: ALUOp + Funct to ALU Control - specifies exact operation.
5. **Two-stage design optimizes timing and modularity**, separating concerns.
6. **Main control signals**: RegDst, Branch, MemRead, MemtoReg, MemWrite, ALUSrc, RegWrite, ALUOp.
7. **Load/Store always use ADD** for address calculation, regardless of other details.
8. **Branch uses SUBTRACT** for comparison, with Zero flag indicating equality.
9. **R-type ALU operation from funct field**, providing operation flexibility.
10. **Instruction format regularity simplifies control**, with consistent field positions.
11. **Register roles vary by instruction type**, especially RT (destination vs. source).
12. **Control signals mutually exclusive** for proper operation - only valid combinations used.
13. **Separate MemRead/MemWrite needed** for no-op state and future pipelining.
14. **Control logic is combinational** (no state), generating signals each cycle.
15. **Truth tables map opcode to control patterns**, enabling systematic design.
16. **"Don't care" values simplify logic minimization**, reducing gate count.
17. **Control unit design uses standard digital logic techniques**, including K-maps and Boolean algebra.
18. **Datapath elements may operate but outputs ignored** if not selected by control signals.



19. **Complete processor integrates datapath and control**, with control signals orchestrating all operations.
20. **Single-cycle design simple but inefficient** - foundation for advanced multi-cycle and pipelined designs.

### 3.2.12 Summary

The control unit completes the single-cycle MIPS processor, generating control signals that orchestrate datapath operations based on instruction opcodes. The two-stage ALU control generation (opcode  $\rightarrow$  ALUOp  $\rightarrow$  ALU Control) elegantly separates concerns, with the main control handling instruction-level decisions and the ALU control handling operation-specific details. Each control signal serves a specific purpose, from selecting multiplexer inputs (RegDst, ALUSrc, MemtoReg) to enabling register and memory operations (RegWrite, MemRead, MemWrite) to handling branches (Branch). Truth tables systematically map instructions to control patterns, with "don't care" values simplifying logic design. While the single-cycle processor provides conceptual clarity and simplicity, its inefficiency (all instructions taking the same time as the slowest) motivates more sophisticated designs. Understanding this foundation prepares us for multi-cycle processors (which break execution into variable-length stages) and pipelined processors (which overlap instruction execution for higher throughput), both building on the control principles established here.